



Department of Computer Science

Lab Manual

of

MACHINE LEARNING

MCA521-4

Class Name: IV MCA

Master of Computer Application

2026

Prepared by: Alok Tayal

Verified by:

Dr. Rupali Sunil Wagh Dr. Helen K Joy Dr. Shivangi Singh

Department Overview

Department of Computer Science of CHRIST (Deemed to be University) strives to shape outstanding computer professionals with ethical and human values to reshape nation's destiny. The training imparted aims to prepare young minds for the challenging opportunities in the IT industry with a global awareness rooted in the Indian soil, nourished and supported by experts in the field.

Vision

The Department of Computer Science endeavours to imbibe the vision of the University “**Excellence and Service**”. The department is committed to this philosophy which pervades every aspect and functioning of the department.

Mission

“To develop IT professionals with ethical and human values”. To accomplish our mission, the department encourages students to apply their acquired knowledge and skills towards professional achievements in their career. The department also moulds the students to be socially responsible and ethically sound.

Introduction to the Programme

Master of Computer Applications (MCA) is a post-graduate programme of CHRIST (Deemed to be University) and approved by the All India Council of Technical Education (AICTE). It is a two-year programme spread over six trimesters to bridge the chasm between computer studies and applications, bringing within reach of enthusiastic youngsters an excellent group of experienced and dedicated staff and experts, and an enviable infrastructure. MCA at CHRIST (Deemed to be University) strives to shape outstanding computer professionals with ethical and human values to reshape the nation's destiny. The training programme aims to prepare young minds for challenging opportunities in the IT industry with a global awareness rooted in the Indian soil, nourished and supported by experts in the field. The Department is committed to the motto “Excellence and Service,” and this philosophy pervades every aspect of its functioning.

Programme Objective

This programme is conceptualised and designed based on the strong commitment of the department to provide better quality education to the students. The principal objectives of this course are to:

- Heighten technological awareness
- Train future industry specialists
- Encourage effective software development
- Provide research training
- Involve students in innovative research
- Produce graduates with a two-year professional education in Computer Science with technical, professional, and communications skills.

Ethics and Human Values

1. Only proprietary or open-source software would be used for academic teaching and learning purposes.
2. Copying of programs from the internet, friends or from other sources is strictly discouraged since it impairs development of programming skills.
3. Unique Practical (Domain based) exercises ensure that the students don't get involved in code plagiarism.
4. Projects undertaken by students during the course are done in teams to improve collaborative work and synergy between team members.
5. Projects involve modularization which initiates students to take individual responsibility for common goals.
6. Passion for excellence is promoted among the students, be it in software development or project documentation.
7. Giving due credit to sources during the seminar and research assignment is promoted among the students
8. The course and its design enforce the practice of good referencing technique to improve the sense of integrity.
9. Courses involving group discussions and debates on ethical practices and human values are designed to sensitize the students in dealing with customers and members within the Organization.

Programme Outcomes:

P01: Foundation Knowledge: Apply knowledge of mathematics, programming logic, and coding fundamentals for solution architecture and problem-solving.

P02: Problem Analysis: Identify, review, formulate, and analyse problems primarily focussing on customer requirements using critical thinking frameworks.

P03: Development of Solutions: Design, develop and investigate problems with as an innovative approach for solutions incorporating ESG/SDG goals.

P04: Modern Tool Usage: Select, adapt, and apply modern computational tools such as the development of algorithms with an understanding of the limitations including human biases.

P05: Individual and Teamwork: Function and communicate effectively as an individual or a team leader in diverse and multidisciplinary groups. Use methodologies such as agile.

P06: Project Management and Finance: Use the principles of project management such as scheduling, and work breakdown structure and be conversant with the principles of Finance for profitable project management.

P07: Ethics: Commit to professional ethics in managing software projects with financial aspects.

Learn to use new technologies for cyber security and insulate customers from malware

P08: Life-long learning: Change management skills and the ability to learn, keep up with contemporary technologies and ways of working.

Programme Educational Objectives(PEO's)

PEO1: Develop innovative and effective software solutions through research that addresses real-world challenges, grounded in a commitment to continuous learning and adaptability.

PEO2: Advance the knowledge and practices in the field of computer applications through lifelong learning and rigorous research, supporting professional growth and academic excellence.

PEO3: Exhibit ethical leadership, social responsibility to contribute and enhance community well-being by innovating solutions.

MCA521-4 – MACHINE LEARNING

Course Name: MACHINE LEARNING Code: MCA521-4

Total Hours: 75

L + P + T: 3 + 4 + 0

Max Marks: 100

Credits: 4

Course Type: Core Course

Course Category: Theo&Prac

Course Description

Machine Learning is a key area in computer applications that focuses on building models capable of making predictions or informed decisions based on data. A strong understanding of machine learning enables students to analyze and interpret complex datasets, develop predictive models, and extract meaningful insights. This course covers a wide range of machine learning concepts and techniques, including supervised and unsupervised learning. It provides a solid theoretical foundation while emphasizing practical, hands-on experience with algorithms and real-world data. The course equips students with the knowledge and skills required to apply machine learning techniques across various domains using modern tools and methodologies.

Course Objectives

This course enables students to understand the fundamental concepts of Machine Learning, including its core principles, terminology, and methodologies. Students will learn to differentiate between various types of learning algorithms such as supervised, unsupervised, and reinforcement learning, and recognize their appropriate applications.

Course Outcomes

After successful completion of the course, students will be able to:

- Explain the fundamental principles and scope of Machine Learning
- Implement modelling practices in machine learning for better generalisation
- Interpret predictive modelling results and performance of supervised machine learning techniques
- Assess the outcomes and effects of unsupervised machine learning processes
- Deploy robust machine learning models for interdisciplinary applications

Unit-1: INTRODUCTION TO MACHINE LEARNING Teaching Hours: 15 Hours

Introduction to AI – Knowledge Representation – Data – Representation of Data, Data Processing, Data Storage, Data Processing, Types of Data – Exploratory Data Analysis, Visualisations and Interpretation, Outliers

Machine Learning: Definition, Objectives, Components, Features, Applications – Traditional Programming v/s Machine Learning, Types of Machine Learning – Predictive Models, Statistical Inferences – Applications of Machine Learning

Challenges in ML: Insufficient Quantity of Data, Non-representative and Poor-Quality Data, Irrelevant Features, Estimation Estimation, Reducing Human Biases in Model Building, Ethical Use of Technology

Lab Exercises:

1. Data Preprocessing and visualisation
2. Data exploration and inferences

Unit-2: PROCESSES IN MACHINE LEARNING Teaching Hours: 15 Hours

Data Preparation and Model Selection: Effect of Data Preparation: Encoding, Nominal and Ordinal Representation, Feature Transformation and Feature Engineering, Generalisation, Normalisation, Sampling, Using Saved Weights, Model Selection Strategies: Overfitting, Underfitting, Bias-Variance Trade-off, Testing and Validation: Performance measures - Confusion matrix, Precision-Recall Trade off, F1 Score, ROC Curve, AUC, Error Analysis, Model Parameters, Hyperparameters, Hyperparameter Tuning, Cross Validation, Decision Functions: Introduction to Supervised Machine Learning Methods, Decision Functions, Decision Thresholds, Distance Measures, Hyperplanes, Regression & Classification Problems, Loss and Cost Functions, Linear Regression using OLS, Multi-class vs Multi-label, KNN Classification

Lab Exercises:

3. Saving weights of a generalised model.
4. Regression and Classification Evaluation Metrics: Illustration through Linear Regression and KNN Classification

Unit-3: SUPERVISED LEARNING ALGORITHMS

Teaching Hours: 15

Hours

Learning through Optimisation: Gradient Descent, Linear Regression, Logistic Regression
Computational Learning: Decision Trees, Naive-bayes classification, Support Vector Machines, Ensemble Learners: Learning, Voting, Bagging, Stacking, Boosting, Random Forests, GridSearchCV

Lab Exercises:

5. Regression through Gradient Descent
6. Comparison of Different Classifiers
7. Illustration of Ensemble Learning Methods

Unit-4: UNSUPERVISED LEARNING AND DIMENSIONALITY REDUCTION

Teaching

Hours: 15 Hours

Introduction: Types of Unsupervised Learning, Challenges in Unsupervised Learning, Concept of Cost Functions in Unsupervised Learning, Clustering Methods: Distance based, Centroid Based, Elbow method to find Optimal Number of Clusters, Silhouette Method, K-Means Clustering, Hierarchical Clustering: Agglomerative and Bottom Up, Applying Supervised Learning after Clustering, Dimensionality Reduction Methods: Principal Component Analysis (PCA), Linear Discriminant Analysis (LDA)

Lab Exercises:

8. KMeans and Elbow Methods
9. Hierarchical Clustering and Dendrograms
10. Dimensionality Reduction
11. Image Compression using Dimensionality Reduction.

Unit-5 Teaching Hours: 15 Hours

MACHINE LEARNING IN REAL-WORLD

Emerging Techniques: Role of ML in attaining Sustainable Development Goals, Time Series Preprocessing, Blackbox Methods: Neural Networks & Deep Learning, Reinforcement Learning/QLearning, Self Supervised Learning, Quantum Machine Learning, Keras and Tensorflow

for designing Multi Layer Perceptron. Case Studies: CNN, RNN, Advances in Deep Learning, Natural Language Processing, Real Time Systems, Computer Vision, Robotics, Expert Systems, Generative Networks, Large Language Models. Model Deployment: Model Deployment: Connecting Models to User Interface, Deployment of ML Models.

Lab Exercises:

12. Training a Multi-layer perceptron

13. Deploying Trained ML-models

Text and Reference Books

[1] Campesato, O. (2020). Artificial Intelligence, Machine Learning and Deep Learning. Mercury Learning and Information.

[2] Andreas C. Müller and Sarah Guido (2017), “Introduction to Machine Learning with Python A Guide For Data Scientists” O’Reilly book

[3] Ethem Alpaydin (2005), “Introduction to Machine Learning”, Prentice Hall of India

[4] Max Kuhn and Kjell Johnson (2018), “Applied Predictive Modelling”, Springer

Essential Reading

[1] Stuart Russell and Peter Norvig(2020), “Artificial Intelligence: A Modern Approach” (4th Edition), Pearson

[2] Aurelien Geron(2019), “Hands on Machine Learning with Scikit-learn, Keras and Tensorflow”, 2nd Edition, O’Reilly Books

[3] Kevin P. Murphy(2012), “Machine Learning: A Probabilistic Perspective”, MIT Press

[4] Hastie, Tibshirani, Friedman(2008), “The Elements of Statistical Learning” (2nd ed), Springer

[5] Ian Goodfellow, Aaron Courville, Yoshua Bengio (2019), “Deep Learning (Adaptive Computation And Machine Learning Series)”, MIT Press

Additional Information (Web Resources):

1. Andrew Ng, Deeplearning.ai, Machine Learning Specialisation, offered through Coursera:

<https://www.deeplearning.ai/courses/machine-learning-specialization/>

Evaluation Pattern: CIA - 50% ESE - 50%

LIST OF PROGRAMS

MCA				
Sl. no	Title of lab Experiment	Page number	RB T	CO
1	To investigate whether worsening air quality across Indian states is linked to declining agricultural output by independently choosing, applying, and justifying appropriate data analysis and visualisation techniques on raw, messy, real-world datasets.	8	L3	CO1, 3
2	To investigate whether worsening air quality across Indian states is linked to declining agricultural output by independently choosing, applying, and justifying appropriate data analysis and visualisation techniques on raw, messy, real-world datasets.	28	L3	CO2, 3
3	Student Survey: Simple Linear Regression using Scikit Learn, OLS, and Parameter Saving.	44	L3	CO3
4	Regression and Classification Evaluation Metrics	66	L3	CO3
5	Linear Regression through Gradient Descent	86	L3	CO3
6	Comparison of Logistic Regression and K Nearest Neighbors (KNN) Classifiers	97	L3	CO3, 4
7	Decision Tree Classifier	115	L4	CO3
8	Implementation and Performance Evaluation of Categorical Naive Bayes Classifier	131	L3	CO3
9	Support Vector Machine (SVM) and Principal Component Analysis (PCA)	144	L3	CO4
10	Learning the XOR Boolean Function Using an MLP	160	L5	CO4

Evaluation Rubrics:

Evaluation Rubrics for each program would be

Submission Guidelines:

- Make a copy of the lab manual template with your <name_reg:no_subject name > ,
- Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as lab number.
- Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
- Create a Git Repository in your profile <ML lab-reg no> . Follow a different branch for each lab <Lab 1, Lab 2...>, and push the code to Git. The link should be provided in Google Classroom along with the PDF of the lab manual.
- Upload the PDF to Google Classroom before the deadline.

Lab 1

Lab Exercise I:

Aim

- To investigate whether worsening air quality across Indian states is linked to declining agricultural output by independently choosing, applying, and justifying appropriate data analysis and visualisation techniques on raw, messy, real-world datasets.
-

Evaluation Rubrics

Submission Guidelines

1. Make a copy of the lab manual template with your `<name_reg:no_subject name>`
2. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
3. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
4. Create a **Git repository** in your profile `<ML lab-reg no>`. Follow a different branch for each lab `<Lab 1, Lab 2 ...>` and push the code to Git.
5. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
6. Upload the PDF to Google Classroom before the deadline.

Code with Results

Task 1

A junior analyst hands you two CSV files and says — "I have no idea what's in these. We need to use them to build a machine learning model next week." Before any analysis can begin, you need to understand what you are working with.

Investigate both files thoroughly. Produce a structured summary that gives a complete first-impression picture of the data — its size, its contents, and where it might be problematic. Decide what information a data scientist would need before trusting this data, and make sure your summary covers all of it.

A structured data profile for each file in a markdown cell

At least one written observation about what concerns you after the inspection

Think: What does a data scientist need to know about a dataset before using it? What functions help you find that? 2 marks Completeness of profile + written concern

```
> 
import pandas as pd
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
[2] ✓ 1.1s
```

```
city = pd.read_csv('city_day.csv')
crop = pd.read_csv('crop_production.csv')
[3] ✓ 0.2s
```

Reading the dataset and converting the dataset into a dataframe

```
print("City DataFrame:")
print(city.head())

print("\nCrop DataFrame:")
print(crop.head())
[4] ✓ 0.0s
```

```

City DataFrame:
   City      Date  PM2.5  PM10    NO    NO2    NOx  NH3    CO    SO2
0 Ahmedabad 2015-01-01   NaN   NaN  0.92  18.22  17.15  NaN  0.92  27.64
1 Ahmedabad 2015-01-02   NaN   NaN  0.97  15.69  16.46  NaN  0.97  24.55
2 Ahmedabad 2015-01-03   NaN   NaN 17.40  19.30  29.70  NaN 17.40  29.07
3 Ahmedabad 2015-01-04   NaN   NaN  1.70  18.48  17.97  NaN  1.70  18.59
4 Ahmedabad 2015-01-05   NaN   NaN 22.10  21.42  37.76  NaN 22.10  39.33

   O3  Benzene  Toluene  Xylene  AQI  AQI_Bucket
0 133.36    0.00    0.02    0.00  NaN          NaN
1  34.06    3.68    5.50    3.77  NaN          NaN
2  30.70    6.80   16.40    2.25  NaN          NaN
3  36.08    4.43   10.14    1.00  NaN          NaN
4  39.31    7.01   18.89    2.78  NaN          NaN

Crop DataFrame:
   State_Name  District_Name  Crop_Year  Season \
0 Andaman and Nicobar Islands  NICOBARS    2000  Kharif
1 Andaman and Nicobar Islands  NICOBARS    2000  Kharif
2 Andaman and Nicobar Islands  NICOBARS    2000  Kharif
3 Andaman and Nicobar Islands  NICOBARS    2000  Whole Year
4 Andaman and Nicobar Islands  NICOBARS    2000  Whole Year

   Crop    Area  Production
0  Arecanut 1254.0    2000.0
1  Other Kharif pulses    2.0     1.0
2      Rice   102.0    321.0
3   Banana   176.0    641.0
4  Cashewnut   720.0    165.0

```

Printing First 5 observations of city and crop dataset

```

print("City")
print("Size: ", city.shape)
print("\n")
print("Columns: ", city.columns)
print("\n")
print("Columns Types: ")
print(city.dtypes)
print("\n")
print(city.describe())
print("\n")
print(city.info())
print("\n")
print("Null values in City DataFrame: ")
print(city.isnull().sum())
print("\n")
print("Duplicates in City DataFrame: ")
print(city.duplicated().sum())
print("\n")
city_pct = (city.isnull().sum() / len(city) * 100).round(2)
null_df = pd.DataFrame({'Missing Count': city.isnull().sum(), 'Missing %': city_pct})
print(null_df[null_df['Missing Count'] > 0])

```

```

City
Size: (29531, 16)

Columns: Index(['City', 'Date', 'PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2',
               'O3', 'Benzene', 'Toluene', 'Xylene', 'AQI', 'AQI_Bucket'],
              dtype='object')

Columns Types:
City           object
Date           object
PM2.5          float64
PM10           float64
NO             float64
NO2            float64
NOx            float64
NH3            float64
CO             float64
SO2            float64
O3             float64
Benzene        float64
Toluene        float64
Xylene         float64
AQI            float64
AQI_Bucket     object
dtype: object

```

	PM2.5	PM10	NO	NO2	NOx \
count	24933.000000	18391.000000	25949.000000	25946.000000	25346.000000
mean	67.450578	118.127103	17.574730	28.560659	32.309123
std	64.661449	90.605110	22.785846	24.474746	31.646011
min	0.040000	0.010000	0.020000	0.010000	0.000000
25%	28.820000	56.255000	5.630000	11.750000	12.820000
50%	48.570000	95.680000	9.890000	21.690000	23.520000
75%	80.590000	149.745000	19.950000	37.620000	40.127500
max	949.990000	1000.000000	390.680000	362.210000	467.630000
	NH3	CO	SO2	O3	Benzene \
count	19203.000000	27472.000000	25677.000000	25509.000000	23908.000000
mean	23.483476	2.248598	14.531977	34.491430	3.280840
std	25.684275	6.962884	18.133775	21.694928	15.811136
min	0.010000	0.000000	0.010000	0.010000	0.000000
25%	8.580000	0.510000	5.670000	18.860000	0.120000
50%	15.850000	0.890000	9.160000	30.840000	1.070000
75%	30.020000	1.450000	15.220000	45.570000	3.080000
max	352.890000	175.810000	193.860000	257.730000	455.030000
	Toluene	Xylene	AQI		
count	21490.000000	11422.000000	24850.000000		
mean	8.700972	3.070128	166.463581		
std	19.969164	6.323247	140.696585		
min	0.000000	0.000000	13.000000		
25%	0.600000	0.140000	81.000000		
50%	2.970000	0.980000	118.000000		
75%	9.150000	3.350000	208.000000		
max	454.850000	170.370000	2049.000000		


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29531 entries, 0 to 29530
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   City            29531 non-null  object
1   Date            29531 non-null  object
2   PM2.5           24933 non-null  float64
3   PM10            18391 non-null  float64
4   NO              25949 non-null  float64
5   NO2             25946 non-null  float64
6   NOx             25346 non-null  float64
7   NH3             19203 non-null  float64
8   CO              27472 non-null  float64
9   SO2             25677 non-null  float64
10  O3              25509 non-null  float64
11  Benzene         23908 non-null  float64
12  Toluene         21490 non-null  float64
13  Xylene          11422 non-null  float64
14  AQI             24850 non-null  float64
15  AQI_Bucket      24850 non-null  object
dtypes: float64(13), object(3)
memory usage: 3.6+ MB
None
```

Null values in City DataFrame:

```
City          0
Date          0
PM2.5        4598
PM10         11140
NO           3582
NO2          3585
NOx          4185
NH3          10328
CO           2059
SO2          3854
O3           4022
Benzene       5623
Toluene       8041
Xylene       18109
AQI           4681
AQI_Bucket    4681
dtype: int64
```

Duplicates in City DataFrame:

0

	Missing Count	Missing %
PM2.5	4598	15.57
PM10	11140	37.72
NO	3582	12.13
NO2	3585	12.14
NOx	4185	14.17
NH3	10328	34.97
CO	2059	6.97
SO2	3854	13.05
O3	4022	13.62
Benzene	5623	19.04
Toluene	8041	27.23
Xylene	18109	61.32
AQI	4681	15.85
AQI_Bucket	4681	15.85

Inference:

The City dataset contain 29531 observations and 16 attributes

The Columns with null values are

PM2.5 = 4598

PM10 = 11140

NO = 3582

NO2 = 3585

NOx = 4185

NH3 = 10328

CO = 2059

SO2 = 3854

O3 = 4022

Benzene = 5623

Toluene = 8041

Xylene = 18109

AQI = 4681

AQI_Bucket = 4681

The City dataset does not contain any duplicate values

Xylene is missing 61.32% values which should be dropped as a modelling feature.

AQI and AQI_Bucket both are missing 15.85% combined values which is a core feature and can not be filled by itself

PM10 is missing 37.72% of null values.

Date is stored as object that must be converted into DateTime before any time series analysis.

Max Value of AQI is 2049 which is impossible because max value of AQI can be 500 leading to a data error

```
print("Crop\n")
print("Size: \n", crop.shape)
print("Columns: \n", crop.columns)
print("\n")
print("Columns Types: ")
print(crop.dtypes)
print("\n")
print(crop.describe())
print("\n")
print(crop.info())
print("\n")
print("Null values in Crop DataFrame: ")
print(crop.isnull().sum())
print("\n")
print("Duplicates in Crop DataFrame: ")
print(crop.duplicated().sum())

crop_pct = (crop.isnull().sum() / len(crop) * 100).round(2)
null_df_crop = pd.DataFrame({'Missing Count': crop.isnull().sum(), 'Missing %': crop_pct})
print(null_df_crop[null_df_crop['Missing Count'] > 0])

print(sorted(crop['State_Name'].unique()))
```

✓ 0.3s

```
Crop
Size:
(246091, 7)
Columns:
Index(['State_Name', 'District_Name', 'Crop_Year', 'Season', 'Crop', 'Area',
      'Production'],
      dtype='object')
```

Columns Types:

```
State_Name      object
District_Name   object
Crop_Year       int64
Season          object
Crop            object
Area            float64
Production      float64
dtype: object
```

	Crop_Year	Area	Production
count	246091.000000	2.460910e+05	2.423610e+05
mean	2005.643018	1.200282e+04	5.825034e+05
std	4.952164	5.052340e+04	1.706581e+07
min	1997.000000	4.000000e-02	0.000000e+00
25%	2002.000000	8.000000e+01	8.800000e+01
50%	2006.000000	5.820000e+02	7.290000e+02
75%	2010.000000	4.392000e+03	7.023000e+03

```
max      2015.000000  8.580100e+06  1.250800e+09
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 246091 entries, 0 to 246090
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   State_Name      246091 non-null object
1   District_Name   246091 non-null object
2   Crop_Year       246091 non-null int64
3   Season         246091 non-null object
4   Crop            246091 non-null object
5   Area            246091 non-null float64
6   Production      242361 non-null float64
dtypes: float64(2), int64(1), object(4)
memory usage: 13.1+ MB
None
```

Null values in Crop DataFrame:

```
State_Name      0
District_Name   0
Crop_Year       0
Season          0
Crop            0
Area            0
Production      3730
dtype: int64
```

Duplicates in Crop DataFrame:

```
0
Missing Count Missing %
Production      3730      1.52
['Andaman and Nicobar Islands', 'Andhra Pradesh', 'Arunachal Pradesh', 'Assam', 'Bihar', 'Chandigarh', 'Chhattisgarh', 'Dadra and Nagar Haveli',
```

Inference:

The Crop dataset contain 246091 observations and 7 attributes

The Columns with null values are

Production = 3730

The Crop dataset does not contain any duplicate values

Production is missing 1.5% it can be filled with median by crop type

Task 2 The data has holes — and not all holes are equal

After the inspection, it is clear that several columns have missing values. A colleague suggests simply deleting all rows with any null value. Another suggests filling everything with the mean. You are not sure either approach is right.

Devise your own missing value treatment strategy. For each column with missing data, decide whether to drop it, drop affected rows, or impute — and explain why that choice is appropriate for that specific column. Apply your strategy and verify it worked.

A markdown cell that justifies each decision column by column
Before and after null counts as evidence the treatment worked

Think: Does the volume of missing data matter? Does the column type matter? What does imputing with mean vs median assume about the data? 2 marks

Task 2

 Generate

 Code

 Markdo

```
print("Null values in City Dataset: ")
print(city.isnull().sum())
print("\n")

city_clean = city.copy()

city_clean = city_clean.drop(columns = ['Xylene', 'AQI_Bucket'])

cols = ['AQI', 'PM2.5', 'PM10', 'NO', 'NO2', 'NOx',
        'NH3', 'CO', 'SO2', 'O3',
        'Benzene', 'Toluene']

for col in cols:
    city_clean[col] = city_clean.groupby('City')[col].transform(
        lambda x: x.fillna(x.median())
    )

    city_clean[col] = city_clean[col].fillna(
        city_clean[col].median()
    )

print(city_clean.isnull().sum())

city_clean['Date'] = pd.to_datetime(city_clean['Date'], errors='coerce')
```

✓ 0.1s

Null values in City Dataset:

City	0
Date	0
PM2.5	4598
PM10	11140
NO	3582
NO2	3585
NOx	4185
NH3	10328
CO	2059
SO2	3854
O3	4022
Benzene	5623
Toluene	8041
Xylene	18109
AQI	4681
AQI_Bucket	4681

dtype: int64

City	0
Date	0
PM2.5	0
PM10	0
NO	0
NO2	0
NOx	0
NH3	0
CO	0
SO2	0

O3	0
Benzene	0
Toluene	0
AQI	0

dtype: int64

PM2.5 filled with median group by city because median is robust to extremes
PM10 filled with median group by city because PM10 is important feature and dropping it will lose too much data
NO, NO2, NOx filled with median group by city as they are correlated
NH3 filled with median group by city as dropping will lose 35% of the observations
CO filled with median group by city
SO2, Benzene, Toluene and O3 filled with median as mean > median which states that there are high outliers in the datasets and median is robust to extreme outliers
Xylene missing 60% of the values filling those null values can lead to inconsistent dataset so that column is dropped
AQI is filled with median by city as different city has different AQI
AQI_Bucket is categorical column which can not be dropped.

```
print("Null values in Crop Dataset: ")
print(crop.isnull().sum())
print("\n")

crop_clean = crop.copy()

crop_clean['Production'] = crop_clean.groupby(['Crop', 'Season'])['Production'].transform(
    lambda x: x.fillna(x.median())
)

crop_clean['Production'] = crop_clean['Production'].fillna(crop_clean['Production'].median())

print('Null values in Crop Dataset After Cleaning:')
print(crop_clean.isnull().sum())
```

✓ 0.1s


```
Null values in Crop Dataset:
```

```
State_Name      0
District_Name   0
Crop_Year       0
Season          0
Crop            0
Area            0
Production      3730
dtype: int64
```

```
Null values in Crop Dataset After Cleaning:
```

```
State_Name      0
District_Name   0
Crop_Year       0
Season          0
Crop            0
Area            0
Production      0
dtype: int64
```

Production is filled with median group by crop and season before taking the median because different crops in different seasons produce very different quantities.

Task 3 The two files disagree on how to spell "Tamil Nadu"

You need to combine both files later in the lab using the State column as the common key. But when you look closely, the same state appears with different spellings, extra spaces, or old names across the two files. There are also rows that appear more than once for no clear reason.

Make both files agree on state names and remove any redundant records. Document every inconsistency you found and how you resolved it. Show that the files are now ready to be merged reliably.

A list of all inconsistencies found and the fix applied for each

Record counts before and after to prove duplicates were removed

Think: What could go wrong in a merge if state names don't match exactly? How would you systematically find all variants of a state name? 2 marks

Task 3

A merge() in pandas uses exact string matching. 'Tamil Nadu' and 'Tamil Nadu ' (trailing space) are treated as different keys. The merge drops all rows from the right dataframe that carry the space producing incorrect NaN filled rows or missing data with no error message.

Strip whitespace from all state-name columns in both files

Title case normalisation to catch capitalisation variants

Map known old names (eg. 'Orissa' to 'Odisha', 'Uttaranchal' to 'Uttarakhand')

Build a city to state lookup for the city file

```
crop_merge = crop_clean.copy()
city_merge = city_clean.copy()

print("Crop states with leading/trailing spaces:")
space_states = [s for s in crop_merge['State_Name'].unique()
                 if s != s.strip()]
print(space_states)

crop_merge['State_Name'] = crop_merge['State_Name'].str.strip().str.title()

state_rename_map = {
    'Orissa': 'Odisha',
    'Uttaranchal': 'Uttarakhand',
    'Pondicherry': 'Puducherry'
}

crop_merge['State_Name'] = crop_merge['State_Name'].replace(state_rename_map)

c_to_s = {
    'Ahmedabad': 'Gujarat', 'Aizawl': 'Mizoram', 'Amaravati': 'Andhra Pradesh', 'Amritsar': 'Punjab', 'Bengaluru': 'Karnataka',
    'Bhopal': 'Madhya Pradesh', 'Brajrajnagar': 'Odisha', 'Chandigarh': 'Chandigarh', 'Chennai': 'Tamil Nadu', 'Coimbatore': 'Tamil Nadu',
    'Delhi': 'Delhi', 'Ernakulam': 'Kerala', 'Gurugram': 'Haryana', 'Guwahati': 'Assam', 'Hyderabad': 'Telangana', 'Jaipur': 'Rajasthan',
    'Jorapokhar': 'Jharkhand', 'Kochi': 'Kerala', 'Kolkata': 'West Bengal', 'Lucknow': 'Uttar Pradesh', 'Mumbai': 'Maharashtra', 'Patna': 'Bihar',
    'Shillong': 'Meghalaya', 'Talcher': 'Odisha', 'Thiruvananthapuram': 'Kerala', 'Visakhapatnam': 'Andhra Pradesh'
}

city_merge['State'] = city_merge['City'].map(c_to_s)

print("Cities without state mapping:")
print(city_merge[city_merge['State'].isna()][ 'City'].unique())

crop_before = len(crop_merge)
crop_merge = crop_merge.drop_duplicates()
crop_after = len(crop_merge)
```

```
print("Cities without state mapping:")
print(city_merge[city_merge['State'].isna()][ 'City'].unique())

crop_before = len(crop_merge)
crop_merge = crop_merge.drop_duplicates()
crop_after = len(crop_merge)

city_before = len(city_merge)
city_merge = city_merge.drop_duplicates()
city_after = len(city_merge)

print(f"Crop records: {crop_before} -> {crop_after}")
print(f"City records: {city_before} -> {city_after}")
```

```
Crop states with leading/trailing spaces:
['Jammu and Kashmir ', 'Telangana ']
Cities without state mapping:
[]
Crop records: 246091 -> 246091
City records: 29531 -> 29531
```

```
print("Unique states consistency check")

print("\nStates in City Dataset:")
print(sorted(city_merge['State'].dropna().unique()))

print("\nStates in Crop Dataset:")
print(sorted(crop_merge['State_Name'].dropna().unique()))

city_states = set(city_merge['State'].dropna().unique())
crop_states = set(crop_merge['State_Name'].dropna().unique())

overlap = city_states & crop_states

print(f"\nStates common to both files: {len(overlap)}")
print(sorted(overlap))

print("\nCity states NOT in crop file:")
print(sorted(city_states - crop_states))

print("\nCrop states NOT in city file:")
print(sorted(crop_states - city_states))
```

✓ 0.0s

```
Unique states consistency check

States in City Dataset:
['Andhra Pradesh', 'Assam', 'Bihar', 'Chandigarh', 'Delhi', 'Gujarat', 'Haryana', 'Jharkhand', 'Karnataka', 'Kerala', 'Madhya Pradesh', 'Maharashtra', 'Me

States in Crop Dataset:
['Andaman And Nicobar Islands', 'Andhra Pradesh', 'Arunachal Pradesh', 'Assam', 'Bihar', 'Chandigarh', 'Chhattisgarh', 'Dadra And Nagar Haveli', 'Goa', 'G

States common to both files: 20
['Andhra Pradesh', 'Assam', 'Bihar', 'Chandigarh', 'Gujarat', 'Haryana', 'Jharkhand', 'Karnataka', 'Kerala', 'Madhya Pradesh', 'Maharashtra', 'Meghalaya',

City states NOT in crop file:
['Delhi']

Crop states NOT in city file:
['Andaman And Nicobar Islands', 'Arunachal Pradesh', 'Chhattisgarh', 'Dadra And Nagar Haveli', 'Goa', 'Himachal Pradesh', 'Jammu And Kashmir', 'Manipur',
```

Task 4 Where do most cities actually sit on the AQI scale?

The pollution control board wants to know — are most Indian cities moderately polluted, or is the problem concentrated in just a few places? They also want to know whether the average AQI is a fair number to report publicly, or whether extreme cities are pulling it up unfairly.

Investigate the shape of the AQI distribution. Decide which visualisation(s) best answer both questions the board is asking, justify your choice, produce the plot(s), and record two specific observations that directly address their concerns.

Your chosen plot(s) with fully labelled axes and a descriptive title

A markdown cell with your justification for the chosen plot type and two observations

Think: Which plot shows where values cluster? Which one reveals extreme values? Do you need one plot or two to answer both questions? 2 marks

Task 4


```
plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
plt.hist(city_clean['AQI'], bins=30)
plt.title("AQI Distribution")
plt.xlabel("AQI")
plt.ylabel("Frequency")

plt.subplot(1,2,2)
plt.boxplot(city_clean['AQI'])
plt.title("AQI Box Plot")
plt.ylabel("AQI")

plt.tight_layout()
plt.show()
```

✓ 0.1s



The histogram shows that most AQI values are concentrated between about 50 and 250, meaning most cities experience moderate to poor air quality. The box plot reveals a large number of extreme outliers, with some AQI values exceeding 1000 and even 2000. These unusually high values can increase the average AQI suggesting that the mean may overstate the pollution level of a typical city. A histogram was chosen because it shows where AQI values are clustered and helps determine whether most cities fall into a particular pollution range. A box plot was chosen because it highlights extreme values and outliers.

Task 5 Something is making the average AQI look worse than it is

A data quality report flags that a few AQI readings in the dataset are implausibly high — values that no monitoring station should realistically record. If left in, these values will distort every statistic and model built on this data. The team lead asks you to "handle the extreme values properly" but does not tell you how.

Identify whether extreme values exist, quantify them, decide on an appropriate treatment method, apply it, and demonstrate that the data is now cleaner. Justify every decision you make.

The method you chose to detect extremes and why

The count of values affected and the treatment applied

A visual comparison of the data before and after treatment

Think: Should extreme values always be deleted? What are the alternatives? How do you prove your treatment actually worked? 2 marks

Task 5

```

q1 = city_clean['AQI'].quantile(0.25)
q3 = city_clean['AQI'].quantile(0.75)

iqr = q3 - q1

lower = q1 - 1.5 * iqr
upper = q3 + 1.5 * iqr

outliers = city_clean[
    (city_clean['AQI'] < lower) |
    (city_clean['AQI'] > upper)
]

print("Number of Outliers:", len(outliers))

city_outlier = city_clean.copy()

city_outlier['AQI'] = city_outlier['AQI'].clip(
    lower=lower,
    upper=upper
)

print("Maximum AQI before treatment:", city_clean['AQI'].max())
print("Maximum AQI after treatment :", city_outlier['AQI'].max())

```

```

Number of Outliers: 1465
Maximum AQI before treatment: 2049.0
Maximum AQI after treatment : 391.0

```

```

plt.figure(figsize=(10,5))


plt.subplot(1,2,1)
plt.boxplot(city_clean['AQI'])
plt.title("Before Outlier Treatment")
plt.ylabel("AQI")

plt.subplot(1,2,2)
plt.boxplot(city_outlier['AQI'])
plt.title("After Outlier Treatment")
plt.ylabel("AQI")

plt.tight_layout()
plt.show()

```

✓ 0.1s



The boxplot before treatment shows a large number of extreme AQI values with some observations exceeding 2000. After using IQR method AQI values and capped at IQR upper limit the extreme values were greatly reduced and the maximum AQI dropped to 391.

The method flags any value below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$ as a statistical outlier.

AQI is heavily right-skewed

IQR is robust to the very outliers it is trying to detect

Conclusion /inference

The raw datasets required extensive cleaning before analysis could begin. Missing values in the city dataset were handled using city wise median to preserve local trends while the heavily missing Xylene column was removed. In the crop dataset production gaps were filled using crop and season wise median values. State name inconsistencies were standardized and city-to-state mapping was corrected. AQI outliers were treated using IQR based clipping making the datasets clean and ready for analysis.

Lab 2

Lab Exercise II:

Aim

- To investigate whether worsening air quality across Indian states is linked to declining agricultural output by independently choosing, applying, and justifying appropriate data analysis and visualisation techniques on raw, messy, real-world datasets.
-

Evaluation Rubrics

Submission Guidelines

7. Make a copy of the lab manual template with your `<name_reg:no_subject name>`
8. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
9. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
10. Create a **Git repository** in your profile `<ML lab-reg no>`. Follow a different branch for each lab `<Lab 1, Lab 2 ...>` and push the code to Git.
11. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
12. Upload the PDF to Google Classroom before the deadline.

Code with Results

Task 6: Is India's air getting better or worse over time?

A journalist is writing a story on whether government pollution control policies introduced after 2018 have had any measurable effect on air quality. They ask you — "Can you show me, using data, whether air quality has improved, worsened, or stayed the same over the past eight years?"

Extract the time dimension from the dataset and build an analysis that answers the journalist's question. Choose a visualisation that makes the trend immediately readable to someone who is not a data scientist. Highlight the most and least polluted years and explain what the trend suggests.

Your plot with the trend clearly visible and key years highlighted

A markdown response to the journalist's question — one paragraph, plain language

Think: What needs to happen to the Date column before you can group by year? Which plot type communicates a trend over time most clearly? 2 marks

Task 6

```
city_trend = city_outlier.copy()

city_trend['Year'] = city_trend['Date'].dt.year

yearly_aqi = city_trend.groupby('Year')['AQI'].mean().reset_index()
yearly_aqi.columns = ['Year', 'Mean_AQI']

most_polluted = yearly_aqi.loc[yearly_aqi['Mean_AQI'].idxmax()]
least_polluted = yearly_aqi.loc[yearly_aqi['Mean_AQI'].idxmin()]

print(yearly_aqi)

print(f"\nMost polluted year : {int(most_polluted['Year'])} "
      f"(AQI {most_polluted['Mean_AQI']:.2f})")

print(f"Least polluted year: {int(least_polluted['Year'])} "
      f"(AQI {least_polluted['Mean_AQI']:.2f})")
```

✓ 0.0s

	Year	Mean_AQI
0	2015	185.398072
1	2016	191.520558
2	2017	175.784816
3	2018	160.005718
4	2019	145.420830
5	2020	111.574042

Most polluted year : 2016 (AQI 191.52)
 Least polluted year: 2020 (AQI 111.57)

```
plt.figure(figsize=(10,5))

plt.plot(yearly_aqi['Year'],
         yearly_aqi['Mean_AQI'],
         marker='o')


plt.scatter(most_polluted['Year'],
           most_polluted['Mean_AQI'],
           s=100,
           label='Most Polluted Year')

plt.scatter(least_polluted['Year'],
           least_polluted['Mean_AQI'],
           s=100,
           label='Least Polluted Year')

plt.title('Average AQI by Year')
plt.xlabel('Year')
plt.ylabel('Mean AQI')
plt.grid(True)
plt.legend()

plt.show()
```

✓ 0.1s



India Air Quality has been improved in last six years although there is a sudden increase in average AQI from 2015 to 2016 but after 2016 the average AQI keeps on constantly decreasing each year recording approx 111.57 in 2020. The most polluted year is 2016 and least is 2020 there is a 42% decrease in average AQI between 2016 and 2020. The decline observed in 2020 may have been influenced by reduced industrial and transportation activity during the COVID-19 lockdowns but the decline had already been started several years earlier suggesting that broader pollution control measures and environmental initiatives may also have contributed to the improvement.

Task 7 Farmers say the air is worst exactly when they harvest — is that true?

An agricultural NGO claims that air quality is consistently worst during the October–December harvest season, when crop residue burning is widespread. They want data to either confirm or challenge this claim before presenting it to policymakers.

Investigate whether AQI follows a seasonal pattern across the year. Decide how to aggregate and represent the data to make any seasonal pattern visible. Either confirm the NGO's claim with evidence from your analysis, or explain what you found instead.

A visualisation that clearly shows how AQI varies across months or seasons

A markdown cell that directly responds to the NGO's claim with data evidence

Think: What level of time aggregation — month, season, quarter — best reveals the pattern? How do you extract that from a date column? 2 marks

Task 7

```

season_data = city_outlier.copy()

season_data['Month'] = season_data['Date'].dt.month

monthly_aqi = season_data.groupby('Month')['AQI'].mean().reset_index()
monthly_aqi.columns = ['Month', 'Mean_AQI']

month_names = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun',
               'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']

monthly_aqi['Month_Name'] = monthly_aqi['Month'].apply(
    lambda x: month_names[x - 1]
)

def season(m):
    if m in [12, 1, 2]:
        return 'Winter (Dec-Feb)'
    elif m in [3, 4, 5]:
        return 'Pre-Monsoon (Mar-May)'
    elif m in [6, 7, 8, 9]:
        return 'Monsoon (Jun-Sep)'
    else:
        return 'Post-Monsoon (Oct-Nov)'

monthly_aqi['Season'] = monthly_aqi['Month'].apply(season)

```

```

seasonal_mean = season_data.copy()
seasonal_mean['Season'] = seasonal_mean['Month'].apply(season)

seasonal_mean = (
    seasonal_mean.groupby('Season')['AQI']
    .mean()
    .reset_index()
    .sort_values('AQI', ascending=False)
)

print("Monthly Mean AQI:")
print(monthly_aqi[['Month_Name', 'Mean_AQI', 'Season']].to_string(index=False))

print("\nSeasonal Mean AQI:")
print(seasonal_mean.to_string(index=False))

```

Monthly Mean AQI:

Month_Name	Mean_AQI	Season
Jan	207.696871	Winter (Dec-Feb)
Feb	184.675824	Winter (Dec-Feb)
Mar	154.527283	Pre-Monsoon (Mar-May)
Apr	139.972243	Pre-Monsoon (Mar-May)
May	138.559731	Pre-Monsoon (Mar-May)
Jun	124.918534	Monsoon (Jun-Sep)
Jul	116.608379	Monsoon (Jun-Sep)
Aug	116.645517	Monsoon (Jun-Sep)
Sep	119.412333	Monsoon (Jun-Sep)
Oct	166.974149	Post-Monsoon (Oct-Nov)
Nov	210.678101	Post-Monsoon (Oct-Nov)
Dec	208.411304	Winter (Dec-Feb)

Seasonal Mean AQI:

Season	AQI
Winter (Dec-Feb)	200.321169
Post-Monsoon (Oct-Nov)	188.522625
Pre-Monsoon (Mar-May)	144.327551
Monsoon (Jun-Sep)	119.765366

```

season_colors = {
    'Winter (Dec-Feb)': 'skyblue',
    'Pre-Monsoon (Mar-May)': 'orange',
    'Monsoon (Jun-Sep)': 'green',
    'Post-Monsoon (Oct-Nov)': 'red'
}

plot_df = monthly_aqi.sort_values('Month')

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

axes[0].bar(
    plot_df['Month_Name'],
    plot_df['Mean_AQI'],
    color=plot_df['Season'].map(season_colors)
)

axes[0].set_title('Average AQI by Month')
axes[0].set_xlabel('Month')
axes[0].set_ylabel('Mean AQI')
axes[0].tick_params(axis='x', rotation=45)

season_order = [
    'Winter (Dec-Feb)',
    'Pre-Monsoon (Mar-May)',
    'Monsoon (Jun-Sep)',
    'Post-Monsoon (Oct-Nov)'
]

season_df = city_outlier.copy()
season_df['Month'] = season_df['Date'].dt.month

season_data = [
    season_df[season_df['Month'].apply(season) == s]['AQI']
    for s in season_order
]

```

```

axes[1].boxplot(season_data, labels=season_order)


axes[1].boxplot(season_data, labels=season_order)

axes[1].set_title('AQI Distribution by Season')
axes[1].set_ylabel('AQI')

plt.tight_layout()
plt.show()

```

✓ 0.3s



The claim is only partly true. Air pollution is high during the October–November harvest season, with November recording the highest average AQI. Pollution remains almost as severe in December and January showing that the problem extends beyond the harvest period. Monsoon months (June–September) have the cleanest air because rainfall helps remove pollutants from the atmosphere. This suggests that while crop-residue burning contributes to poor air quality winter weather conditions that trap pollutants also play a major role.

Task 8: Can the two datasets talk to each other?

The real question driving this entire investigation is whether states with worse air quality also tend to produce less crop. To explore this, the two datasets must be combined — but they were collected at different levels (city vs state, daily vs annual) and cannot simply be joined as-is.

Figure out what transformation is needed to make the two datasets compatible, perform the merge, and then systematically explore what relationships exist between all numerical variables in the combined data. Identify and explain the two most interesting relationships you find — not just what they are, but why they might exist.

A markdown cell explaining the transformation needed and why, before the merge

A visualisation of relationships across all numerical features

Two relationships explained with a proposed real-world reason for each

Think: If one file has city-day rows and the other has state-year rows, what needs to happen before they can be joined? What does a correlation matrix actually tell you — and what doesn't it tell you? 3 marks

Task 8

```
city_outlier['State'] = city_outlier['City'].map(c_to_s)

city_state = city_outlier.dropna(subset=['State']).copy()
city_state = city_state.groupby('State').agg(
    Mean_AQI = ('AQI', 'mean'),
    Mean_PM25 = ('PM2.5', 'mean'),
    Mean_NO2 = ('NO2', 'mean'),
    Mean_SO2 = ('SO2', 'mean'),
    Mean_PM10 = ('PM10', 'mean')
).reset_index()

crop_state = crop_merge.copy()
crop_state = crop_state.groupby('State_Name').agg(
    Total_Production = ('Production', 'sum'),
    Total_Area = ('Area', 'sum')
).reset_index()
crop_state['Yield'] = crop_state['Total_Production'] / crop_state['Total_Area']
crop_state.rename(columns={'State_Name': 'State'}, inplace=True)

merged = pd.merge(city_state, crop_state, on='State', how='inner')
print(f"Merged dataset shape: {merged.shape}")
print(merged[['State', 'Mean_AQI', 'Mean_PM25', 'Yield', 'Total_Production']].to_string(index=False))
```

Merged dataset shape: (20, 9)

	State	Mean_AQI	Mean_PM25	Yield	Total_Production
Andhra Pradesh	106.463324	42.500396	131.702442	1.732491e+10	
Assam	137.151394	63.651753	30.005695	2.111764e+09	
Bihar	229.776103	117.904634	2.857217	3.665010e+08	
Chandigarh	96.259868	40.742401	5.195376	6.505650e+04	
Gujarat	336.380040	64.889706	3.383944	5.243219e+08	
Haryana	217.997618	114.631751	4.292970	3.846385e+08	
Jharkhand	148.783576	61.837348	1.147661	1.077774e+07	
Karnataka	93.908412	35.463743	4.255281	8.634396e+08	
Kerala	80.608635	27.890564	3067.605766	9.788139e+10	
Madhya Pradesh	132.339100	49.897024	1.363072	4.495591e+08	
Maharashtra	96.536585	28.214617	3.922346	1.263804e+09	
Meghalaya	51.832258	28.057742	3.001837	1.211250e+07	
Mizoram	34.557522	17.012301	1.702848	1.692180e+06	
Odisha	151.891841	60.744557	1.456157	1.609543e+08	
Punjab	117.200655	54.275520	4.627288	5.863957e+08	
Rajasthan	133.371634	54.449533	1.037128	2.821247e+08	
Tamil Nadu	106.986221	46.756344	125.950833	1.207842e+10	
Telangana	108.333500	46.746690	4.120369	3.352358e+08	
Uttar Pradesh	213.469886	108.514629	7.459280	3.234580e+09	
West Bengal	136.584767	62.956560	6.489690	1.397913e+09	


```

num_cols = ['Mean_AQI', 'Mean_PM25', 'Mean_NO2', 'Mean_SO2', 'Mean_PM10',
            'Total_Production', 'Total_Area', 'Yield']

corr = merged[num_cols].corr()

plt.figure(figsize=(10, 7))
sns.heatmap(corr, annot=True, fmt='.2f', cmap='coolwarm',
            center=0, square=True, linewidths=0.5)
plt.title('Correlation Matrix: Air Quality vs Crop Features (State Level)')
plt.tight_layout()
plt.show()

plt.figure(figsize=(14, 6))

sorted_merged = merged.sort_values('Yield', ascending=False)

plt.bar(sorted_merged['State'], sorted_merged['Yield'], color='steelblue', edgecolor='white')


plt.title("Crop Yield by State", fontsize=14, fontweight='bold')
plt.xlabel("State", fontsize=11)
plt.ylabel("Yield (Production / Area)", fontsize=11)
plt.xticks(rotation=45, ha='right', fontsize=9)

plt.tight_layout()
plt.show()

```

✓ 0.4s





Mean_AQI and Mean_NO2 ($r = 0.73$, strong positive)
 States with higher AQI also tend to have higher NO2 levels because NO2 comes mainly from vehicles and power plants states like Delhi, UP, and Bihar have dense traffic and heavy industry so both AQI and NO2 rise together.

Pollutants vs Yield ($r = -0.20$ to -0.30 weak negative)
 States with worse air quality tend to have slightly lower crop yield. High PM2.5 and PM10 can reduce sunlight reaching crops, and SO2/NO2 contribute to acid rain which degrades soil over time. The correlation is weak though factors like irrigation and rainfall likely matter more. Kerala is an outlier with the lowest AQI 80.6 but highest yield 3067.

A correlation matrix only shows linear relationships between pairs of variables it can not tell if pollution causes lower yield. Correlation does not imply causation.

Task 9 : The minister needs to act — what do you tell her?

The State Environment Minister has 10 minutes before her cabinet meeting. She has never opened a Jupyter notebook. Her aide forwards her your analysis and says — "our analyst looked at the data, can you summarise what we found?" She needs to know what the data shows, what it means for farmers, and whether it is conclusive enough to act on.

Write a clear, honest briefing addressed directly to the minister. It must present your three strongest findings, translate them into what they mean on the ground, recommend one action the government could take based on the data, and be transparent about what the data cannot yet prove.

150–200 words in a markdown cell, addressed to the minister. Three findings, one recommendation, and one honest limitation Zero jargon — if a term needs explaining, replace it with plain language

Think: Which of your findings actually matters to a farmer or a policymaker? What is the difference between correlation and proof? Would the minister act on what you found? 3 marks

Task 9

1. India air is getting cleaner: The average pollution level across monitored cities fell by 45% between 2016 and 2020 from a reading of 191 down to 111 on the national scale. This is a meaningful improvement and suggests that measures taken after 2018 are working though the 2020 also played a role.
2. October and November are the worst months every year consistently: Pollution spikes sharply in November average more than double those in July. This pattern repeats every year and aligns with the period of widespread crop residue burning across northern states.
3. states with higher pollution tend to show slightly lower farm output per unit of land: The relationship is weak and not conclusive but it points in the direction one would expect if air pollution is affecting crop health.

Recommended action: Expand real-time air-quality monitoring in agricultural districts during October–November and link it to existing alerts so government can deploy enforcement resources where the data shows the worst spikes.

Limitation: The two datasets cover different time periods and different geographic scales. We cannot prove that pollution causes lower yields in these states other factors such as rainfall, soil type, and crop mix also matter.

Conclusion /inference

Conclusion of Lab1 and Lab2

Air Quality Is Improving, But Not Uniformly: India national average AQI dropped from 191 in 2016 to 111 in 2020 a 45% improvement at a rate of roughly 17 AQI points per year.

Air quality follows a clear and repeating seasonal pattern as shown in the AQI distribution plot by season

The data confirms that worsening air quality and lower agricultural yield tend to co-occur at the state level but does not prove that one causes the other. The association is statistically weak ($r = -0.21$) and confounded by crop type, geography, and agricultural practices.

Visualization showed that AQI is not uniform across cities or time. Pollution spikes sharply during October and November, coinciding with post-harvest crop residue burning in northern India. Winter months (December to February) consistently record the worst air quality levels. After merging both datasets using state attribute a weak negative correlation ($r = -0.20$ to -0.30) was observed between pollution and crop yield. The relationship is not strong enough to conclude causation as other factors like rainfall, irrigation, and crop type also significantly influence yield.

Build the case — or break it

Attempt only after completing Tasks 1–9. The central hypothesis of this entire investigation is: states with higher pollution have lower crop yields. Your job here is to either build the strongest possible case for this hypothesis — or find evidence that challenges it.

Task A: The two extremes — do they tell the same story?

If the hypothesis is true, you would expect the most polluted states and the least polluted states to show a clear difference in crop output. But the data might surprise you.

Identify the states that represent the two extremes of pollution. Compare their agricultural output using a visualisation of your choice. Analyse what you see — does the pattern support the hypothesis, contradict it, or is it more complicated than that?

Optional Task**Task A**

```
merged_sorted = merged.sort_values('Mean_AQI', ascending=False).reset_index(drop=True)


top5 = merged_sorted.head(5)
bot5 = merged_sorted.tail(5)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 6))

ax1.bar(top5['State'], top5['Yield'], color='#d62728', edgecolor='white', width=0.6)
ax1.set_title("Most Polluted States", fontsize=12, fontweight='bold', color='#d62728')
ax1.set_xlabel("State")
ax1.set_ylabel("Yield (Production / Area)")
ax1.tick_params(axis='x', rotation=45)

ax2.bar(bot5['State'], bot5['Yield'], color='#2ca02c', edgecolor='white', width=0.6)
ax2.set_yscale('log')
ax2.set_title("Least Polluted States", fontsize=12, fontweight='bold', color='#2ca02c')
ax2.set_xlabel("State")
ax2.set_ylabel("Yield - log scale (Production / Area)")
ax2.tick_params(axis='x', rotation=45)

fig.suptitle("Crop Yield: Most Polluted vs Least Polluted States",
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()
```



The hypothesis state with high pollution cause lower crop yield is not completely rejected or completely accepted because states like (Gujarat, Bihar, Haryana, UP, odisha) are top 5 most polluted states but the crop yield is low to moderate. While states like kerala with low AQI shows an extreme good crop yield but other states like meghalya, mizoram, karnataka, chandigarh does not follow the same approach.

Task B: Put a number on the relationship

A research team wants to know not just whether a relationship exists between AQI and crop production, but how strong it is and in which direction. They will use your output to decide whether to fund a full causal study.

Quantify the relationship between average state AQI and average crop production. Choose a method that both measures the strength and makes the pattern visible in a single output. Interpret your result for the research team — be honest about what it proves and what it does not.

Think: What does a correlation value of 0.2 vs 0.8 mean in practice? Does a strong correlation mean pollution is causing yield loss?

Task B

```
r, _ = stats.pearsonr(merged['Mean_AQI'], merged['Yield'])
print(f"Correlation (r): {round(r, 3)}")


slope, intercept, _, _, _ = stats.linregress(merged['Mean_AQI'], merged['Yield'])
x_line = np.linspace(merged['Mean_AQI'].min(), merged['Mean_AQI'].max(), 100)
y_line = slope * x_line + intercept

plt.figure(figsize=(10, 6))
plt.scatter(merged['Mean_AQI'], merged['Yield'], color='steelblue', s=80)

for _, row in merged.iterrows():
    plt.text(row['Mean_AQI'] + 2, row['Yield'], row['State'], fontsize=7)

plt.plot(x_line, y_line, color='red', linestyle='--', label=f'r = {r:.2f}')

plt.title("Mean AQI vs Crop Yield by State")
plt.xlabel("Mean AQI")
plt.ylabel("Yield (Production / Area)")
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



Task C: One plot to rule them all

You have produced many visualisations across this lab. A data journalism outlet wants to publish just one chart from your work to accompany a story on India's air-agriculture crisis.

Choose the single visualisation from your entire lab that you believe most powerfully tells the story of this dataset. Write a 3–4 sentence caption for it — not a description of what the chart shows, but an explanation of why it matters and what it reveals that no other chart in your analysis does.

Caption reveals insight, not just description

Across all optional tasks, your writing must acknowledge that correlation is not causation and name at least one factor — rainfall, irrigation, soil type, economic conditions — that this dataset cannot account for but which could explain any relationship observed.

Task C

The heatmap shows all air quality and crop variables together in one view. Pollutants like AQI, NO₂, and PM_{2.5} are strongly linked to each other ($r = 0.73$ to 0.76), but their connection to crop yield is weak ($r = -0.20$ to -0.30). The pattern points in the expected direction but is not strong enough to prove causation rainfall, irrigation and soil type are not in this dataset and could equally explain the yield differences seen across states.

Lab 3

Lab Exercise III:

Aim

- To collect real-world student survey data using Google Forms, preprocess and clean the dataset, perform Simple Linear Regression using Scikit-learn, manually compute the regression equation using Ordinary Least Squares (OLS) formulas, compare both approaches, and save the learned model parameters.
-

Evaluation Rubrics

Submission Guidelines

13. Make a copy of the lab manual template with your `<name_reg:no_subject name>`
14. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
15. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
16. Create a **Git repository** in your profile `<ML lab-reg no>`. Follow a different branch for each lab `<Lab 1, Lab 2 ...>` and push the code to Git.
17. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
18. Upload the PDF to Google Classroom before the deadline.

Code with Results

Part A: Data Collection and Preprocessing

Perform the following steps:

1. Load the dataset using Pandas


```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import re
import pickle
import warnings
warnings.filterwarnings('ignore')

```

```
df = pd.read_csv('Student_Awareness_Survey.csv')
```

2. Display the first 5 rows

```

print("First 5 rows of the dataset:")
df.head()

```

First 5 rows of the dataset:

	Timestamp	Registration Number	Email	Job role that you are interested in	What is the minimum salary of students placed through campus (In LPA..respond as a number)	What is the maximum salary of students placed through campus (In LPA..respond as a number)	What is the median salary of students placed through campus (In LPA..respond as a number)	Which is the highest paying company that recruits from campus?	Rate your contribution towards extra curricular activities	Rate your technical competencies
0	6/15/2026 9:25:39	2547231	kunnal.kunnal@mca.christuniversity.in	Software Development Engineer (SDE)	3.5	12	6.8	Akasa air	4.0	3.0
1	6/15/2026 9:53:54	2547237	omkaar.chakraborty@mca.christuniversity.in	Software Development Engineer (SDE)	6	20	10	Fractal	4.0	4.0
2	6/15/2026 9:54:56	2547203	abhinav.jain@mca.christuniversity.in	Full Stack Developer	4	12	6.3	Akasa Air	5.0	4.0
3	6/15/2026 9:55:17	2547228	jai.pareek@mca.christuniversity.in	Full Stack Developer	600000	1400000	800000	Akasa airlines	5.0	4.0
4	6/15/2026 9:55:42	2547241	r.karan@mca.christuniversity.in	Software Development Engineer (SDE)	4	4	6	12	2.0	3.0

```
df.describe()
```

[188]

	Registration Number	Rate your contribution towards extra curricular activities	Rate your technical competencies	Your GPA of last semester
count	5.000000e+01	49.000000	49.000000	50.000000
mean	2.547232e+06	3.489796	3.530612	3.497000
std	1.792671e+01	1.209725	0.738863	0.690865
min	2.547201e+06	1.000000	2.000000	2.740000
25%	2.547217e+06	3.000000	3.000000	3.305000
50%	2.547232e+06	4.000000	3.000000	3.400000
75%	2.547246e+06	4.000000	4.000000	3.600000
max	2.547262e+06	5.000000	5.000000	8.000000

```
df.info()
```

[189]

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 15 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Timestamp                                                            50 non-null    object
1   Registration Number                                                  50 non-null    int64
2   Email                                                                50 non-null    object
3   Job role that you are interested in                                 50 non-null    object
4   What is the minimum salary of students placed through campus (In LPA..respond as a number) 50 non-null    object
5   What is the maximum salary of students placed through campus (In LPA..respond as a number) 50 non-null    object
6   What is the median salary of students placed through campus (In LPA..respond as a number) 50 non-null    object
7   Which is the highest paying company that recruits from campus?     49 non-null    object
8   Rate your contribution towards extra curricular activities          49 non-null    float64
9   Rate your technical competencies                                    49 non-null    float64
10  What are your package expectations (LPA)                             49 non-null    object
11  Your CIA % of last semester                                          50 non-null    object
12  Your GPA of last semester                                            50 non-null    float64
13  Your maximum attendance % till last semester                       50 non-null    object
14  Internships Interests                                                49 non-null    object
dtypes: float64(3), int64(1), object(11)
memory usage: 6.0+ KB
```

3. Check dataset dimensions

```
print("Dataset Dimensions")
df.shape
```

[190]

```
Dataset Dimensions

(50, 15)
```

The Dataset contain 50 observations and 15 attributes

4. Identify missing values

```
print("Null Values in The dataset")
df.isnull().sum()
```

1]

```
Null Values in The dataset
```

Timestamp	0
Registration Number	0
Email	0
Job role that you are interested in	0
What is the minimum salary of students placed through campus (In LPA..respond as a number)	0
What is the maximum salary of students placed through campus (In LPA..respond as a number)	0
What is the median salary of students placed through campus (In LPA..respond as a number)	0
Which is the highest paying company that recruits from campus?	1
Rate your contribution towards extra curricular activities	1
Rate your technical competencies	1
What are your package expectations (LPA)	1
Your CIA % of last semester	0
Your GPA of last semester	0
Your maximum attendance % till last semester	0
Internships Interests	1

dtype: int64

```
df.isna().sum()
```

[192]

```
... Timestamp 0
Registration Number 0
Email 0
Job role that you are interested in 0
What is the minimum salary of students placed through campus (In LPA..respond as a number) 0
What is the maximum salary of students placed through campus (In LPA..respond as a number) 0
What is the median salary of students placed through campus (In LPA..respond as a number) 0
Which is the highest paying company that recruits from campus? 1
Rate your contribution towards extra curricular activities 1
Rate your technical competencies 1
What are your package expectations (LPA) 1
Your CIA % of last semester 0
Your GPA of last semester 0
Your maximum attendance % till last semester 0
Internships Interests 1
dtype: int64
```

```
total_null = df.isnull().sum().sum()
print(f"Total Null Values in the dataset: {total_null}")
```

[193]

```
... Total Null Values in the dataset: 5
```

There are total 5 null values in the dataset

5. Remove or handle null values

Handling Missing Values

```
key_cols = [
    "Your CIA % of last semester",
    "Your GPA of last semester",
    "Your maximum attendance % till last semester",
]

df.dropna(subset=key_cols, inplace=True)
print(f"Rows after dropping key-column nulls: {len(df)}")
print("\nRemaining nulls in key columns:")
print(df[key_cols].isnull().sum().to_string())
```

[194]

```
... Rows after dropping key-column nulls: 50

Remaining nulls in key columns:
Your CIA % of last semester      0
Your GPA of last semester        0
Your maximum attendance % till last semester  0
```

6. Convert required columns into numerical datatype

Convert required columns into numerical datatype

```
print(f"Your CIA % of last semester: {df['Your CIA % of last semester'].dtype}")
print(f"Your GPA of last semester: {df['Your GPA of last semester'].dtype}")
print(f"Your maximum attendance % till last semester: {df['Your maximum attendance % till last semester'].dtype}")
```

[195]

```
... Your CIA % of last semester: object
Your GPA of last semester: float64
Your maximum attendance % till last semester: object
```

[196]

```
def clean_numeric(val):
    s = str(val).strip().replace("%", "").replace(" ", "")
    match = re.search(r"[\d.]+", s)
    return float(match.group()) if match else np.nan
```

```
df["CIA_pct"] = df["Your CIA % of last semester"].apply(clean_numeric)
df["Attendance_pct"] = df["Your maximum attendance % till last semester"].apply(clean_numeric)
df["GPA"] = pd.to_numeric(df["Your GPA of last semester"], errors="coerce")

print("\nDatatypes AFTER conversion:")
print(f"CIA_pct: {df['CIA_pct'].dtype}")
print(f"Attendance_pct: {df['Attendance_pct'].dtype}")
print(f"GPA: {df['GPA'].dtype}")

df.dropna(subset=["CIA_pct", "Attendance_pct", "GPA"], inplace=True)
print(f"\nRows after numeric conversion + NaN drop: {len(df)}")
```

Datatypes AFTER conversion:

CIA_pct: float64

Attendance_pct: float64

GPA: float64

Rows after numeric conversion + NaN drop: 50

[198]

```
outlier = df[df["GPA"] > 5][["CIA_pct", "GPA", "Attendance_pct"]]
print(outlier.to_string())
df = df[df["GPA"] < 5].reset_index(drop=True)
print(f"Rows after outlier removal: {len(df)}")
```

```
...
   CIA_pct  GPA  Attendance_pct
13      7.0  8.0             85.0
Rows after outlier removal: 49
```

7. Remove duplicate records if present

Remove duplicate records if present

```
> dup_count = df.duplicated().sum()
print(f"\nNumber of duplicate rows found: {dup_count}")

if dup_count > 0:
    df.drop_duplicates(inplace=True)
    df.reset_index(drop=True, inplace=True)
    print(f"Duplicates removed. Rows remaining: {len(df)}")
else:
    print("No duplicates found. Dataset unchanged.")
```

```
Number of duplicate rows found: 0
No duplicates found. Dataset unchanged.
```

```
No duplicates in the dataset
```

8. Generate statistical summary

Generate statistical summary

```
> ~
summary = df[["CIA_pct", "Attendance_pct", "GPA"]].describe()
print()
print(summary.to_string())
```

[200]

..

	CIA_pct	Attendance_pct	GPA
count	49.000000	49.000000	49.000000
mean	72.825102	94.004898	3.405102
std	6.378355	3.674066	0.237023
min	60.000000	85.000000	2.740000
25%	69.780000	92.000000	3.300000
50%	71.780000	95.000000	3.400000
75%	77.000000	96.000000	3.600000
max	91.000000	100.000000	3.900000

9. Select appropriate dependent and independent variables

Select appropriate dependent and independent variables

```
▷ ~
print("""
    Dependent Variable - GPA: GPA is the outcome we want to predict.

    Independent Variables (X):
    | Experiment 1 - CIA_pct (CIA % of last semester)
    | Experiment 2 - Attendance_pct (Maximum attendance %)
    """)
```

[201]

...

Dependent Variable - GPA: GPA is the outcome we want to predict.

Independent Variables (X):

Experiment 1 - CIA_pct (CIA % of last semester)

Experiment 2 - Attendance_pct (Maximum attendance %)


```
fig, axes = plt.subplots(1, 3, figsize=(15, 4))

cols = {"CIA_pct": "CIA %", "Attendance_pct": "Attendance %", "GPA": "GPA"}
colors = ["steelblue", "darkorange", "seagreen"]

for ax, (col, label), color in zip(axes, cols.items(), colors):
    ax.hist(df[col], bins=10, color=color, edgecolor="white", alpha=0.85)
    ax.set_title(f"Distribution of {label}")
    ax.set_xlabel(label)
    ax.set_ylabel("Frequency")
    ax.grid(True, linestyle="--", alpha=0.4)

plt.suptitle("Feature Distributions", fontsize=13, fontweight="bold")
plt.tight_layout()
plt.show()
```

94] ✓ 0.3s



The Cia% values lies from 60 to 90 with with maximum number between 70 to 73
Attendance lies from 85 to 100 with maximum between 94 to 95
GPA values lies from 2.8 to 4 with maximum concetration from 3.3 to 3.4


```
fig, axes = plt.subplots(1, 2, figsize=(12, 4))

for ax, (x_col, xlabel, color) in zip(axes, [
    ("CIA_pct", "CIA %", "steelblue"),
    ("Attendance_pct", "Attendance %", "darkorange")
]):
    ax.scatter(df[x_col], df["GPA"], color=color, edgecolors="white", alpha=0.8)
    ax.set_title(f"{xlabel} vs GPA")
    ax.set_xlabel(xlabel)
    ax.set_ylabel("GPA")
    ax.grid(True, linestyle="--", alpha=0.4)

plt.suptitle("Scatter Plots – Independent vs Dependent Variable", fontsize=13, fontweight="bold")
plt.tight_layout()
plt.show()
```

95] ✓ 0.2s

Scatter Plots — Independent vs Dependent Variable



The data is spread all over in CIA vs GPA and Attendance vs GPA

In first graph p spread across the full X range (60–92), giving the regression line enough variation to work with
The second graph no clear directional trend points are scattered almost horizontally across the GPA range

Saveing to cleaned dataset

```
df.to_csv("Student_Survey_Cleaned.csv", index=False)
```

Part B: Simple Linear Regression using Scikit-learn

Perform the following:

1. Split the dataset into training and testing sets
2. Train the model using LinearRegression from Scikit learn
3. Obtain:
 - Slope
 - Intercept
4. Predict output values using the trained mod

```

df_clean = pd.read_csv("Student_Survey_Cleaned.csv")

X1 = df_clean[["CIA_pct"]]
y1 = df_clean["GPA"]

X1_train, X1_test, y1_train, y1_test = train_test_split(
    X1, y1, test_size=0.2, random_state=42
)
print(f"\nTraining samples : {len(X1_train)}")
print(f"Testing samples : {len(X1_test)}")

model1 = LinearRegression()
model1.fit(X1_train, y1_train)

print(f"\nSlope: {model1.coef_[0]:.4f}")
print(f"Intercept: {model1.intercept_:.4f}")
print(f"Equation: GPA = {model1.coef_[0]:.4f} * CIA_pct + {model1.intercept_:.4f}")

```

```

Training samples : 39
Testing samples : 10

Slope: 0.0128
Intercept: 2.4834
Equation: GPA = 0.0128 * CIA_pct + 2.4834

```

```

y1_pred = model1.predict(X1_test)

print("\nPredicted vs Actual GPA (first 10 test samples):")
comparison1 = pd.DataFrame({
    "CIA_pct": X1_test["CIA_pct"].values,
    "Actual GPA": y1_test.values,
    "Predicted GPA": np.round(y1_pred, 4)
})
print(comparison1.to_string(index=False))

```

Predicted vs Actual GPA (first 10 test samples):

CIA_pct	Actual GPA	Predicted GPA
70.0	3.49	3.3794
70.0	3.69	3.3794
66.0	2.74	3.3282
75.0	3.66	3.4434
64.0	3.11	3.3026
71.0	3.33	3.3922
70.0	3.00	3.3794
70.0	3.40	3.3794
80.0	3.60	3.5074
70.0	3.34	3.3794

```
mse1 = mean_squared_error(y1_test, y1_pred)
rmse1 = np.sqrt(mse1)
r2_1 = r2_score(y1_test, y1_pred)
print(f"\nMSE: {mse1:.4f}")
print(f"RMSE: {rmse1:.4f}")
print(f"R2: {r2_1:.4f}")
```

```
MSE: 0.0697
RMSE: 0.2640
R2: 0.1771
```

There is a weak positive relation between CIA and GPA which is 17.7%.
RMSE of 0.26 suggest that model GPA prediction deviate by 0.26.
MSE of 0.06 suggest that error are there and model is not accurate.

Experiment 2

```

X2 = df_clean[["Attendance_pct"]]
y2 = df_clean["GPA"]

X2_train, X2_test, y2_train, y2_test = train_test_split(
    X2, y2, test_size=0.2, random_state=42
)
print(f"\nTraining samples: {len(X2_train)}")
print(f"Testing samples: {len(X2_test)}")

model2 = LinearRegression()
model2.fit(X2_train, y2_train)

print(f"\nSlope: {model2.coef_[0]:.4f}")
print(f"Intercept: {model2.intercept_:.4f}")
print(f"Equation: GPA = {model2.coef_[0]:.4f} * Attendance_pct + {model2.intercept_:.4f}")

```

[207]

Training samples: 39

Testing samples: 10

Slope: 0.0244

Intercept: 1.1407

Equation: GPA = 0.0244 * Attendance_pct + 1.1407

```

y2_pred = model2.predict(X2_test)

print("\nPredicted vs Actual GPA (first 10 test samples):")
comparison2 = pd.DataFrame({
    "Attendance_pct": X2_test["Attendance_pct"].values,
    "Actual GPA": y2_test.values,
    "Predicted GPA": np.round(y2_pred, 4)
})
print(comparison2.to_string(index=False))

mse2 = mean_squared_error(y2_test, y2_pred)
rmse2 = np.sqrt(mse2)
r2_2 = r2_score(y2_test, y2_pred)
print(f"\nMSE: {mse2:.4f}")
print(f"RMSE: {rmse2:.4f}")
print(f"R2: {r2_2:.4f}")

```

[208]

...

Predicted vs Actual GPA (first 10 test samples):

Attendance_pct	Actual GPA	Predicted GPA
100.0	3.49	3.5816
97.0	3.69	3.5084
90.0	2.74	3.3375
95.0	3.66	3.4596
100.0	3.11	3.5816
99.0	3.33	3.5572
93.0	3.00	3.4108
96.0	3.40	3.4840
95.0	3.60	3.4596
95.0	3.34	3.4596

MSE: 0.0922

RMSE: 0.3037

R2: -0.0891

There is a weak negative relation between Attendance and GPA which is -8.9%.
RMSE of 0.26 suggest that model GPA prediction deviate by 0.30.
MSE of 0.09 suggest that error are there and model is not accurate.

```


fig, axes = plt.subplots(1, 2, figsize=(14, 5))

ax1 = axes[0]
ax1.scatter(X1_test, y1_test, color="steelblue", label="Actual", zorder=3)
ax1.plot(X1_test, y1_pred, color="red", linewidth=2, label="Regression line")
ax1.set_title("Exp 1: CIA % vs GPA", fontsize=13)
ax1.set_xlabel("CIA Percentage")
ax1.set_ylabel("GPA")
ax1.legend()
ax1.grid(True, linestyle="--", alpha=0.5)
ax1.text(0.05, 0.92,
        f"y = {model1.coef_[0]:.4f}x + {model1.intercept_:.4f}\nR2 = {r2_1:.4f}",
        transform=ax1.transAxes, fontsize=9,
        bbox=dict(boxstyle="round", facecolor="lightyellow", alpha=0.8))

ax2 = axes[1]
ax2.scatter(X2_test, y2_test, color="darkorange", label="Actual", zorder=3)
ax2.plot(X2_test, y2_pred, color="red", linewidth=2, label="Regression line")
ax2.set_title("Exp 2: Attendance % vs GPA", fontsize=13)
ax2.set_xlabel("Attendance Percentage")
ax2.set_ylabel("GPA")
ax2.legend()
ax2.grid(True, linestyle="--", alpha=0.5)
ax2.text(0.05, 0.92,
        f"y = {model2.coef_[0]:.4f}x + {model2.intercept_:.4f}\nR2 = {r2_2:.4f}",
        transform=ax2.transAxes, fontsize=9,
        bbox=dict(boxstyle="round", facecolor="lightyellow", alpha=0.8))

plt.suptitle("Simple Linear Regression - Student Survey", fontsize=14, fontweight="bold")
plt.tight_layout()
plt.show()

```



Part C: Manual Computation using Ordinary Least Squares (OLS)

Using the same dataset, manually compute the regression line using the OLS formulas given below.

Manual Regression Equation

Construct the regression equation manually using the computed slope and intercept values

```
def ols_manual(x, y):
    x, y = x.values, y.values
    b1 = np.sum((x - x.mean()) * (y - y.mean())) / np.sum((x - x.mean()) ** 2)
    return b1, y.mean() - b1 * x.mean()
```

```
def run_experiment(exp_num, x_col, y_col, df):
    X_train, X_test, y_train, y_test = train_test_split(
        df[[x_col]], df[y_col], test_size=0.2, random_state=42
    )

    sk = LinearRegression().fit(X_train, y_train)
    y_pred_sk = sk.predict(X_test)

    m_slope, m_intercept = ols_manual(X_train[x_col], y_train)
    y_pred_man = m_slope * X_test[x_col].values + m_intercept

    xv, yv = X_train[x_col].values, y_train.values
    xm, ym = xv.mean(), yv.mean()
    print(f"x'={xm:.4f}, y'={ym:.4f}")
    print(f"sum(xi-x')(yi-y') = {np.sum((xv-xm)*(yv-ym)):.4f}")
    print(f"sum(xi-x')^2 = {np.sum((xv-xm)**2):.4f}")

    print(f"\nScikit-learn: {y_col} = {sk.coef_[0]:.6f}*{x_col} + {sk.intercept_:.6f}")
    print(f"Manual OLS: {y_col} = {m_slope:.6f}*{x_col} + {m_intercept:.6f}")

    diff = np.abs(y_pred_sk - y_pred_man)
    print(f"\nMax |diff| = {diff.max():.2e} | Mean |diff| = {diff.mean():.2e}")

    return {
        "x_col": x_col, "y_col": y_col,
        "X_test": X_test, "y_test": y_test,
        "sk_slope": sk.coef_[0], "sk_intercept": sk.intercept_,
        "man_slope": m_slope, "man_intercept": m_intercept,
        "y_pred_sk": y_pred_sk, "y_pred_manual": y_pred_man, "diff": diff
    }
```

0.0s

```

r1 = run_experiment(1,"CIA_pct","GPA",df_clean)
print("\n")
r2 = run_experiment(2,"Attendance_pct","GPA",df_clean)

```

```

x'=73.3956, y`=3.4228
sum(xi-x`)(yi-y`) = 21.9668
sum(xi-x`)2 = 1716.1978

Scikit-learn: GPA = 0.012800*CIA_pct + 2.483380
Manual OLS: GPA = 0.012800*CIA_pct + 2.483380

Max |diff| = 4.44e-16 | Mean |diff| = 1.33e-16

x'=93.4933, y`=3.4228
sum(xi-x`)(yi-y`) = 12.3985
sum(xi-x`)2 = 507.9299

Scikit-learn: GPA = 0.024410*Attendance_pct + 1.140655
Manual OLS: GPA = 0.024410*Attendance_pct + 1.140655

Max |diff| = 4.44e-16 | Mean |diff| = 4.44e-17

```

The manual and sickit learn shows the same results as both use the same forulas so the difference is negligible.

Comparison Task

Compare the following:

1. Predictions from Scikit learn
2. Predictions from manually computed OLS equation
3. Difference between both outputs

Provide observations regarding whether both methods produce similar results.

Comparison Task

```

print(
    """Experiment 1 – CIA % vs GPA:
    Both methods produced identical regression equations:

    Scikit-learn: GPA = 0.012800 * CIA_pct + 2.483380
    Manual OLS: GPA = 0.012800 * CIA_pct + 2.483380

    The prediction values are exactly the same for all 10 test samples.
    For example for CIA_pct = 70.0 both predict 3.3794 and for CIA_pct = 80.0 both predict 3.5074.
    """
)
print()
print(
    """Experiment 2 – Attendance % vs GPA:
    Again both methods produced identical regression equations:

    Scikit-learn: GPA = 0.024410 * Attendance_pct + 1.140655
    Manual OLS: GPA = 0.024410 * Attendance_pct + 1.140655

    All 10 predictions match perfectly.
    For example for Attendance = 100% both predict 3.5816 and for Attendance = 90% both predict 3.3375.
    """
)

```

```

Experiment 1 – CIA % vs GPA:
    Both methods produced identical regression equations:

    Scikit-learn: GPA = 0.012800 * CIA_pct + 2.483380
    Manual OLS: GPA = 0.012800 * CIA_pct + 2.483380

    The prediction values are exactly the same for all 10 test samples.
    For example for CIA_pct = 70.0 both predict 3.3794 and for CIA_pct = 80.0 both predict 3.5074.

Experiment 2 – Attendance % vs GPA:
    Again both methods produced identical regression equations:

    Scikit-learn: GPA = 0.024410 * Attendance_pct + 1.140655
    Manual OLS: GPA = 0.024410 * Attendance_pct + 1.140655

    All 10 predictions match perfectly.
    For example for Attendance = 100% both predict 3.5816 and for Attendance = 90% both predict 3.3375.

```

Parameter Saving Task

Save the parameters learned by the model using Pickle.

The following parameters must be saved:

1. Slope
2. Intercept

Suggested file name:

linear_regression_weights.pkl

Demonstrate:

1. Saving parameters into Pickle file
2. Loading the Pickle file
3. Using loaded parameters for prediction

```
trained = {
    "Experiment 1": {
        "x_col": r1["x_col"],
        "slope": r1["sk_slope"],
        "intercept": r1["sk_intercept"],
        "X_test": r1["X_test"],
        "y_test": r1["y_test"],
    },
    "Experiment 2": {
        "x_col": r2["x_col"],
        "slope": r2["sk_slope"],
        "intercept": r2["sk_intercept"],
        "X_test": r2["X_test"],
        "y_test": r2["y_test"],
    },
}

parameters_to_save = {
    exp_name: {
        "slope": info["slope"],
        "intercept": info["intercept"],
        "x_col": info["x_col"],
    }
    for exp_name, info in trained.items()
}

pickle_file = "linear_regression_weights.pkl"
```

```
pickle_file = "linear_regression_weights.pkl"

with open(pickle_file, "wb") as f:
    pickle.dump(parameters_to_save, f)

print(f"\nParameters saved successfully to '{pickle_file}'\n")
print("Parameters saved:")
for exp_name, params in parameters_to_save.items():
    print(f"\n[{exp_name}] ({params['x_col']} - GPA)")
    print(f"Slope: {params['slope']:.6f}")
    print(f"Intercept: {params['intercept']:.6f}")
```

```
***
Parameters saved successfully to 'linear_regression_weights.pkl'

Parameters saved:

[Experiment 1] (CIA_pct - GPA)
    Slope: 0.012800
    Intercept: 2.483380

[Experiment 2] (Attendance_pct - GPA)
    Slope: 0.024410
    Intercept: 1.140655
```

```

for exp_name, params in loaded_params.items():
    b1 = params["slope"]
    b0 = params["intercept"]
    col = params["x_col"]

    X_test = trained[exp_name]["X_test"]
    y_test = trained[exp_name]["y_test"]

    y_pred = b1 * X_test[col].values + b0

    print(f"\n[{exp_name}] {col} - GPA")
    print(f"Equation: GPA = {b1:.6f} * {col} + {b0:.6f}")
    print()

    result_df = pd.DataFrame({
        col: X_test[col].values,
        "Actual GPA": np.round(y_test.values, 4),
        "Predicted GPA": np.round(y_pred, 4),
        "Error (Actual - Pred)": np.round(y_test.values - y_pred, 4),
    })
    print(result_df.to_string(index=False))

```

[Experiment 1] CIA_pct - GPA
Equation: GPA = 0.012800 * CIA_pct + 2.483380

CIA_pct	Actual GPA	Predicted GPA	Error (Actual - Pred)
70.0	3.49	3.3794	0.1106
70.0	3.69	3.3794	0.3106
66.0	2.74	3.3282	-0.5882
75.0	3.66	3.4434	0.2166
64.0	3.11	3.3026	-0.1926
71.0	3.33	3.3922	-0.0622
70.0	3.00	3.3794	-0.3794
70.0	3.40	3.3794	0.0206
80.0	3.60	3.5074	0.0926
70.0	3.34	3.3794	-0.0394

[Experiment 2] Attendance_pct - GPA
Equation: GPA = 0.024410 * Attendance_pct + 1.140655

Attendance_pct	Actual GPA	Predicted GPA	Error (Actual - Pred)
100.0	3.49	3.5816	-0.0916
97.0	3.69	3.5084	0.1816
90.0	2.74	3.3375	-0.5975
95.0	3.66	3.4596	0.2004
100.0	3.11	3.5816	-0.4716
99.0	3.33	3.5572	-0.2272
93.0	3.00	3.4108	-0.4108
96.0	3.40	3.4840	-0.0840
95.0	3.60	3.4596	0.1404
95.0	3.34	3.4596	-0.1196

Conclusion /inference

Neither CIA percentage nor Attendance percentage individually can predict GPA with acceptable accuracy. CIA percentage performs slightly better (positive R^2) compared to Attendance percentage (negative R^2), but both models are insufficient on their own. There is no difference between Manual and Sickit learn calculations.

Lab 4

Lab Exercise IV:

Aim

- To implement KNN classification on the Breast Cancer dataset and analyze model performance using train-test split, heuristic K selection, cross-validation, ROC-AUC, and classification metrics. Also, to compare classification metrics with regression metrics studied in Linear Regression (Lab 3).
-

Evaluation Rubrics

Submission Guidelines

19. Make a copy of the lab manual template with your `<name_reg:no_subject name>`
20. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
21. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
22. Create a **Git repository** in your profile `<ML lab-reg no>`. Follow a different branch for each lab `<Lab 1, Lab 2 ...>` and push the code to Git.
23. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
24. Upload the PDF to Google Classroom before the deadline.

Code with Results

Problem Statement

A healthcare analytics team is developing a predictive model for early cancer detection. You are required to build a KNN classifier, optimize its performance using different validation techniques, and compare classification evaluation metrics with regression evaluation metrics from Lab 3.

Tasks

Task 1: Data Preparation

1. Load the dataset using sklearn.

2. Convert into DataFrame and explore structure.
3. Check missing values and duplicates.
4. Apply feature scaling using StandardScaler and justify its importance.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LinearRegression
from sklearn.decomposition import PCA
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, confusion_matrix, roc_curve, auc,
                             mean_absolute_error, mean_squared_error, r2_score)

import warnings
warnings.filterwarnings('ignore')
```

✓ 1.6s

Task 1: Data Preparation
Load the dataset using sklearn.

```
df = pd.read_csv('brca.csv')
df.head()
```

Unnamed: 0	x.radius_mean	x.texture_mean	x.perimeter_mean	x.area_mean	x.smoothness_mean	x.compactness_mean	x.concavity_mean	x.concave_pts_mean	x.symmetry_mean	...	x.texture_worst	x.perimet
0	1	13.540	14.36	87.46	566.3	0.09779	0.08129	0.06664	0.047810	0.1885	...	19.26
1	2	13.080	15.71	85.63	520.0	0.10750	0.12700	0.04568	0.031100	0.1967	...	20.49
2	3	9.504	12.44	60.34	273.9	0.10240	0.06492	0.02956	0.020760	0.1815	...	15.66
3	4	13.030	18.42	82.61	523.8	0.08983	0.03766	0.02562	0.029230	0.1467	...	22.81
4	5	8.196	16.84	51.71	201.9	0.08600	0.05943	0.01588	0.005917	0.1769	...	21.96

5 rows x 32 columns

```

> print("Dataset Shape:", df.shape)
> df.info()
[4]
... Dataset Shape: (569, 32)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            569 non-null    int64
1   x.radius_mean                         569 non-null    float64
2   x.texture_mean                        569 non-null    float64
3   x.perimeter_mean                      569 non-null    float64
4   x.area_mean                           569 non-null    float64
5   x.smoothness_mean                     569 non-null    float64
6   x.compactness_mean                    569 non-null    float64
7   x.concavity_mean                      569 non-null    float64
8   x.concave_pts_mean                   569 non-null    float64
9   x.symmetry_mean                       569 non-null    float64
10  x.fractal_dim_mean                    569 non-null    float64
11  x.radius_se                            569 non-null    float64
12  x.texture_se                           569 non-null    float64
13  x.perimeter_se                         569 non-null    float64
14  x.area_se                             569 non-null    float64
15  x.smoothness_se                       569 non-null    float64
16  x.compactness_se                      569 non-null    float64
17  x.concavity_se                        569 non-null    float64
18  x.concave_pts_se                     569 non-null    float64
19  x.symmetry_se                         569 non-null    float64
20  x.fractal_dim_se                      569 non-null    float64
21  x.radius_worst                        569 non-null    float64
22  x.texture_worst                       569 non-null    float64
23  x.perimeter_worst                     569 non-null    float64

```

```

24  x.area_worst                          569 non-null    float64
25  x.smoothness_worst                   569 non-null    float64
26  x.compactness_worst                  569 non-null    float64
27  x.concavity_worst                    569 non-null    float64
28  x.concave_pts_worst                  569 non-null    float64
29  x.symmetry_worst                     569 non-null    float64
30  x.fractal_dim_worst                  569 non-null    float64
31  y                                    569 non-null    object
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB

```


Convert into DataFrame and explore structure.

```
df.drop(columns=['Unnamed: 0'], inplace=True)
print("Columns after dropping index column:", df.shape[1])
```

[5]

```
... Columns after dropping index column: 31
```

Check missing values and duplicates.

 Ge

```
print("Total Missing Values:", df.isnull().sum().sum())
print("Total Duplicate Rows:", df.duplicated().sum())
```

[6]

```
... Total Missing Values: 0
    Total Duplicate Rows: 0
```


No missing values and no duplicate rows in the dataset.

```
print(df['y'].value_counts())

df['y'].value_counts().plot(kind='bar', color=['steelblue', 'darkorange'])
plt.title('Class Distribution (B = Benign, M = Malignant)')
plt.xlabel('Diagnosis')
plt.ylabel('Count')
plt.show()
```

[7]

```
... y
B    357
M    212
Name: count, dtype: int64
```



```
df['target'] = df['y'].map({'B': 0, 'M': 1})

X = df.drop(columns=['y', 'target'])
y = df['target']

print("Features shape:", X.shape)
print("Target value counts:\n", y.value_counts())
```

Features shape: (569, 30)
Target value counts:
target
0 357
1 212
Name: count, dtype: int64

Apply feature scaling using StandardScaler and justify its importance.

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled = pd.DataFrame(X_scaled, columns=X.columns)

X_scaled.describe().loc[['mean', 'std']]
```

	x.radius_mean	x.texture_mean	x.perimeter_mean	x.area_mean	x.smoothness_mean	x.compactness_mean	x.concavity_mean	x.concave_pts_mean	x.symmetry_mean	x.fractal_dim_mean	...	x.radius_worst
mean	-1.998011e-16	-7.492542e-16	-9.990056e-17	-1.998011e-16	1.498508e-16	1.998011e-16	0.000000	9.990056e-17	1.498508e-16	4.745277e-16	...	-9.990056e-17
std	1.000880e+00	1.000880e+00	1.000880e+00	1.000880e+00	1.000880e+00	1.000880e+00	1.000880e+00	1.000880e+00	1.000880e+00	1.000880e+00	...	1.000880e+00

2 rows x 30 columns

KNN classifies points based on distance. Features like area_mean range up to 2500, while features like smoothness_mean range only between 0.05 and 0.2. Without scaling, distance would be dominated by large-range features like area, drowning out the contribution of small-range features. StandardScaler transforms every feature to mean 0 and standard deviation 1, so all 30 features contribute equally to the distance calculation.

Task 2: Train-Test Split Analysis

1. Split dataset into training and testing sets (80:20).
2. Repeat with 70:30 and 90:10 splits.
3. Compare performance variations across splits.
4. Analyze how dataset splitting affects model stability and generalization.


```
splits = {'80:20': 0.2, '70:30': 0.3, '90:10': 0.1}
split_results = {}

for name, test_size in splits.items():
    X_train_s, X_test_s, y_train_s, y_test_s = train_test_split(
        X_scaled, y, test_size=test_size, random_state=42, stratify=y
    )
    knn = KNeighborsClassifier(n_neighbors=5)
    knn.fit(X_train_s, y_train_s)
    acc = accuracy_score(y_test_s, knn.predict(X_test_s))
    split_results[name] = acc
    print(f"{name} split - Train: {len(X_train_s)}, Test: {len(X_test_s)}, Accuracy: {acc:.4f}")
```

80:20 split - Train: 455, Test: 114, Accuracy: 0.9561
70:30 split - Train: 398, Test: 171, Accuracy: 0.9649
90:10 split - Train: 512, Test: 57, Accuracy: 0.9649

```
plt.figure(figsize=(6, 4))
plt.bar(split_results.keys(), split_results.values(), color='seagreen')
plt.title('Accuracy across Train-Test Splits (K=5)')
plt.xlabel('Split Ratio')
plt.ylabel('Accuracy')
plt.ylim(0.90, 1.0)
plt.show()
```

[11]



All three splits give close accuracy since the two classes are well separated in this dataset. The 90:10 split leaves very few test samples, so its accuracy is less reliable as an estimate. The 70:30 split leaves less data for training, which can slightly hurt a distance-based model like KNN. The 80:20 split offers a good balance between training data size and a reliable test estimate, so it is used for the remaining tasks.

```
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42, stratify=y
)
print("Training samples:", len(X_train))
print("Testing samples :", len(X_test))
```

[12]

```
Training samples: 455
Testing samples : 114
```

Task 3: KNN Model with Heuristic K Selection

3.1 Heuristic Method for K Selection

1. Compute initial K using heuristic rule ($K = \sqrt{n}$); where n = number of training samples
2. Use this value as a baseline K.

Task 3: KNN Model with Heuristic K Selection

3.1 Heuristic Method for K Selection

Compute initial K using heuristic rule ($K = \sqrt{n}$); where n = number of training samples
Use this value as a baseline K

```
> ~
n_train = len(X_train)
k_heuristic = int(np.sqrt(n_train))

if k_heuristic % 2 == 0:
    k_heuristic += 1

print(f"Training samples (n): {n_train}")
print(f"sqrt(n) = {np.sqrt(n_train):.2f}")
print(f"Heuristic K (rounded to nearest odd): {k_heuristic}")

[13]

... Training samples (n): 455
      sqrt(n) = 21.33
      Heuristic K (rounded to nearest odd): 21
```

3.2 Model Training

1. Train KNN classifier using heuristic K value.
2. Experiment with nearby values of K ($K \pm 5$).
3. Plot accuracy vs K values.
4. Identify optimal K based on performance trend.

3.2 Model Training

Train KNN classifier using heuristic K value.

Experiment with nearby values of K ($K \pm 5$).

Plot accuracy vs K values.

Identify optimal K based on performance trend.

```
k_values = list(range(max(1, k_heuristic - 5), k_heuristic + 6))
accuracies = []

for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    acc = accuracy_score(y_test, knn.predict(X_test))
    accuracies.append(acc)
    print(f"K={k} - Accuracy: {acc:.4f}")
```


[14]

```
... K=16 - Accuracy: 0.9386
    K=17 - Accuracy: 0.9386
    K=18 - Accuracy: 0.9386
    K=19 - Accuracy: 0.9386
    K=20 - Accuracy: 0.9386
    K=21 - Accuracy: 0.9386
    K=22 - Accuracy: 0.9386
    K=23 - Accuracy: 0.9386
    K=24 - Accuracy: 0.9474
    K=25 - Accuracy: 0.9474
    K=26 - Accuracy: 0.9474
```

```
plt.figure(figsize=(7, 4))
plt.plot(k_values, accuracies, marker='o', color='steelblue')
plt.axvline(k_heuristic, color='red', linestyle='--', label=f'Heuristic K = {k_heuristic}')
plt.title('Accuracy vs K (Train-Test Split)')
plt.xlabel('K value')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

best_k_test = k_values[int(np.argmax(accuracies))]
print(f"Best K based on test-split accuracy: {best_k_test}")
```

[15]



... Best K based on test-split accuracy: 24

Accuracy stays close to its peak around the heuristic K and nearby values, then gradually drops as K moves away from this range, confirming the heuristic gives a reasonable starting point.

3.3 Distance Matrix and Decision Boundary Mapping

1. Explain any two distance metrics used in KNN: **Euclidean Distance** and **Manhattan Distance**. Also mention when each distance metric is suitable.
2. Plot the **decision boundary** of the KNN classifier for different values of K, such as K = 1, K = 5, K = 10, and K = 20. Analyze how the decision boundary changes as K increase

3.3 Distance Matrix and Decision Boundary Mapping

Explain any two distance metrics used in KNN: Euclidean Distance and Manhattan Distance. Also mention when each distance metric is suitable.

Plot the decision boundary of the KNN classifier for different values of K, such as K = 1, K = 5, K = 10, and K = 20. Analyze how the decision boundary changes as K increase

Euclidean Distance: straight-line distance between two points, $d(x, y) = \sqrt{\sum (x_i - y_i)^2}$. It is suitable for continuous, low-to-moderate dimensional, normalized numeric data where the actual magnitude of difference across every feature matters - which fits this scaled cancer dataset well, and is the default metric in KNN.

Manhattan Distance: sum of absolute differences along each axis, $d(x, y) = \sum |x_i - y_i|$, like distance travelled on a city grid. It is suitable for high-dimensional or sparse data, and is less sensitive to large differences in a single feature (outliers), since differences are not squared.

```
pca = PCA(n_components=2)
X_train_2d = pca.fit_transform(X_train)
X_test_2d = pca.transform(X_test)

print("Explained variance by first 2 components:", pca.explained_variance_ratio_.sum().round(4))
```

[16]

... Explained variance by first 2 components: 0.6355

```
def plot_decision_boundary(ax, X, y_labels, k):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X, y_labels)


    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 200), np.linspace(y_min, y_max, 200))

    Z = knn.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    ax.contourf(xx, yy, Z, alpha=0.3, cmap='coolwarm')
    ax.scatter(X[:, 0], X[:, 1], c=y_labels, cmap='coolwarm', edgecolors='k', s=15)
    ax.set_title(f'K = {k}')
    ax.set_xlabel('PC1')
    ax.set_ylabel('PC2')

fig, axes = plt.subplots(1, 4, figsize=(20, 4))
for ax, k in zip(axes, [1, 5, 10, 20]):
    plot_decision_boundary(ax, X_train_2d, y_train.values, k)

plt.suptitle('KNN Decision Boundary on first 2 PCA components')
plt.tight_layout()
plt.show()
```



At K=1 the boundary is jagged and tightly wraps individual training points - it overfits and is sensitive to noise (high variance, low bias). As K increases to 5 and 10, the boundary becomes noticeably smoother and generalizes better. At K=20, the boundary is very smooth and simple, starting to ignore local patterns near the class overlap region (higher bias, lower variance).

Task 4: Cross Validation

1. Apply K-Fold Cross Validation (k = 5 or 10).
2. Compute mean accuracy for different K values.
3. Compare cross-validation results with train-test split results.

4. Select best K based on both heuristic + validation results.

Task 4: Cross Validation

Apply K-Fold Cross Validation (k = 5 or 10).

Compute mean accuracy for different K values.

Compare cross-validation results with train-test split results.

Select best K based on both heuristic + validation results

```
[18]
cv_results_5, cv_results_10 = {}, {}

for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)


    scores5 = cross_val_score(knn, X_scaled, y, cv=KFold(n_splits=5, shuffle=True, random_state=42))
    cv_results_5[k] = scores5.mean()

    scores10 = cross_val_score(knn, X_scaled, y, cv=KFold(n_splits=10, shuffle=True, random_state=42))
    cv_results_10[k] = scores10.mean()

    print(f"K={k} - 5-Fold CV Acc: {scores5.mean():.4f} | 10-Fold CV Acc: {scores10.mean():.4f}")

...
K=16 - 5-Fold CV Acc: 0.9525 | 10-Fold CV Acc: 0.9561
K=17 - 5-Fold CV Acc: 0.9578 | 10-Fold CV Acc: 0.9596
K=18 - 5-Fold CV Acc: 0.9543 | 10-Fold CV Acc: 0.9578
K=19 - 5-Fold CV Acc: 0.9508 | 10-Fold CV Acc: 0.9578
K=20 - 5-Fold CV Acc: 0.9508 | 10-Fold CV Acc: 0.9543
K=21 - 5-Fold CV Acc: 0.9490 | 10-Fold CV Acc: 0.9561
K=22 - 5-Fold CV Acc: 0.9508 | 10-Fold CV Acc: 0.9508
K=23 - 5-Fold CV Acc: 0.9473 | 10-Fold CV Acc: 0.9526
K=24 - 5-Fold CV Acc: 0.9490 | 10-Fold CV Acc: 0.9526
K=25 - 5-Fold CV Acc: 0.9455 | 10-Fold CV Acc: 0.9526
K=26 - 5-Fold CV Acc: 0.9473 | 10-Fold CV Acc: 0.9508
```

```
[19]
plt.figure(figsize=(8, 5))
plt.plot(k_values, accuracies, marker='o', label='Train-Test Split (80:20)')
plt.plot(list(cv_results_5.keys()), list(cv_results_5.values()), marker='s', label='5-Fold CV')
plt.plot(list(cv_results_10.keys()), list(cv_results_10.values()), marker='^', label='10-Fold CV')
plt.xlabel('K value')
plt.ylabel('Accuracy')
plt.title('Train-Test Split vs Cross Validation Accuracy')
plt.legend()
plt.show()
```



```
best_k_cv = max(cv_results_10, key=cv_results_10.get)
print(f"Heuristic K      : {k_heuristic}")
print(f"Best K (Test split) : {best_k_test}")
print(f"Best K (10-Fold CV)   : {best_k_cv} -> Accuracy: {cv_results_10[best_k_cv]:.4f}")

final_k = best_k_cv
print(f"\nFinal K selected for evaluation: {final_k}")
```

[20]

```
... Heuristic K      : 21
     Best K (Test split) : 24
     Best K (10-Fold CV) : 17 -> Accuracy: 0.9596

Final K selected for evaluation: 17
```

Cross-validation accuracy is more stable across K values than a single train-test split, since it averages performance over multiple folds instead of depending on one lucky/unlucky split. The final K is chosen using the 10-fold CV result, which agrees closely with the heuristic K range, confirming the heuristic gave a good starting estimate.

Task 5: Classification Evaluation

Evaluate final model using:

- Accuracy

- Precision
- Recall
- F1 Score
- Confusion Matrix
- ROC Curve and AUC Score

Task 5: Classification Evaluation

Evaluate final model using:

Accuracy

Precision

Recall

F1 Score

Confusion Matrix

ROC Curve and AUC Score

▷ ▾

```
final_knn = KNeighborsClassifier(n_neighbors=final_k)
final_knn.fit(X_train, y_train)
```

```
y_pred = final_knn.predict(X_test)
y_proba = final_knn.predict_proba(X_test)[:, 1]
```

```
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall    : {rec:.4f}")
print(f"F1 Score  : {f1:.4f}")
```

[21]

```
... Accuracy : 0.9386
Precision: 0.9730
Recall    : 0.8571
F1 Score  : 0.9114
```

```
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", cm)


plt.figure(figsize=(4, 4))
plt.imshow(cm, cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.xticks([0, 1], ['Benign', 'Malignant'])
plt.yticks([0, 1], ['Benign', 'Malignant'])

for i in range(2):
    for j in range(2):
        plt.text(j, i, cm[i, j], ha='center', va='center')

plt.colorbar()
plt.show()
```

[22]

```
... Confusion Matrix:
[[71  1]
 [ 6 36]]
```




```

fpr, tpr, thresholds = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(5, 5))
plt.plot(fpr, tpr, color='darkorange', label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()

print(f"AUC Score: {roc_auc:.4f}")

```



```
... AUC Score: 0.9950
```

Accuracy and AUC are both high, but Recall is noticeably lower than Precision here - the confusion matrix shows several malignant cases predicted as benign (false negatives). This is exactly the failure mode that matters most in cancer detection, and it is a direct consequence of choosing a fairly large, smoothed-out a smaller K, or a K chosen specifically to maximize recall instead of overall accuracy, would likely catch more of these malignant cases at a small cost in precision.

Task 6: Comparative Study with Regression (Lab 3 Integration)

Reference: Recall Lab 3 (Linear Regression) where you evaluated:

- MAE
- MSE
- RMSE
- R^2 Score

Tasks:

1. Compare regression evaluation metrics with classification metrics.
2. Explain differences between:
 - Error-based evaluation (Regression)
 - Decision-based evaluation (Classification)
3. Compare:
 - R^2 Score vs Accuracy
 - RMSE vs F1 Score
 - MAE vs Confusion Matrix
4. Discuss differences in evaluation logic for:
 - Continuous prediction tasks
 - Classification tasks

Task 6: Comparative Study with Regression (Lab 3 Integration)

```

lr = LinearRegression()
lr.fit(X_train, y_train)

y_pred_lr = lr.predict(X_test)

mae_lr = mean_absolute_error(y_test, y_pred_lr)
mse_lr = mean_squared_error(y_test, y_pred_lr)
rmse_lr = np.sqrt(mse_lr)
r2_lr = r2_score(y_test, y_pred_lr)

print("Linear Regression on brca.csv (continuous prediction of 0/1 target)")
print(f"MAE: {mae_lr:.4f}")
print(f"MSE: {mse_lr:.4f}")
print(f"RMSE: {rmse_lr:.4f}")
print(f"R2: {r2_lr:.4f}")

```

[27] ✓ 0.0s

```

... Linear Regression on brca.csv (continuous prediction of 0/1 target)
MAE: 0.2010
MSE: 0.0627
RMSE: 0.2503
R2: 0.7307

```

```

print("REGRESSION (Linear) vs CLASSIFICATION (KNN)")
print(f"{'Regression Metric':<25} {'Value':>10} {'Classification Metric':<22} {'Value':>8}")
print(f"{'MAE':<25} {mae_lr:>10.4f} {'Accuracy':<22} {acc:>8.4f}")
print(f"{'MSE':<25} {mse_lr:>10.4f} {'Precision':<22} {prec:>8.4f}")
print(f"{'RMSE':<25} {rmse_lr:>10.4f} {'Recall':<22} {rec:>8.4f}")
print(f"{'R2 Score':<25} {r2_lr:>10.4f} {'F1 Score':<22} {f1:>8.4f}")
print(f"{'-':<25} {'-':>10} {'AUC':<22} {roc_auc:>8.4f}")

```

[28] ✓ 0.0s

Regression Metric	Value	Classification Metric	Value
MAE	0.2010	Accuracy	0.9386
MSE	0.0627	Precision	0.9730
RMSE	0.2503	Recall	0.8571
R2 Score	0.7307	F1 Score	0.9114
-	-	AUC	0.9950

```

fig, axes = plt.subplots(1, 2, figsize=(12, 4))


axes[0].scatter(y_test, y_pred_lr, color='steelblue', alpha=0.6)
axes[0].axhline(0.5, color='red', linestyle='--', label='Decision threshold 0.5')
axes[0].set_xlabel('Actual (0=Benign, 1=Malignant)')
axes[0].set_ylabel('Predicted (continuous)')
axes[0].set_title('Linear Regression: Predicted vs Actual')
axes[0].legend()

reg_vals = [mae_lr, mse_lr, rmse_lr, r2_lr]
clf_vals = [acc, prec, rec, f1]
x = np.arange(4)
w = 0.35
axes[1].bar(x - w/2, reg_vals, w, label='Regression', color='steelblue')
axes[1].bar(x + w/2, clf_vals, w, label='Classification (KNN)', color='darkorange')
axes[1].set_xticks(x)
axes[1].set_xticklabels(['MAE / Accuracy', 'MSE / Precision', 'RMSE / Recall', 'R² / F1'])
axes[1].set_title('Regression Metrics vs Classification Metrics')
axes[1].legend()

plt.tight_layout()
plt.show()

```

[25] ✓ 0.2s



Error-based vs Decision-based evaluation:

Regression metrics (MAE, MSE, RMSE, R^2) are error-based - they measure how far a predicted continuous number is from the actual number, in the same or squared units as the target. Classification metrics (Accuracy, Precision, Recall, F1, ROC-AUC) are decision-based - they measure whether a predicted class label matches the true label, and break the result down by the type of mistake made (false positive vs false negative).

R^2 Score vs Accuracy: R^2 shows how much of the variance in a continuous target the model explains, and can go negative if the model performs worse than simply predicting the mean (as seen in Lab 3 Experiment 2). Accuracy is the fraction of correctly classified samples and is bounded between 0 and 1. They are not on the same scale, but both summarize "how good is the model" within their own task type.

RMSE vs F1 Score: RMSE expresses the average prediction error in the same units as the target and penalizes large errors more heavily; a lower RMSE means a better numeric fit. F1 Score is the harmonic mean of precision and recall and is useful when class distribution is imbalanced, as is the case here (357 Benign vs 212 Malignant).

MAE vs Confusion Matrix: MAE compresses regression error into one average absolute number. A confusion matrix instead breaks classification results into True Positives, True Negatives, False Positives and False Negatives, giving a much more detailed picture of what kind of mistake the model is making - which matters a lot in cancer detection, where missing a malignant case (false negative) is far more costly than a false alarm.

Continuous vs Classification evaluation logic: continuous prediction tasks compare predicted numbers to actual numbers and quantify the size of the gap (error magnitude). Classification tasks compare predicted labels to true labels and quantify correctness of a discrete decision, usually broken down per class rather than reduced to a single distance value.

Inference

Regression metrics (MAE, MSE, RMSE, R^2) measure numeric prediction error magnitude, while classification metrics (Accuracy, Precision, Recall, F1, ROC-AUC) measure decision correctness broken down by error type. In medical diagnosis, Recall and ROC-AUC matter more than Accuracy since missing a true malignant case

Task 7: Analytical Questions

1. Why is KNN called a lazy learning algorithm?
2. Why is feature scaling required in KNN?
3. Explain heuristic K selection using $\sqrt{n}$ rule.
4. Why is cross-validation more reliable than a single train-test split?
5. How does K affect bias-variance trade-off?
6. Why is recall more important than accuracy in cancer prediction?
7. What is the limitation of very large K values?

Task 7: Analytical Questions**1. Why is KNN called a lazy learning algorithm?**

KNN does not build an explicit model during training - it simply stores the training data and only performs computation (distance calculation and voting) at prediction time, hence it is called "lazy".

2. Why is feature scaling required in KNN?

KNN relies entirely on distance calculations. Unscaled features with larger numeric ranges would dominate the distance metric and bias results toward those features, so scaling brings every feature to a comparable range.

3. Explain heuristic K selection using the $\sqrt{n}$ rule.

A common rule of thumb sets K equal to the square root of the number of training samples (n). It gives a reasonable starting K that is large enough to smooth out noise but not so large that it averages over the whole dataset.

4. Why is cross-validation more reliable than a single train-test split?

A single split's accuracy depends on which exact samples happen to land in the test set. Cross-validation trains and tests on multiple different folds and averages the results, giving a more stable, less luck-dependent estimate.

5. How does K affect the bias-variance trade-off?

Small K (e.g. K=1) gives low bias but high variance - it overfits to noise and individual points. Large K gives high bias but low variance - it oversmooths the boundary and can underfit, ignoring local pattern

6. Why is recall more important than accuracy in cancer prediction?

Missing an actual cancer case (false negative) is far more costly than raising a false alarm (false positive). Recall directly measures how many true cancer cases were correctly identified, so optimizing for recall reduces the risk of missing real patients.

7. What is the limitation of very large K values?

Very large K oversmooths the decision boundary and can pull in points from the opposite class near the boundary, increasing bias and risking underfitting. If K becomes too large it starts approaching the majority-class baseline and loses its discriminative power.

Conclusion

```

print("Summary")
print(f"Optimal K (heuristic, sqrt(n)): {k_heuristic}")
print(f"Optimal K (test-split validation): {best_k_test}")
print(f"Optimal K (10-fold cross validation): {best_k_cv}")
print(f"Final K used for evaluation: {final_k}")
print()
print(f"Final Model Performance (K={final_k}):")
print(f"Accuracy: {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall: {rec:.4f}")
print(f"F1 Score: {f1:.4f}")
print(f"AUC: {roc_auc:.4f}")

```

```

[25]
...
Summary
Optimal K (heuristic, sqrt(n)): 21
Optimal K (test-split validation): 24
Optimal K (10-fold cross validation): 17
Final K used for evaluation: 17

Final Model Performance (K=17):
Accuracy: 0.9386
Precision: 0.9730
Recall: 0.8571
F1 Score: 0.9114
AUC: 0.9950

```

Conclusion /inference

All three split ratios (80:20, 70:30, 90:10) gave similar accuracy at $K=5$, confirming $K=\text{root}(n)$ as a good, cross-validation-backed choice, with 80:20 kept as the most reliable balance. The final model scored high on accuracy, precision, F1 and AUC but lower recall missing some malignant cases showing that unlike Lab 3's error-based regression metrics, this classification task must weight recall and ROC-AUC extra heavily since false negatives are far costlier than false alarms.

Lab 5

Lab Exercise V:

Aim

- To implement Linear Regression using the Gradient Descent optimization algorithm and evaluate its performance on a real world dataset.
-

Evaluation Rubrics

Submission Guidelines

25. Make a copy of the lab manual template with your <name_reg:no_subject name>
26. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
27. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
28. Create a **Git repository** in your profile <ML lab-reg no>. Follow a different branch for each lab <Lab 1, Lab 2 ...> and push the code to Git.
29. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
30. Upload the PDF to Google Classroom before the deadline.

Code with Results

Dataset

Dataset Name: Student Performance Dataset

Source: UCI Machine Learning Repository

Dataset Link:

<https://archive.ics.uci.edu/dataset/320/student+performance>

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import warnings
warnings.filterwarnings('ignore')

```

[30] ✓ 0.0s

Tasks to Perform

1. Download and load the Student Performance dataset.

1. Download and load the Student Performance dataset.

```

df = pd.read_csv('student-por.csv', sep=';')
print("Shape:", df.shape)
df.head()

```

✓ 0.0s

Shape: (649, 33)

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	4	3	4	1	1	3	4	0	11	11
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	5	3	3	1	1	3	2	9	11	11
2	GP	F	15	U	LE3	T	1	1	at_home	other	...	4	3	2	2	3	3	6	12	13	12
3	GP	F	15	U	GT3	T	4	2	health	services	...	3	2	2	1	1	5	0	14	14	14
4	GP	F	16	U	GT3	T	3	3	other	other	...	4	3	2	1	2	5	0	11	13	13

5 rows x 33 columns

```

df.info()

```

✓ 0.0s

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 649 entries, 0 to 648
Data columns (total 33 columns):
#   Column          Non-Null Count  Dtype
---  -
0   school          649 non-null    object
1   sex             649 non-null    object
2   age            649 non-null    int64
3   address         649 non-null    object
4   famsize         649 non-null    object
5   Pstatus         649 non-null    object
6   Medu            649 non-null    int64
7   Fedu            649 non-null    int64
8   Mjob            649 non-null    object
9   Fjob            649 non-null    object
10  reason          649 non-null    object
11  guardian        649 non-null    object
12  traveltime      649 non-null    int64
13  studytime       649 non-null    int64
14  failures        649 non-null    int64
15  schoolsup       649 non-null    object
16  famsup          649 non-null    object
17  paid            649 non-null    object
18  activities      649 non-null    object
19  nursery         649 non-null    object
20  higher          649 non-null    object
21  internet        649 non-null    object
22  romantic        649 non-null    object
23  famrel          649 non-null    int64
24  freetime        649 non-null    int64
```

```
25  goout           649 non-null    int64
26  Dalc            649 non-null    int64
27  Walc            649 non-null    int64
28  health          649 non-null    int64
29  absences        649 non-null    int64
30  G1              649 non-null    int64
31  G2              649 non-null    int64
32  G3              649 non-null    int64
dtypes: int64(16), object(17)
memory usage: 167.4+ KB
```

2. Perform data preprocessing, including handling missing values (if any), encoding categorical variables, and feature scaling.

```
print("Missing values per column:")
print(df.isnull().sum().sum(), "total missing values")
df.isnull().sum()[df.isnull().sum() > 0]
```

[33] ✓ 0.0s

... Missing values per column:
0 total missing values

... Series([], dtype: int64)

Inference: The dataset has no missing values, so no imputation is required.

```
binary_maps = {
    'school': {'GP': 0, 'MS': 1},
    'sex': {'F': 0, 'M': 1},
    'address': {'U': 0, 'R': 1},
    'famsize': {'LE3': 0, 'GT3': 1},
    'Pstatus': {'A': 0, 'T': 1},
    'schoolsup': {'no': 0, 'yes': 1},
    'famsup': {'no': 0, 'yes': 1},
    'paid': {'no': 0, 'yes': 1},
    'activities': {'no': 0, 'yes': 1},
    'nursery': {'no': 0, 'yes': 1},
    'higher': {'no': 0, 'yes': 1},
    'internet': {'no': 0, 'yes': 1},
    'romantic': {'no': 0, 'yes': 1},
}

df_enc = df.copy()
for col, mapping in binary_maps.items():
    df_enc[col] = df_enc[col].map(mapping)

nominal_cols = ['Mjob', 'Fjob', 'reason', 'guardian']
df_enc = pd.get_dummies(df_enc, columns=nominal_cols, drop_first=True)

bool_cols = df_enc.select_dtypes(include='bool').columns
df_enc[bool_cols] = df_enc[bool_cols].astype(int)

print("Encoded shape:", df_enc.shape)
df_enc.head()
```

✓ 0.0s

Encoded shape: (649, 42)

Encoded shape: (649, 42)

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	traveltime	studytime	...	Mjob_teacher	Fjob_health	Fjob_other	Fjob_services	Fjob_teacher	reason_home	reason_other	reason_reputation	guardian
0	0	0	18	0	1	0	4	4	2	2	...	0	0	0	0	0	1	0	0	0
1	0	0	17	0	1	1	1	1	1	2	...	0	0	1	0	0	0	0	0	0
2	0	0	15	0	0	1	1	1	1	2	...	0	0	1	0	0	0	1	0	0
3	0	0	15	0	1	1	4	2	1	3	...	0	0	0	1	0	1	0	0	0
4	0	0	16	0	1	1	3	3	1	2	...	0	0	1	0	0	1	0	0	0

5 rows x 42 columns

Inference: Binary yes/no columns (schoolsup, famsup, paid, activities, nursery, higher, internet, romantic) and binary categorical columns (school, sex, address, famsize, Pstatus) are label-encoded to 0/1.
Multi-category nominal columns (Mjob, Fjob, reason, guardian) are one-hot encoded.

Feature Scaling

```
mu = X_train.mean(axis=0)
sigma = X_train.std(axis=0)
sigma[sigma == 0] = 1.0

X_train_scaled = (X_train - mu) / sigma
X_test_scaled = (X_test - mu) / sigma

print("Train mean (should be ~0):", np.round(X_train_scaled.mean(axis=0)[:5], 3))
print("Train std (should be ~1):", np.round(X_train_scaled.std(axis=0)[:5], 3))
```

✓ 0.0s Python

Train mean (should be ~0): [0. 0. 0. 0. 0.]
Train std (should be ~1): [1. 1. 1. 1. 1.]

Inference: Gradient Descent converges much faster (and more reliably) when features are on comparable scales. We standardize features (zero mean, unit variance) using statistics computed only from the training set, then apply the same transform to the test set to avoid data leakage.

3. Select appropriate input features and the target variable.

```
target_col = 'G3'
X = df_enc.drop(columns=[target_col]).values.astype(float)
y = df_enc[target_col].values.astype(float)
feature_names = df_enc.drop(columns=[target_col]).columns.tolist()

print("Number of features:", X.shape[1])
print("Feature names:", feature_names)
```

✓ 0.0s Python

Number of features: 41
Feature names: ['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu', 'traveltime', 'studytime', 'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery', 'higher', 'romantic', 'internet', 'Mjob', 'Fjob', 'reason', 'guardian', 'G1', 'G2']

Inference: Target: G3 as the final grade (0–20).
Features: all remaining columns after encoding (demographics, school-related attributes, social factors, and the first-period and second-period grades G1, G2).

4. Split the dataset into training and testing sets.

4. Split the dataset into training and testing sets.

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
print("Train shape:", X_train.shape, "Test shape:", X_test.shape)
```

✓ 0.0s

Train shape: (519, 41) Test shape: (130, 41)

5. Implement Linear Regression using the Gradient Descent algorithm.

```
def compute_cost(X, y, w, b):
    m = X.shape[0]
    y_pred = X.dot(w) + b
    cost = (1 / (2 * m)) * np.sum((y_pred - y) ** 2)
    return cost

def gradient_descent(X, y, w, b, learning_rate, n_iterations):
    m = X.shape[0]
    cost_history = np.zeros(n_iterations)

    for i in range(n_iterations):
        y_pred = X.dot(w) + b
        error = y_pred - y

        dw = (1 / m) * X.T.dot(error)
        db = (1 / m) * np.sum(error)

        w -= learning_rate * dw
        b -= learning_rate * db

        cost_history[i] = compute_cost(X, y, w, b)

    return w, b, cost_history
```

✓ 0.0s

6. Experiment with different learning rates and observe their effect on model convergence.

```
learning_rates = [0.001, 0.01, 0.05, 0.1, 0.5]
n_iterations = 1000
n_features = X_train_scaled.shape[1]

lr_results = {}

for lr in learning_rates:
    w_init = np.zeros(n_features)
    b_init = 0.0
    w, b, cost_history = gradient_descent(X_train_scaled, y_train, w_init, b_init, lr, n_iterations)
    lr_results[lr] = {'w': w, 'b': b, 'cost_history': cost_history}
    print(f"learning_rate={lr:<6} final_cost={cost_history[-1]:.4f}")
```

✓ 0.1s


```
learning_rate=0.001  final_cost=10.4588
learning_rate=0.01   final_cost=0.7551
learning_rate=0.05   final_cost=0.7456
learning_rate=0.1    final_cost=0.7456
learning_rate=0.5    final_cost=770470986015490805598093097649058496013343545101115883257856.0000
```

```

plt.figure(figsize=(9, 6))
for lr in learning_rates:
    ch = lr_results[lr]['cost_history']
    ch_clipped = np.clip(ch, None, np.nanpercentile(ch[np.isfinite(ch)], 100) if np.any(np.isfinite(ch)) else 1e6)
    plt.plot(ch_clipped, label=f'lr = {lr}')

plt.xlabel('Iteration')
plt.ylabel('Cost J(w, b)')
plt.title('Cost vs. Iterations for Different Learning Rates')
plt.legend()
plt.yscale('log')
plt.grid(alpha=0.3)
plt.show()

```



Inference: very small learning rates (e.g. 0.001) converge slowly and have not reached the minimum within 1000 iterations. Moderate learning rates (0.01 to 0.1) converge smoothly and quickly. Learning rates that are too large (e.g. 0.5, depending on the feature scaling) can cause oscillation or divergence of the cost function as gradient descent needs to take more number of steps to reach minimum.

7. Plot the loss (cost) versus the number of iterations.

```
best_lr = min(learning_rates, key=lambda lr: lr_results[lr]['cost_history'][-1])
print("Selected best learning rate:", best_lr)


n_iterations_final = 3000
w_final = np.zeros(n_features)
b_final = 0.0
w_final, b_final, final_cost_history = gradient_descent(
    X_train_scaled, y_train, w_final, b_final, best_lr, n_iterations_final
)

plt.figure()
plt.plot(final_cost_history, color='darkred')
plt.xlabel('Iteration')
plt.ylabel('Cost J(w, b)')
plt.title(f'Loss vs. Iterations (learning_rate = {best_lr})')
plt.grid(alpha=0.3)
plt.show()

print(f"Final training cost: {final_cost_history[-1]:.4f}")
```

✓ 0.2s

Selected best learning rate: 0.1



8. Evaluate the trained model using:

- Mean Absolute Error (MAE)
- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- R^2 Score

8. Evaluate the trained model using:

Mean Absolute Error (MAE)

Mean Squared Error (MSE)

Root Mean Squared Error (RMSE)

R2 Score

```
y_pred_test = X_test_scaled.dot(w_final) + b_final
```

```
mae = mean_absolute_error(y_test, y_pred_test)
```

```
mse = mean_squared_error(y_test, y_pred_test)
```

```
rmse = np.sqrt(mse)
```

```
r2 = r2_score(y_test, y_pred_test)
```

```
print(f"Mean Absolute Error (MAE): {mae:.4f}")
```

```
print(f"Mean Squared Error (MSE): {mse:.4f}")
```

```
print(f"Root Mean Squared Error: {rmse:.4f}")
```

```
print(f"R2 Score: {r2:.4f}")
```

✓ 0.0s

Mean Absolute Error (MAE): 0.7651


Mean Squared Error (MSE): 1.4759

Root Mean Squared Error: 1.2149

R2 Score: 0.8487

```
plt.figure()
plt.scatter(y_test, y_pred_test, alpha=0.6, edgecolor='k')
lims = [min(y_test.min(), y_pred_test.min()), max(y_test.max(), y_pred_test.max())]
plt.plot(lims, lims, 'r--', label='Perfect prediction')
plt.xlabel('Actual G3')
plt.ylabel('Predicted G3')
plt.title('Actual vs. Predicted Final Grades (Test Set)')
plt.legend()
plt.grid(alpha=0.3)
plt.show()
```

✓ 0.1s



```
coef_df = pd.DataFrame({'feature': feature_names, 'weight': w_final})
coef_df['abs_weight'] = coef_df['weight'].abs()
coef_df = coef_df.sort_values('abs_weight', ascending=False).drop(columns='abs_weight')
coef_df.head(10)
```

✓ 0.0s

	feature	weight
27	G2	2.471735
26	G1	0.460262
34	Fjob_services	-0.240538
6	Medu	-0.166254
33	Fjob_other	-0.163351
0	school	-0.145509
37	reason_other	-0.128880
10	failures	-0.124303
20	freetime	-0.113908
35	Fjob_teacher	-0.107944

Feature Influence (Standardized Coefficients)

Because all features were standardized before training, the magnitude of each weight reflects its relative influence on the prediction.

9. Interpret the results and comment on the convergence behavior and prediction performance of the model.

9. Interpretation of Results

Convergence behavior:

- Learning rates that were too small led to slow convergence — the cost decreased steadily but had not yet reached its minimum within a limited number of iterations.
- Moderate learning rates converged quickly and smoothly toward a stable minimum cost, thanks to feature standardization, which keeps gradients well-scaled across all features.
- Excessively large learning rates risked overshooting the minimum, producing oscillations or even divergence of the cost.
- The final model was trained with the empirically best learning rate for a larger number of iterations, and the loss curve flattens out, indicating convergence.

Prediction performance:

- The **R2 score** indicates the proportion of variance in the final grade (G_3) explained by the model; a value close to 1 indicates strong predictive power, largely because the first- and second-period grades (G_1 , G_2) are highly correlated with the final grade.
- **MAE** gives the average absolute error in grade points, while **RMSE** penalizes larger errors more heavily; comparing the two indicates whether large outlier errors are present (RMSE noticeably higher than MAE suggests some larger misses).
- The scatter plot of actual vs. predicted grades shows points clustering closely around the diagonal (perfect-prediction) line, confirming that the linear model captures the dominant relationship between the features and the target.

Conclusion /inference

Gradient Descent successfully found a good approximation of the closed-form (normal equation) linear regression solution, and the choice of learning rate is critical: too small wastes iterations, too large risks instability, while feature scaling greatly widens the range of workable learning rates.

Lab 6

Lab Exercise VI:

Aim

- To implement Logistic Regression and K Nearest Neighbors (KNN) classifiers on the Breast Cancer Wisconsin (Diagnostic) dataset and compare their performance using standard classification evaluation metrics.
-

Evaluation Rubrics

Submission Guidelines

31. Make a copy of the lab manual template with your `<name_reg:no_subject name>`
32. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
33. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
34. Create a **Git repository** in your profile `<ML lab-reg no>`. Follow a different branch for each lab `<Lab 1, Lab 2 ...>` and push the code to Git.
35. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
36. Upload the PDF to Google Classroom before the deadline.

Code with Results

Dataset

Dataset Name: Breast Cancer Wisconsin (Diagnostic)

Source: UCI Machine Learning Repository

Dataset Link:

<https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic>

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import seaborn as sns
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings('ignore')

```

[4] ✓ 0.5s

1. Download and load the Breast Cancer Wisconsin (Diagnostic) dataset.

1. Download and load the Breast Cancer Wisconsin (Diagnostic) dataset.

```

LOCAL_PATH = "wdbc.data"

feature_stats = ["mean", "se", "worst"]
feature_names = ["radius", "texture", "perimeter", "area", "smoothness",
                 "compactness", "concavity", "concave_points", "symmetry", "fractal_dimension"]

columns = ["id", "diagnosis"] + [f"{stat}_{name}" for stat in feature_stats for name in feature_names]

df = pd.read_csv(LOCAL_PATH, header=None, names=columns)
print("Loaded dataset from local wdbc.data file.")

df.head()

```

✓ 0.0s

Loaded dataset from local wdbc.data file.

	id	diagnosis	mean_radius	mean_texture	mean_perimeter	mean_area	mean_smoothness	mean_compactness	mean_concavity	mean_concave_points	...	worst_radius	worst_texture	worst_perimeter
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	...	25.38	17.33	184.60
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...	24.99	23.41	158.80
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...	23.57	25.53	152.50
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	...	14.91	26.50	98.87
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	...	22.54	16.67	152.20

5 rows × 32 columns

Inference: The dataset loaded successfully and matches the documented structure, one `id` column, one `diagnosis` target column (M/B), and 30 numeric predictor columns derived from digitized images of fine needle aspirates (FNA) of breast masses. Each row represents one patient sample.

2. Perform exploratory data analysis and data preprocessing.

2. Perform exploratory data analysis and data preprocessing.

```

print("Shape:", df.shape)
df.info()

```

[6] ✓ 0.0s

```

Shape: (569, 32)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   id                                    569 non-null    int64
 1   diagnosis                            569 non-null    object
 2   mean_radius                          569 non-null    float64
 3   mean_texture                         569 non-null    float64
 4   mean_perimeter                      569 non-null    float64
 5   mean_area                           569 non-null    float64
 6   mean_smoothness                     569 non-null    float64
 7   mean_compactness                    569 non-null    float64
 8   mean_concavity                      569 non-null    float64
 9   mean_concave_points                569 non-null    float64
10   mean_symmetry                       569 non-null    float64
11   mean_fractal_dimension              569 non-null    float64
12   se_radius                           569 non-null    float64
13   se_texture                          569 non-null    float64
14   se_perimeter                       569 non-null    float64
15   se_area                             569 non-null    float64
16   se_smoothness                       569 non-null    float64
17   se_compactness                      569 non-null    float64
18   se_concavity                       569 non-null    float64
19   se_concave_points                  569 non-null    float64
20   se_symmetry                         569 non-null    float64
21   se_fractal_dimension                569 non-null    float64
22   worst_radius                        569 non-null    float64
23   worst_texture                       569 non-null    float64

```

```

24   worst_perimeter                    569 non-null    float64
25   worst_area                         569 non-null    float64
26   worst_smoothness                   569 non-null    float64
27   worst_compactness                  569 non-null    float64
28   worst_concavity                    569 non-null    float64
29   worst_concave_points               569 non-null    float64
30   worst_symmetry                     569 non-null    float64
31   worst_fractal_dimension             569 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB

```

Inference: The dataset contains 569 samples and 32 columns (id, diagnosis, and 30 numeric features). All feature columns are numeric (float64), which confirms no type-conversion issues before modeling.

3. Check for missing values and prepare the dataset for classification.

3. Check for missing values and prepare the dataset for classification.

```
missing = df.isnull().sum()
print("Missing values per column:")
print(missing[missing > 0] if missing.sum() > 0 else "No missing values found in any column.")
```

✓ 0.0s

Missing values per column:
No missing values found in any column.

Inference: There are no missing values anywhere in the dataset.

```
print(df["diagnosis"].value_counts())
print()
print(df["diagnosis"].value_counts(normalize=True).round(3) * 100)
```

[8] ✓ 0.0s

```
diagnosis
B    357
M    212
Name: count, dtype: int64
```

```
diagnosis
B    62.7
M    37.3
Name: proportion, dtype: float64
```

Inference: The classes are moderately imbalanced roughly 63% Benign vs. 37% Malignant. This is not severe enough to require resampling techniques, but it does mean accuracy alone can be misleading, and metrics like precision, recall, and F1-score (computed for the malignant/positive class) are important for a fair evaluation.

```
df.describe().T
```

✓ 0.0s

	count	mean	std	min	25%	50%	75%	max
id	569.0	3.037183e+07	1.250206e+08	8670.000000	869218.000000	906024.000000	8.813129e+06	9.113205e+08
mean_radius	569.0	1.412729e+01	3.524049e+00	6.981000	11.700000	13.370000	1.578000e+01	2.811000e+01
mean_texture	569.0	1.928965e+01	4.301036e+00	9.710000	16.170000	18.840000	2.180000e+01	3.928000e+01
mean_perimeter	569.0	9.196903e+01	2.429898e+01	43.790000	75.170000	86.240000	1.041000e+02	1.885000e+02
mean_area	569.0	6.548891e+02	3.519141e+02	143.500000	420.300000	551.100000	7.827000e+02	2.501000e+03
mean_smoothness	569.0	9.636028e-02	1.406413e-02	0.052630	0.086370	0.095870	1.053000e-01	1.634000e-01
mean_compactness	569.0	1.043410e-01	5.281276e-02	0.019380	0.064920	0.092630	1.304000e-01	3.454000e-01
mean_concavity	569.0	8.879932e-02	7.971981e-02	0.000000	0.029560	0.061540	1.307000e-01	4.268000e-01
mean_concave_points	569.0	4.891915e-02	3.880284e-02	0.000000	0.020310	0.033500	7.400000e-02	2.012000e-01
mean_symmetry	569.0	1.811619e-01	2.741428e-02	0.106000	0.161900	0.179200	1.957000e-01	3.040000e-01
mean_fractal_dimension	569.0	6.279761e-02	7.060363e-03	0.049960	0.057700	0.061540	6.612000e-02	9.744000e-02
se_radius	569.0	4.051721e-01	2.773127e-01	0.111500	0.232400	0.324200	4.789000e-01	2.873000e+00
se_texture	569.0	1.216853e+00	5.516484e-01	0.360200	0.833900	1.108000	1.474000e+00	4.885000e+00
se_perimeter	569.0	2.866059e+00	2.021855e+00	0.757000	1.606000	2.287000	3.357000e+00	2.198000e+01
se_area	569.0	4.033708e+01	4.549101e+01	6.802000	17.850000	24.530000	4.519000e+01	5.422000e+02
se_smoothness	569.0	7.040979e-03	3.002518e-03	0.001713	0.005169	0.006380	8.146000e-03	3.113000e-02
se_compactness	569.0	2.547814e-02	1.790818e-02	0.002252	0.013080	0.020450	3.245000e-02	1.354000e-01
se_concavity	569.0	3.189372e-02	3.018606e-02	0.000000	0.015090	0.025890	4.205000e-02	3.960000e-01
se_concave_points	569.0	1.179614e-02	6.170285e-03	0.000000	0.007638	0.010930	1.471000e-02	5.279000e-02
se_symmetry	569.0	2.054230e-02	8.266372e-03	0.007882	0.015160	0.018730	2.348000e-02	7.895000e-02
se_fractal_dimension	569.0	3.794904e-03	2.646071e-03	0.000895	0.002248	0.003187	4.558000e-03	2.984000e-02
worst_radius	569.0	1.626919e+01	4.833242e+00	7.930000	13.010000	14.970000	1.879000e+01	3.604000e+01
worst_texture	569.0	2.567722e+01	6.146258e+00	12.020000	21.080000	25.410000	2.972000e+01	4.954000e+01
worst_perimeter	569.0	1.072612e+02	3.360254e+01	50.410000	84.110000	97.660000	1.254000e+02	2.512000e+02
worst_area	569.0	8.805831e+02	5.693570e+02	185.200000	515.300000	686.500000	1.084000e+03	4.254000e+03
worst_smoothness	569.0	1.323686e-01	2.283243e-02	0.071170	0.116600	0.131300	1.460000e-01	2.226000e-01
worst_compactness	569.0	2.542650e-01	1.573365e-01	0.027290	0.147200	0.211900	3.391000e-01	1.058000e+00
worst_concavity	569.0	2.721885e-01	2.086243e-01	0.000000	0.114500	0.226700	3.829000e-01	1.252000e+00
worst_concave_points	569.0	1.146062e-01	6.573234e-02	0.000000	0.064930	0.099930	1.614000e-01	2.910000e-01
worst_symmetry	569.0	2.900756e-01	6.186747e-02	0.156500	0.250400	0.282200	3.179000e-01	6.638000e-01
worst_fractal_dimension	569.0	8.394582e-02	1.806127e-02	0.055040	0.071460	0.080040	9.208000e-02	2.075000e-01

Inference: The features span very different scales (e.g. `mean_area` is in the hundreds/thousands while `mean_smoothness` is a small fraction). This confirms that feature scaling (standardization) is necessary, especially for a distance-based algorithm like KNN, where features with larger numeric ranges would otherwise dominate the distance calculation.

```
sns.set_style("whitegrid")


fig, axes = plt.subplots(1, 2, figsize=(12, 4.5))

df["diagnosis"].value_counts().plot(kind="bar", color=[ "#4C72B0", "#D08452"], ax=axes[0])
axes[0].set_title("Class Distribution")
axes[0].set_xlabel("Diagnosis")
axes[0].set_ylabel("Count")
axes[0].set_xticklabels(["Benign (B)" if l.get_text()=="B" else "Malignant (M)" for l in axes[0].get_xticklabels()], rotation=0)

corr = df.drop(columns=["id"]).corr(numeric_only=True)
top_corr_features = corr[["diagnosis"]].abs().sort_values(ascending=False).index[1:11] if "diagnosis" in corr.columns else None

sns.heatmap(df[["mean_radius", "mean_texture", "mean_perimeter", "mean_area",
                "mean_smoothness", "mean_compactness", "mean_concavity"]].corr(),
            annot=True, fmt=".2f", cmap="coolwarm", ax=axes[1])
axes[1].set_title("Correlation Among Selected Mean Features")

plt.tight_layout()
plt.show()
```



Inference: The class distribution plot confirms the moderate imbalance noted earlier. The correlation heatmap shows very strong positive correlation among `mean_radius`, `mean_perimeter`, and `mean_area` (as expected, since these are geometrically related measurements of the same tumor). This multicollinearity mainly affects the interpretability of Logistic Regression coefficients (not its predictive power), and is naturally handled by KNN since it relies on overall distance rather than individual coefficients.

```

y = df["diagnosis"].map({"M": 1, "B": 0})
X = df.drop(columns=["id", "diagnosis"])

print("Feature matrix shape:", X.shape)
print("Target vector shape:", y.shape)
print("\nEncoded target distribution:")
print(y.value_counts().rename({1: "Malignant (1)", 0: "Benign (0)"}))

```

1] ✓ 0.0s

```

Feature matrix shape: (569, 30)
Target vector shape: (569,)

Encoded target distribution:
diagnosis
Benign (0)      357
Malignant (1)   212
Name: count, dtype: int64

```

Inference: The target is encoded so that **Malignant = 1** (the positive/at-risk class) and **Benign = 0**. This choice matters clinically: in cancer diagnosis, recall on the malignant class (i.e., not missing a cancer case) is usually the more critical metric than raw accuracy, since a false negative (missing a malignant tumor) is far more costly than a false positive.

4. Split the dataset into training and testing sets.

4. Split the dataset into training and testing sets.

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print("Training set size:", X_train.shape[0])
print("Test set size:", X_test.shape[0])
print("\nTraining class ratio:\n", y_train.value_counts(normalize=True).round(3))
print("\nTest class ratio:\n", y_test.value_counts(normalize=True).round(3))

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
Training set size: 455
Test set size: 114

Training class ratio:
diagnosis
0    0.626
1    0.374
Name: proportion, dtype: float64

Test class ratio:
diagnosis
0    0.632
1    0.368
Name: proportion, dtype: float64
```

Inference: Stratified splitting confirms the class ratio is preserved almost identically in both sets, so performance differences between train and test won't be an artifact of an accidental imbalance shift. Standardization is fit only on the training set and then applied to the test set this avoids **data leakage**, ensuring the test set truly represents unseen data. Scaling is essential for KNN (distance-based) and also improves numerical stability and convergence for Logistic Regression's gradient-based solver.

5. Train a Logistic Regression classifier.

5. Train a Logistic Regression classifier.

```
from sklearn.linear_model import LogisticRegression

log_reg = LogisticRegression(max_iter=10000, random_state=42)
log_reg.fit(X_train_scaled, y_train)

y_pred_log = log_reg.predict(X_test_scaled)
print("Logistic Regression trained successfully.")
```

3] ✓ 0.3s

Inference: Logistic Regression models the log-odds of malignancy as a linear combination of the (scaled) features. Because it is a linear, well-regularized model, it tends to generalize well on datasets like this one where classes are largely linearly separable.

6. Train a K Nearest Neighbors (KNN) classifier.

6. Train a K Nearest Neighbors (KNN) classifier

```
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train_scaled, y_train)

y_pred_knn = knn.predict(X_test_scaled)
print("KNN (k=5) trained successfully.")
```

4] ✓ 0.5s

KNN (k=5) trained successfully.

```

from sklearn.metrics import accuracy_score


k_range = range(1, 21)
k_scores = []
for k in k_range:
    model = KNeighborsClassifier(n_neighbors=k)
    model.fit(X_train_scaled, y_train)
    pred = model.predict(X_test_scaled)
    k_scores.append(accuracy_score(y_test, pred))

plt.figure(figsize=(7, 4))
plt.plot(list(k_range), k_scores, marker="o")
plt.xlabel("k (Number of Neighbors)")
plt.ylabel("Test Accuracy")
plt.title("KNN: Test Accuracy vs. k")
plt.xticks(list(k_range))
plt.grid(True)
plt.show()

best_k = list(k_range)[int(np.argmax(k_scores))]
print(f"Best k found: {best_k} with accuracy {max(k_scores):.4f}")

```

✓ 0.6s



Best k found: 5 with accuracy 0.9561

Inference: The accuracy-vs-k curve shows how sensitive KNN is to the choice of neighborhood size very small k can overfit to noise, while very large k oversmooths the decision boundary and can blur the boundary between classes. The plot lets us pick a k that balances bias and variance. We proceed with the original $k=5$ model for the main comparison below.

7. Evaluate both classifiers using:

- Accuracy
- Precision
- Recall
- F1 Score
- Confusion Matrix

7. Evaluate both classifiers using:

Accuracy
Precision
Recall
F1 Score
Confusion Matrix

```
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, confusion_matrix, classification_report)

def evaluate_model(name, y_true, y_pred):
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    cm = confusion_matrix(y_true, y_pred)
    print(f"{name}")
    print(f"Accuracy: {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall: {rec:.4f}")
    print(f"F1 Score: {f1:.4f}")
    print(f"\nConfusion Matrix:")
    print(cm)
    print(f"\nClassification Report:")
    print(classification_report(y_true, y_pred, target_names=["Benign (0)", "Malignant (1)"]))
    return {"Model": name, "Accuracy": acc, "Precision": prec, "Recall": rec, "F1 Score": f1, "Confusion Matrix": cm}

results_log = evaluate_model("Logistic Regression", y_test, y_pred_log)
results_knn = evaluate_model("K-Nearest Neighbors (k=5)", y_test, y_pred_knn)
```

16] ✓ 0.0s

```
Logistic Regression
Accuracy: 0.9649
Precision: 0.9750
Recall: 0.9286
F1 Score: 0.9512
```

```
Confusion Matrix:
[[71  1]
 [ 3 39]]
```

```

Classification Report:
              precision    recall  f1-score   support

   Benign (0)       0.96      0.99      0.97        72
  Malignant (1)       0.97      0.93      0.95        42

   accuracy              0.96        114
  macro avg       0.97      0.96      0.96        114
 weighted avg       0.97      0.96      0.96        114

```

K-Nearest Neighbors (k=5)

Accuracy: 0.9561

Precision: 0.9744

Recall: 0.9048

F1 Score: 0.9383

Confusion Matrix:

```

[[71  1]
 [ 4 38]]

```

```

Classification Report:
              precision    recall  f1-score   support

   Benign (0)       0.95      0.99      0.97        72
  Malignant (1)       0.97      0.90      0.94        42

   accuracy              0.96        114
  macro avg       0.96      0.95      0.95        114
 weighted avg       0.96      0.96      0.96        114

```

```


fig, axes = plt.subplots(1, 2, figsize=(11, 4.5))

for ax, res in zip(axes, [results_log, results_knn]):
    sns.heatmap(res["Confusion Matrix"], annot=True, fmt="d", cmap="Blues",
                xticklabels=["Benign", "Malignant"], yticklabels=["Benign", "Malignant"], ax=ax)
    ax.set_title(res["Model"])
    ax.set_xlabel("Predicted")
    ax.set_ylabel("Actual")

plt.tight_layout()
plt.show()

```

✓ 0.2s



Inference: The confusion matrices let us directly see the count of false negatives (malignant cases predicted as benign the most clinically dangerous error) and false positives (benign cases predicted as malignant leads to unnecessary follow-up but is less dangerous) for each model. A strong classifier for this problem should keep the false-negative count as close to zero as possible.

```
corrs = X.corrwith(y).abs().sort_values(ascending=False)
print(corrs.head(5))

best_feature = corrs.index[0]
print(f"\nUsing '{best_feature}' for the logistic curve visualization")

X_single = X[[best_feature]].values
y_single = y.values

simple_log_reg = LogisticRegression()
simple_log_reg.fit(X_single, y_single)

x_range = np.linspace(X_single.min(), X_single.max(), 300).reshape(-1, 1)
y_prob = simple_log_reg.predict_proba(x_range)[:, 1]

plt.figure(figsize=(8, 5))

jitter = np.random.normal(0, 0.02, size=len(y_single))
plt.scatter(X_single, y_single + jitter, c=y_single, cmap="coolwarm",
            alpha=0.5, edgecolor="k", linewidth=0.3,
            label="Actual data (0=Benign, 1=Malignant)")

plt.plot(x_range, y_prob, color="black", linewidth=2.5,
         label="Logistic Regression curve (P(Malignant))")


plt.axhline(0.5, color="gray", linestyle="--", linewidth=1,
            label="Decision threshold (0.5)")

plt.xlabel(best_feature)
plt.ylabel("Probability of Malignant")
plt.title(f"Logistic Regression Curve - {best_feature}")
plt.legend(loc="center right")
plt.tight_layout()
plt.show()
```



```
worst_concave_points    0.793566
worst_perimeter         0.782914
mean_concave_points     0.776614
worst_radius            0.776454
mean_perimeter          0.742636
dtype: float64
```

Using 'worst_concave_points' for the logistic curve visualization



Inference: The S-shaped curve shows how the predicted probability of malignancy rises as {best_feature} increases. Data points cluster near 0 (Benign) at low values and near 1 (Malignant) at high values, and the curve crosses the 0.5 decision threshold at the point that separates the two classes.

markd

```

from sklearn.neighbors import KNeighborsClassifier
import numpy as np
import matplotlib.pyplot as plt

f1, f2 = "mean_radius", "mean_texture"
X_2d = X[[f1, f2]].values
y_2d = y.values

knn_2d = KNeighborsClassifier(n_neighbors=5)
knn_2d.fit(X_2d, y_2d)


x_min, x_max = X_2d[:, 0].min() - 1, X_2d[:, 0].max() + 1
y_min, y_max = X_2d[:, 1].min() - 1, X_2d[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

Z = knn_2d.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.figure(figsize=(7, 5))
plt.contourf(xx, yy, Z, alpha=0.4, cmap="coolwarm")
plt.scatter(X_2d[:, 0], X_2d[:, 1], c=y_2d, cmap="coolwarm", edgecolor="k")
plt.xlabel(f1)
plt.ylabel(f2)
plt.title("KNN Decision Boundary (k=5)")
plt.show()

```

✓ 1.6s



Inference: The KNN decision boundary is irregular and non-linear — unlike a straight line, it bends and curves around clusters of points, since KNN classifies each region based on the majority class among its nearest neighbors rather than a fixed formula. This makes KNN flexible enough to capture complex, local patterns in the data, but it also means the boundary can become jagged near overlapping points, making the model more sensitive to noise in that region.

8. Compare the performance of the two classifiers using a comparison table

8. Compare the performance of the two classifiers using a comparison table.

```
comparison_df = pd.DataFrame([
    {"Model": "Logistic Regression",
     "Accuracy": results_log["Accuracy"],
     "Precision": results_log["Precision"],
     "Recall": results_log["Recall"],
     "F1 Score": results_log["F1 Score"]},
    {"Model": "K-Nearest Neighbors (k=5)",
     "Accuracy": results_knn["Accuracy"],
     "Precision": results_knn["Precision"],
     "Recall": results_knn["Recall"],
     "F1 Score": results_knn["F1 Score"]},
]).set_index("Model").round(4)
```


comparison_df

✓ 0.0s

	Accuracy	Precision	Recall	F1 Score
Model				
Logistic Regression	0.9649	0.9750	0.9286	0.9512
K-Nearest Neighbors (k=5)	0.9561	0.9744	0.9048	0.9383

```
comparison_df.plot(kind="bar", figsize=(8, 5), ylim=(0.85, 1.01))
plt.title("Logistic Regression vs. KNN – Metric Comparison")
plt.ylabel("Score")
plt.xticks(rotation=0)
plt.legend(loc="lower right")
plt.tight_layout()
plt.show()
```

✓ 0.1s



Inference: The comparison table and chart summarize all four metrics side by side. Look specifically at which model has higher Recall on the malignant class since missing a cancer diagnosis (false negative) is the costliest type of error in this domain, the model with higher recall (even if its overall accuracy is marginally lower) may be the more clinically preferable choice.

9. Interpret the results and comment on the better-performing classifier for the given dataset.

9. Interpret the results and comment on the better-performing classifier for the given dataset.

Accuracy tells us overall correctness, but with ~63/37 class imbalance it can slightly overstate performance — a model that leans toward predicting the majority class (Benign) can still score reasonably well on accuracy alone. **Precision** (for Malignant) tells us: of all cases the model flagged as malignant, how many truly were. Low precision means more unnecessary alarm/follow-up testing for healthy patients. **Recall** (for Malignant) tells us: of all truly malignant cases, how many the model actually caught. This is the single most important metric clinically, since a missed malignant case (false negative) can have severe consequences. **F1-score** balances precision and recall into one number, useful when you care about both types of error jointly.

Typical outcome for this dataset: Logistic Regression usually performs at least as well as, and often slightly better than, KNN on WDBC. This matches the original UCI documentation, which notes the classes are **linearly separable** using the 30 features a condition that favors a linear model like Logistic Regression. Logistic Regression also tends to be more robust here because:

1. It naturally handles the correlated/redundant features (e.g. radius, perimeter, area) without being thrown off by local distance distortions the way KNN can be in high-dimensional, correlated feature spaces (the "curse of dimensionality").
2. It produces well-calibrated probability outputs, useful for setting a custom decision threshold if, for example, we want to favor recall over precision (as is clinically desirable here).
3. It is more interpretable — the sign and magnitude of coefficients indicate which measurements (e.g. worst concavity, worst radius) push predictions toward malignant.

KNN, meanwhile, is simple and non-parametric (makes no assumption about the decision boundary shape), and can still perform competitively once features are properly scaled but it is more sensitive to the choice of k , more computationally expensive at prediction time, and less interpretable.

Self Learning

```
from sklearn.model_selection import StratifiedKFold, cross_val_score, cross_validate
import numpy as np

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

✓ 0.0s

```
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier

log_reg_pipeline = Pipeline([
    ("scaler", StandardScaler()),
    ("clf", LogisticRegression(max_iter=10000, random_state=42))
])

knn_pipeline = Pipeline([
    ("scaler", StandardScaler()),
    ("clf", KNeighborsClassifier(n_neighbors=5))
])
```

✓ 0.0s

```
scoring = ["accuracy", "precision", "recall", "f1"]

log_cv_results = cross_validate(log_reg_pipeline, X, y, cv=cv, scoring=scoring)
knn_cv_results = cross_validate(knn_pipeline, X, y, cv=cv, scoring=scoring)
```

✓ 0.1s

```
def summarize_cv(name, results):
    print(f"==== {name} (5-Fold CV) =====")
    for metric in scoring:
        scores = results[f"test_{metric}"]
        print(f"{metric.capitalize():10s}: {scores.mean():.4f} ± {scores.std():.4f}")
    print()

summarize_cv("Logistic Regression", log_cv_results)
summarize_cv("KNN (k=5)", knn_cv_results)
```

✓ 0.0s

```

===== Logistic Regression (5-Fold CV) =====
Accuracy : 0.9737 ± 0.0166
Precision : 0.9863 ± 0.0182
Recall    : 0.9436 ± 0.0525
F1        : 0.9633 ± 0.0240

===== KNN (k=5) (5-Fold CV) =====
Accuracy : 0.9631 ± 0.0179
Precision : 0.9755 ± 0.0217
Recall    : 0.9246 ± 0.0458
F1        : 0.9487 ± 0.0255

```

```

comparison_cv_df = pd.DataFrame({
    "Logistic Regression": [log_cv_results[f"test_{m}"].mean() for m in scoring],
    "KNN (k=5)": [knn_cv_results[f"test_{m}"].mean() for m in scoring],
}, index=[m.capitalize() for m in scoring])

```

```
comparison_cv_df.round(4)
```

✓ 0.0s

	Logistic Regression	KNN (k=5)
Accuracy	0.9737	0.9631
Precision	0.9863	0.9755
Recall	0.9436	0.9246
F1	0.9633	0.9487

Inference: Across 5-fold cross-validation, Logistic Regression outperforms KNN (k=5) on every single metric accuracy (97.4% vs 96.3%), precision (98.6% vs 97.6%), recall (94.4% vs 92.5%), and F1-score (96.3% vs 94.9%) — confirming that this advantage is consistent across different subsets of the data, not just a result of one lucky train/test split.

Conclusion /inference

Logistic Regression outperformed KNN across every metric (Accuracy 96.5% vs 95.6%, Recall 92.9% vs 90.5%, F1 95.1% vs 93.8%), and since it also caught more true malignant cases (higher recall), Logistic Regression is the better classifier for this dataset both in overall performance and in the clinically critical goal of minimizing missed cancer diagnoses.

Lab 7

Lab Exercise VII:

Aim

- Implement Decision Tree Classifier.
-

Evaluation Rubrics

Submission Guidelines

37. Make a copy of the lab manual template with your `<name_reg:no_subject name>`
38. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
39. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
40. Create a **Git repository** in your profile `<ML lab-reg no>`. Follow a different branch for each lab `<Lab 1, Lab 2 ...>` and push the code to Git.
41. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
42. Upload the PDF to Google Classroom before the deadline.

Code with Results

Dataset

Dataset Name: Iris Dataset (`sklearn.datasets.load_iris()`)

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

Task 1: Dataset Exploration

Load the Iris dataset and answer the following:

1. How many samples and features are present in the dataset?
2. List the feature names and target classes.
3. Display the first five records.
4. Display the class distribution.

Task 1: Dataset Exploration

```
iris = load_iris()
X = iris.data
y = iris.target
feature_names = iris.feature_names
target_names = iris.target_names

df = pd.DataFrame(X, columns=feature_names)
df["target"] = y
df["species"] = df["target"].apply(lambda i: target_names[i])

print(f"Number of samples : {X.shape[0]}")
print(f"Number of features: {X.shape[1]}")
print(f"Feature names      : {feature_names}")
print(f"Target classes     : {list(target_names)}")
```

... Number of samples : 150
 Number of features: 4
 Feature names : ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
 Target classes : [np.str_('setosa'), np.str_('versicolor'), np.str_('virginica')]

The dataset contains of 150 rows and 4 columns, and is about iris dataset which contains 3 type of species that are setosa, versicolor and virginica.

[4] ✓ Os

```
print("First five records:")
df.head()
```

... First five records:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	species
0	5.1	3.5	1.4	0.2	0	setosa
1	4.9	3.0	1.4	0.2	0	setosa
2	4.7	3.2	1.3	0.2	0	setosa
3	4.6	3.1	1.5	0.2	0	setosa
4	5.0	3.6	1.4	0.2	0	setosa

Next steps: [Generate code with df](#) [New interactive sheet](#)

Displaying first 5 rows from the dataset.


```
[5] print("Class distribution:")
df["species"].value_counts()
```

Class distribution:

species	count
setosa	50
versicolor	50
virginica	50

dtype: int64

All the three type of species are equally distributed with each species containing 50 samples.

Task 2: Data Preparation

1. Split the dataset into training (80%) and testing (20%) sets using `random_state=42`.
2. Briefly explain why the dataset is divided into training and testing sets.

Task 2: Data Preparation

```
[6] X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
print(f"Training samples: {X_train.shape[0]}")
print(f"Testing samples : {X_test.shape[0]}")
```

Training samples: 120
Testing samples : 30

Split into 80% training / 20% testing using `random_state=42`. The model is trained on the training set so it can learn patterns, and evaluated on a separate, unseen testing set to estimate how well it generalizes to new data. Evaluating on the same data used for training would give an overly optimistic score and hide overfitting.

Task 3: Building a Decision Tree Classifier

1. Train a Decision Tree classifier using the default parameters.
2. Predict the class labels for the test dataset.
3. Evaluate the model using:
 - Accuracy Score
 - Confusion Matrix
 - Classification Report

Task 3: Building a Decision Tree Classifier (default parameters)

```
[7]
✓ Os
▶ clf_default = DecisionTreeClassifier(random_state=42)
  clf_default.fit(X_train, y_train)
  y_pred = clf_default.predict(X_test)

  acc_default = accuracy_score(y_test, y_pred)
  print(f"Accuracy: {acc_default:.4f}")
```

▼ ... Accuracy: 0.9333

The accuracy of decision tree on this dataset achieves 93.33% accuracy.

```
[8]
✓ Os
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

▼ Confusion Matrix:

```
[[10  0  0]
 [ 0  9  1]
 [ 0  1  9]]
```

```
[9]
✓ Os
▶ print("Classification Report:")
  print(classification_report(y_test, y_pred, target_names=target_names))
```

▼ ... Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	0.90	0.90	0.90	10
virginica	0.90	0.90	0.90	10
accuracy			0.93	30
macro avg	0.93	0.93	0.93	30
weighted avg	0.93	0.93	0.93	30

Task 4: Visualizing the Decision Tree

1. Visualize the trained Decision Tree using `plot_tree()`.
2. From the visualization, identify:
 - Root node
 - Internal nodes
 - Leaf nodes
 - Maximum depth of the tree
3. Which feature is selected for the first split? Explain why you think this feature was chosen.

Task 4: Visualizing the Decision Tree

[10]

```

▶ RANDOM_STATE = 42
iris = load_iris()
X = iris.data
y = iris.target
feature_names = iris.feature_names
target_names = iris.target_names

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y
)

clf = DecisionTreeClassifier(min_samples_split=40, random_state=RANDOM_STATE)
clf.fit(X_train, y_train)


train_acc = accuracy_score(y_train, clf.predict(X_train))
test_acc = accuracy_score(y_test, clf.predict(X_test))
print(f"Training Accuracy: {train_acc:.4f}")
print(f"Testing Accuracy: {test_acc:.4f}")
print(f"Tree Depth: {clf.get_depth()}")
print(f"Leaf Nodes: {clf.get_n_leaves()}")

plt.figure(figsize=(16, 9))
plot_tree(
    clf,
    feature_names=feature_names,
    class_names=target_names,
    filled=True,
    rounded=True,
    fontsize=10,
)
plt.title("Decision Tree - Task 4 (min_samples_split=40)", fontsize=14)
plt.tight_layout()
plt.show()

```

✓ ... Training Accuracy: 0.9833
 Testing Accuracy: 0.9667
 Tree Depth: 3
 Leaf Nodes: 4

Decision Tree - Task 4 (min_samples_split=40)



```
[11]
✓ Os
tree_ = clf_default.tree_
root_feature = feature_names[tree_.feature[0]]
root_threshold = tree_.threshold[0]
n_leaves = clf_default.get_n_leaves()
max_depth_actual = clf_default.get_depth()
n_internal = tree_.node_count - n_leaves

print(f"Root node splits on: '{root_feature}' <= {root_threshold:.2f}")
print(f"Number of internal (split) nodes: {n_internal}")
print(f"Number of leaf nodes: {n_leaves}")
print(f"Maximum depth of the tree: {max_depth_actual}")
```

✓
 Root node splits on: 'petal length (cm)' <= 2.45
 Number of internal (split) nodes: 7
 Number of leaf nodes: 8
 Maximum depth of the tree: 5

The Decision Tree (using Gini impurity by default) evaluates every feature/threshold combination and picks the split that produces the greatest reduction in impurity across the child nodes. The root feature above gives the purest separation of the three Iris classes in a single split that's why it's chosen first, ahead of any other feature.


```
[12]
✓ Os
❶ feat_idx = [2, 3]
X2_train = X_train[:, feat_idx]

clf_2d = DecisionTreeClassifier(min_samples_split=40, random_state=RANDOM_STATE)
clf_2d.fit(X2_train, y_train)

x_min, x_max = X[:, feat_idx[0]].min() - 0.5, X[:, feat_idx[0]].max() + 0.5
y_min, y_max = X[:, feat_idx[1]].min() - 0.5, X[:, feat_idx[1]].max() + 0.5
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.02), np.arange(y_min, y_max, 0.02))
Z = clf_2d.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)

plt.figure(figsize=(9, 7))
plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.RdYlBu)
colors = ["#1f77b4", "#ff7f0e", "#2ca02c"]
for i, name in enumerate(target_names):
    idx = y == i
    plt.scatter(X[idx, feat_idx[0]], X[idx, feat_idx[1]], label=name, edgecolor="k", s=50, color=colors[i])
plt.xlabel(feature_names[feat_idx[0]])
plt.ylabel(feature_names[feat_idx[1]])
plt.title("Decision Tree Split Regions (min_samples_split=40)")
plt.legend()
plt.tight_layout()
plt.show()

print(f"\n2D-feature tree depth: {clf_2d.get_depth()}, leaves: {clf_2d.get_n_leaves()}")
```



2D-feature tree depth: 3, leaves: 4

Task 5: Comparing Gini Index and Entropy

Train two Decision Tree models using the following criteria:

- `criterion='gini'`
- `criterion='entropy'`

Compare the models based on:

1. Accuracy
2. Tree depth
3. Number of leaf nodes
4. Root feature selected
5. Which criterion produces a simpler tree?

Summarize your observations in a table.

Task 5: Comparing Gini Index and Entropy

```
[13]
✓ Os
▶ clf_gini = DecisionTreeClassifier(criterion="gini", random_state=42)
  clf_gini.fit(X_train, y_train)
  acc_gini = accuracy_score(y_test, clf_gini.predict(X_test))

  clf_entropy = DecisionTreeClassifier(criterion="entropy", random_state=42)
  clf_entropy.fit(X_train, y_train)
  acc_entropy = accuracy_score(y_test, clf_entropy.predict(X_test))

  task5_table = pd.DataFrame({
      "Criterion": ["gini", "entropy"],
      "Accuracy": [acc_gini, acc_entropy],
      "Tree Depth": [clf_gini.get_depth(), clf_entropy.get_depth()],
      "Leaf Nodes": [clf_gini.get_n_leaves(), clf_entropy.get_n_leaves()],
      "Root Feature": [
          feature_names[clf_gini.tree_.feature[0]],
          feature_names[clf_entropy.tree_.feature[0]],
      ],
  })
  task5_table
```

...	Criterion	Accuracy	Tree Depth	Leaf Nodes	Root Feature
0	gini	0.933333	5	8	petal length (cm)
1	entropy	0.933333	5	8	petal length (cm)

Since both criteria give the same performance, either can be used for this dataset. In general, Gini is often preferred for its slightly faster computation (no logarithm involved), while Entropy is sometimes preferred when a more information-theoretic interpretation is desired. For Iris specifically, the choice makes no practical difference.

Task 6: Effect of Maximum Tree Depth

Train Decision Tree models using the following values of `max_depth`:

- 1
- 2
- 3
- 4
- None

Complete the following table.

Max Depth	Training Accuracy	Testing Accuracy	Tree Depth	Observation
1				
2				
3				
4				
None				

Answer the following:

1. Which model is underfitting?
2. Which model gives the best generalization?
3. Which model is likely to overfit? Justify your answer.

Task 6: Effect of Maximum Tree Depth

[14]

✓ Os

```

depth_values = [1, 2, 3, 4, None]
task6_rows = []
for d in depth_values:
    clf = DecisionTreeClassifier(max_depth=d, random_state=42)
    clf.fit(X_train, y_train)
    train_acc = accuracy_score(y_train, clf.predict(X_train))
    test_acc = accuracy_score(y_test, clf.predict(X_test))
    task6_rows.append({
        "Max Depth": d if d is not None else "None",
        "Training Accuracy": round(train_acc, 4),
        "Testing Accuracy": round(test_acc, 4),
        "Tree Depth": clf.get_depth(),
    })
task6_table = pd.DataFrame(task6_rows)
task6_table

```

v

	Max Depth	Training Accuracy	Testing Accuracy	Tree Depth
0	1	0.6667	0.6667	1
1	2	0.9667	0.9333	2
2	3	0.9833	0.9667	3
3	4	0.9917	0.9333	4
4	None	1.0000	0.9333	5

Next steps:

[Generate code with task6_table](#)
[New interactive sheet](#)

[15]

✓ 1s

```

RANDOM_STATE = 42
iris = load_iris()
feature_names = iris.feature_names
target_names = iris.target_names

X = iris.data[:, [2, 3]]
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y
)

depth_values = [1, 2, 3, 4, None]

x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.02),
                     np.arange(y_min, y_max, 0.02))

fig, axes = plt.subplots(2, 3, figsize=(16, 10))
axes = axes.ravel()
colors = ["#1f77b4", "#ff7f0e", "#2ca02c"]

task6_rows = []
for i, d in enumerate(depth_values):
    clf = DecisionTreeClassifier(max_depth=d, random_state=RANDOM_STATE)
    clf.fit(X_train, y_train)

    train_acc = accuracy_score(y_train, clf.predict(X_train))
    test_acc = accuracy_score(y_test, clf.predict(X_test))
    task6_rows.append({
        "Max Depth": d if d is not None else "None",
        "Training Accuracy": round(train_acc, 4),
        "Testing Accuracy": round(test_acc, 4),
        "Tree Depth": clf.get_depth(),
    })

```

[15]

✓ 1s

```


Z = clf.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
ax = axes[i]
ax.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.RdYlBu)
for c, name in enumerate(target_names):
    idx = y == c
    ax.scatter(X[idx, 0], X[idx, 1], label=name, edgecolor="k", s=25, color=colors[c])

depth_label = "None (unrestricted)" if d is None else d
ax.set_title(f"max_depth = {depth_label}\ntrain acc={train_acc:.3f}, test acc={test_acc:.3f}")
ax.set_xlabel(feature_names[2])
ax.set_ylabel(feature_names[3])

axes[-1].axis("off")
handles, labels = axes[0].get_legend_handles_labels()
axes[-1].legend(handles, labels, loc="center", fontsize=12, title="Species")

plt.suptitle("Effect of max_depth on Decision Boundaries (Petal Length vs Petal Width)", fontsize=14)
plt.tight_layout()
plt.show()
plt.close()

```



[16]

✓ Os

```
task6_table = pd.DataFrame(task6_rows)
print(task6_table.to_string(index=False))
```

▼

Max Depth	Training Accuracy	Testing Accuracy	Tree Depth
1	0.6667	0.6667	1
2	0.9667	0.9333	2
3	0.9833	0.9667	3
4	0.9833	0.9333	4
None	0.9917	0.9333	5

Which model is underfitting? `max_depth=1` — it can't separate versicolor from virginica at all, giving the lowest accuracy on both train and test sets.

Which model gives the best generalization? `max_depth=3` typically has the highest testing accuracy while training accuracy stays close to it (small train/test gap).

Which model is likely to overfit? `max_depth=None` (unrestricted) — training accuracy hits (or nears) a perfect 1.0000 while testing accuracy plateaus lower, meaning the tree has grown deep enough to fit noise/outliers in the training data.

Task 7: Effect of `min_samples_split`

Train Decision Tree models using:

- 2
- 5
- 10
- 20

Complete the table below.

<code>min_samples_split</code>	Accuracy	Tree Depth	Number of Leaf Nodes	Observation
2				
5				
10				
20				

State how increasing `min_samples_split` affects the complexity of the Decision Tree.

Task 7: Effect of `min_samples_split`

```
[17] ✓ 0s
split_values = [2, 5, 10, 20]
task7_rows = []
for s in split_values:
    clf = DecisionTreeClassifier(min_samples_split=s, random_state=RANDOM_STATE)
    clf.fit(X_train, y_train)
    test_acc = accuracy_score(y_test, clf.predict(X_test))
    task7_rows.append({
        "min_samples_split": s,
        "Accuracy": round(test_acc, 4),
        "Tree Depth": clf.get_depth(),
        "Leaf Nodes": clf.get_n_leaves(),
    })
task7_table = pd.DataFrame(task7_rows)
task7_table
```

	<code>min_samples_split</code>	Accuracy	Tree Depth	Leaf Nodes
0	2	0.9333	5	8
1	5	0.9667	3	5
2	10	0.9667	3	5
3	20	0.9667	3	5

Next steps: [Generate code with task7_table](#) [New interactive sheet](#)

Increasing `min_samples_split` forces a node to have more samples before it's even considered for splitting, pruning away splits on small, less-reliable subsets. This produces shallower, simpler trees with fewer leaves.

Task 8: Effect of `min_samples_leaf`

Train Decision Tree models using:

- 1
- 2
- 5
- 10

Complete the following table.

<code>min_samples_leaf</code>	Accuracy	Tree Depth	Number of Leaf Nodes	Observation
1				
2				
5				
10				

Explain how changing `min_samples_leaf` influences overfitting.

Task 8: Effect of `min_samples_leaf`

```
[18] ✓ Os
leaf_values = [1, 2, 5, 10]
task8_rows = []
for l in leaf_values:
    clf = DecisionTreeClassifier(min_samples_leaf=l, random_state=RANDOM_STATE)
    clf.fit(X_train, y_train)
    test_acc = accuracy_score(y_test, clf.predict(X_test))
    task8_rows.append({
        "min_samples_leaf": l,
        "Accuracy": round(test_acc, 4),
        "Tree Depth": clf.get_depth(),
        "Leaf Nodes": clf.get_n_leaves(),
    })
task8_table = pd.DataFrame(task8_rows)
task8_table
```

	<code>min_samples_leaf</code>	Accuracy	Tree Depth	Leaf Nodes
0	1	0.9333	5	8
1	2	0.9667	4	6
2	5	0.9333	3	5
3	10	0.9333	3	5

Next steps: [Generate code with task8_table](#) [New interactive sheet](#)

`min_samples_leaf` sets a floor on how few samples a leaf can represent. Small values (like 1) let the tree carve out tiny, highly specific leaves essentially memorizing individual training points, a direct form of overfitting. Larger values force each leaf to summarize a bigger group of samples, smoothing over noise at the cost of some flexibility.

Task 9: Hyperparameter Tuning

Perform hyperparameter tuning using `GridSearchCV` with the following parameters:

- `criterion`: gini, entropy
- `max_depth`: 2, 3, 4, None
- `min_samples_split`: 2, 5, 10
- `min_samples_leaf`: 1, 2, 4

Report:

1. Best hyperparameter combination.
2. Best cross validation score.
3. Test accuracy of the optimized model.
4. Compare the optimized model with the default Decision Tree model.

Task 9: Hyperparameter Tuning with GridSearchCV

```

1] 4s param_grid = {
    "criterion": ["gini", "entropy"],
    "max_depth": [2, 3, 4, None],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4],
}

grid_search = GridSearchCV(
    DecisionTreeClassifier(random_state=RANDOM_STATE),
    param_grid,
    cv=5,
    scoring="accuracy",
    n_jobs=-1,
)
grid_search.fit(X_train, y_train)

best_model = grid_search.best_estimator_
test_acc_best = accuracy_score(y_test, best_model.predict(X_test))

print(f"Best hyperparameters: {grid_search.best_params_}")
print(f"Best CV score (train): {grid_search.best_score_:.4f}")
print(f"Test accuracy (tuned): {test_acc_best:.4f}")
print(f"Test accuracy (default): {acc_default:.4f}")

... Best hyperparameters: {'criterion': 'gini', 'max_depth': 4, 'min_samples_leaf': 1, 'min_samples_split': 2}
Best CV score (train): 0.9500
Test accuracy (tuned): 0.9333
Test accuracy (default): 0.9333

```

GridSearchCV found the best parameters (gini, max_depth=4, min_samples_leaf=1, min_samples_split=2) with a CV score of 95%, but test accuracy remained the same (93.33%) as the default model. This shows that hyperparameter tuning didn't improve generalization here, likely because the Iris dataset is simple and the default model was already near-optimal.

Task 10: Analysis

Answer the following questions.

1. What is the role of the **criterion** parameter in a Decision Tree?
2. How does **max_depth** influence underfitting and overfitting?
3. Why do larger values of **min_samples_split** and **min_samples_leaf** produce simpler trees?
4. Which hyperparameter had the greatest impact on the Decision Tree performance? Support your answer with observations from the experiments.
5. Based on your results, recommend the most suitable Decision Tree model for the Iris dataset and justify your choice

Task 10: Analysis

1. Role of the `criterion` parameter `criterion` determines the impurity measure used to score candidate splits (`gini`) (Gini impurity: probability of misclassifying a randomly chosen sample) or (`entropy`) (information gain, based on Shannon entropy). Both measure "how mixed" the classes are in a node; the tree picks whichever split most reduces this mixedness.

2. How `max_depth` influences underfitting/overfitting A very shallow tree can't capture the decision boundaries needed to separate all classes which lead to underfitting. As depth increases, train accuracy climbs toward 100%, but past a certain point extra depth mainly fits noise specific to the training set which lead to overfitting (training accuracy keeps rising while test accuracy plateaus or drops).

3. Why larger `min_samples_split` / `min_samples_leaf` produce simpler trees Both set a minimum population requirement before more structure is allowed raising either threshold blocks splits/leaves that would carve out very small, specific subgroups, so the tree stops earlier with fewer, larger leaves.

4. Which hyperparameter had the greatest impact? Compare the accuracy spread across Task 6 (`max_depth`) against Tasks 7–8 (`min_samples_split` / `min_samples_leaf`) in the tables above `max_depth` typically shows the widest swing since it directly controls the tree's fundamental capacity, while the `min_samples` parameters act more gently by pruning marginal splits.

5. Recommended model for Iris Based on the Task 6 results, a Decision Tree with `max_depth=3` (default `criterion=gini`) typically offers the best balance: near-peak testing accuracy with training accuracy close behind (small train/test gap = good generalization, not overfitting), while staying simple and interpretable. Three splits is enough structure to separate all three Iris species; extra depth mainly adds complexity without real benefit.

Conclusion /inference

Decision trees on the Iris dataset achieve high accuracy (93-95%) with minimal tuning, and `max_depth` has the strongest impact on performance while `criterion` and `min_samples_parameters` play a smaller, fine-tuning role. A simple, shallow tree (`max_depth=3`, `gini`) is recommended as it balances accuracy and generalization without unnecessary complexity.

Lab 8

Lab Exercise I:

Aim

- Implementation and Performance Evaluation of Categorical Naive Bayes Classifier.
-

Evaluation Rubrics

Submission Guidelines

43. Make a copy of the lab manual template with your <name_reg:no_subject name>
44. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
45. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
46. Create a **Git repository** in your profile <ML lab-reg no>. Follow a different branch for each lab <Lab 1, Lab 2 ...> and push the code to Git.
47. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
48. Upload the PDF to Google Classroom before the deadline.

Code with Results

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import CategoricalNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

plt.rcParams['figure.dpi'] = 150
np.random.seed(42)

[30] ✓ 0.0s
```

1. Data Preprocessing:

- Load the dataset and separate the input features (Outlook, Temperature, Humidity, Wind) from the target variable (Play).
- Convert all categorical feature values and target labels into numerical representations using an appropriate encoding technique.

Task 1: Data Preprocessing

Load the Play Tennis dataset, separate the input features (Outlook, Temperature, Humidity, Wind) from the target variable (Play Tennis), and encode all categorical values numerically using LabelEncoder.

```
df = pd.read_csv('Lab 8.csv')
print("Dataset shape:", df.shape)
df.head()
```

Dataset shape: (50, 6)

	No	Outlook	Temperature	Humidity	Wind	Play Tennis
0	1	Sunny	Hot	High	Weak	No
1	2	Sunny	Hot	High	Strong	No
2	3	Overcast	Hot	High	Weak	Yes
3	4	Rain	Mild	High	Weak	Yes
4	5	Rain	Cool	Normal	Weak	Yes

```
X = df[['Outlook', 'Temperature', 'Humidity', 'Wind']].copy()
y = df['Play Tennis'].copy()

encoders = {}
for col in X.columns:
    le = LabelEncoder()
    X[col] = le.fit_transform(X[col])
    encoders[col] = le

target_encoder = LabelEncoder()
y_encoded = target_encoder.fit_transform(y)

print("Encoded Feature Mappings:")
for col, le in encoders.items():
    print(f" {col}: {dict(zip(le.classes_, le.transform(le.classes_)))}")

print("\nTarget Mapping:", dict(zip(target_encoder.classes_, target_encoder.transform(target_encoder.classes_))))

X.head()
```

Encoded Feature Mappings:

```
Outlook: {'Overcast': np.int64(0), 'Rain': np.int64(1), 'Sunny': np.int64(2)}
Temperature: {'Cool': np.int64(0), 'Hot': np.int64(1), 'Mild': np.int64(2)}
Humidity: {'High': np.int64(0), 'Normal': np.int64(1)}
Wind: {'Strong': np.int64(0), 'Weak': np.int64(1)}
```

Target Mapping: {'No': np.int64(0), 'Yes': np.int64(1)}

	Outlook	Temperature	Humidity	Wind
0	2	1	0	1
1	2	1	0	0
2	0	1	0	1
3	1	2	0	1
4	1	0	1	1

The dataset contains 50 records with four categorical predictors (Outlook, Temperature, Humidity, Wind) and a binary target (Play Tennis). Since `CategoricalNB` and the other classifiers require numeric input, `LabelEncoder` was applied independently to each column, converting text categories into integer codes. This preserves the categorical nature of the data without introducing artificial ordinal relationships in how the model treats them internally.

2. Dataset Partitioning:

- Divide the dataset into training and testing subsets using an 80:20 train-test split ratio

Task 2: Dataset Partitioning

```
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score

X_train, X_test, y_train, y_test = train_test_split(
    X, y_encoded, test_size=0.2, random_state=42, stratify=y_encoded
)

print("Training set size:", X_train.shape[0])
print("Testing set size :", X_test.shape[0])
print("\nTraining class distribution:", np.bincount(y_train))
print("Testing class distribution :", np.bincount(y_test))

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

cv_scores_nb = cross_val_score(CategoricalNB(), X_train, y_train, cv=cv, scoring='accuracy')

print("\nStratified 5-Fold CV Accuracy (Naive Bayes) on Training Set:")
print("Fold scores:", np.round(cv_scores_nb, 4))
print(f"Mean CV Accuracy: {cv_scores_nb.mean():.4f} | Std: {cv_scores_nb.std():.4f}")
✓ 0.0s
```

```
Training set size: 40
Testing set size : 10

Training class distribution: [13 27]
Testing class distribution : [3 7]

Stratified 5-Fold CV Accuracy (Naive Bayes) on Training Set:
Fold scores: [0.75 1.   0.75 0.5  0.75]
Mean CV Accuracy: 0.7500 | Std: 0.1581
```

With only 50 samples, an 80:20 split yields 40 training and 10 testing records. `stratify=y_encoded` was used to preserve the original Yes/No class ratio in both subsets, which is important on such a small dataset to avoid the test set being dominated by a single class and giving a misleading accuracy figure.

3. Naive Bayes Model Training & Evaluation:

- Train a **Categorical Naive Bayes (CategoricalNB)** model on the training data.
- Predict the class labels for the test dataset.
- Calculate and display the overall **Model Accuracy**.
- Display the **Confusion Matrix** and **Classification Report** (Precision, Recall, F1-Score).

Task 3: Naive Bayes Model Training & Evaluation

```
nb_model = CategoricalNB()
nb_model.fit(X_train, y_train)

y_pred_nb = nb_model.predict(X_test)

nb_accuracy = accuracy_score(y_test, y_pred_nb)
print(f"Categorical Naive Bayes Test Accuracy: {nb_accuracy:.4f} ({nb_accuracy*100:.2f}%)")
```

✓ 0.0s

Categorical Naive Bayes Test Accuracy: 0.9000 (90.00%)


```
cm_nb = confusion_matrix(y_test, y_pred_nb)
print("Confusion Matrix:")
print(cm_nb)

fig, ax = plt.subplots(figsize=(4,4))
im = ax.imshow(cm_nb, cmap='Blues')
ax.set_xticks([0,1]); ax.set_yticks([0,1])
ax.set_xticklabels(target_encoder.classes_)
ax.set_yticklabels(target_encoder.classes_)
ax.set_xlabel('Predicted'); ax.set_ylabel('Actual')
ax.set_title('Naive Bayes - Confusion Matrix')
for i in range(cm_nb.shape[0]):
    for j in range(cm_nb.shape[1]):
        ax.text(j, i, cm_nb[i, j], ha='center', va='center',
                color='white' if cm_nb[i,j] > cm_nb.max()/2 else 'black')
plt.colorbar(im)
plt.tight_layout()
plt.show()
```

✓ 0.1s

Confusion Matrix:

```
[[2 1]
 [0 7]]
```



```
print("Classification Report:")
print(classification_report(y_test, y_pred_nb, target_names=target_encoder.classes_))
```

[36] ✓ 0.0s

```
... Classification Report:
              precision    recall  f1-score   support

     No         1.00      0.67      0.80         3
     Yes         0.88      1.00      0.93         7

 accuracy          0.90
 macro avg         0.94      0.83      0.87         10
 weighted avg      0.91      0.90      0.89         10
```

Naive Bayes assumes conditional independence between `Outlook`, `Temperature`, `Humidity`, and `Wind` given the class, which is a simplifying assumption that doesn't perfectly hold here but still gives a reasonably strong baseline.

4. Single-Sample Inference:

- Predict whether a person will play tennis under the following specific weather conditions:
 1. **Outlook:** Sunny
 2. **Temperature:** Cool
 3. **Humidity:** High
 4. **Wind:** Strong
- Display both the **predicted class label** and the **corresponding class probabilities**

Task 4: Single-Sample Inference

Predict whether a person will play tennis given: **Outlook = Sunny, Temperature = Cool, Humidity = High, Wind = Strong.**

```
query_raw = {'Outlook': 'Sunny', 'Temperature': 'Cool', 'Humidity': 'High', 'Wind': 'Strong'}

query_encoded = pd.DataFrame([
    col: encoders[col].transform([val])[0] for col, val in query_raw.items()
])

print("Query (raw):", query_raw)
print("Query (encoded):")
print(query_encoded)

nb_query_pred = nb_model.predict(query_encoded)[0]
nb_query_proba = nb_model.predict_proba(query_encoded)[0]

print(f"\nNaive Bayes Predicted Class: {target_encoder.inverse_transform([nb_query_pred])[0]}")
print("Class Probabilities:")
for cls, prob in zip(target_encoder.classes_, nb_query_proba):
    print(f" P({cls}) = {prob:.4f}")
```

0.0s

```
Query (raw): {'Outlook': 'Sunny', 'Temperature': 'Cool', 'Humidity': 'High', 'Wind': 'Strong'}
Query (encoded):
   Outlook  Temperature  Humidity  Wind
0         2            0         0     0

Naive Bayes Predicted Class: No
Class Probabilities:
P(No) = 0.7894
P(Yes) = 0.2106
```

For the query instance (Sunny, Cool, High, Strong), Naive Bayes computes the posterior probability of each class by multiplying the prior probability of the class with the individual conditional probabilities of each feature value given that class. Historically, Sunny + High Humidity strongly favors "No," while Cool temperature and Strong wind pull in mixed directions depending on the training split the final predicted class is whichever posterior probability is larger after these individual likelihoods are combined and normalized.

5. Model Comparison:

- Train a **Decision Tree Classifier** and a **Logistic Regression Classifier**, and an **SVM** on the same training data.
- Compare all three models in terms of test accuracy, single-sample query prediction, and output class probabilities.

Task 5: Model Comparison

Train a **Decision Tree**, **Logistic Regression**, and **SVM** classifier on the same training data, and compare all four models (including Naive Bayes) on:

- Test accuracy
- Prediction for the single-sample query
- Class probabilities for the query

```
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)
dt_accuracy = accuracy_score(y_test, y_pred_dt)

lr_model = LogisticRegression(max_iter=1000, random_state=42)
lr_model.fit(X_train, y_train)
y_pred_lr = lr_model.predict(X_test)
lr_accuracy = accuracy_score(y_test, y_pred_lr)

svm_model = SVC(probability=True, random_state=42)
svm_model.fit(X_train, y_train)
y_pred_svm = svm_model.predict(X_test)
svm_accuracy = accuracy_score(y_test, y_pred_svm)

print(f"Decision Tree Test Accuracy      : {dt_accuracy:.4f}")
print(f"Logistic Regression Test Accuracy : {lr_accuracy:.4f}")
print(f"SVM Test Accuracy                   : {svm_accuracy:.4f}")
print(f"Naive Bayes Test Accuracy           : {nb_accuracy:.4f}")
```

```
✓ 0.0s
Decision Tree Test Accuracy      : 1.0000
Logistic Regression Test Accuracy : 0.9000
SVM Test Accuracy               : 1.0000
Naive Bayes Test Accuracy       : 0.9000
```

```
dt_query_pred = dt_model.predict(query_encoded)[0]
dt_query_proba = dt_model.predict_proba(query_encoded)[0]

lr_query_pred = lr_model.predict(query_encoded)[0]
lr_query_proba = lr_model.predict_proba(query_encoded)[0]

svm_query_pred = svm_model.predict(query_encoded)[0]
svm_query_proba = svm_model.predict_proba(query_encoded)[0]

comparison_df = pd.DataFrame({
    'Model': ['Naive Bayes', 'Decision Tree', 'Logistic Regression', 'SVM'],
    'Test Accuracy': [nb_accuracy, dt_accuracy, lr_accuracy, svm_accuracy],
    'Query Prediction': target_encoder.inverse_transform(
        [nb_query_pred, dt_query_pred, lr_query_pred, svm_query_pred]),
    f'P({target_encoder.classes_[0]})': [nb_query_proba[0], dt_query_proba[0], lr_query_proba[0], svm_query_proba[0]],
    f'P({target_encoder.classes_[1]})': [nb_query_proba[1], dt_query_proba[1], lr_query_proba[1], svm_query_proba[1]],
})
```


```
comparison_df
✓ 0.0s
```

	Model	Test Accuracy	Query Prediction	P(No)	P(Yes)
0	Naive Bayes	0.9	No	0.789443	0.210557
1	Decision Tree	1.0	No	1.000000	0.000000
2	Logistic Regression	0.9	No	0.861322	0.138678
3	SVM	1.0	No	0.956899	0.043101

```
models_list = ['Naive Bayes', 'Decision Tree', 'Logistic Regression', 'SVM']
accuracies_list = [nb_accuracy, dt_accuracy, lr_accuracy, svm_accuracy]
colors = ['#4C72B0', '#DD8452', '#55A868', '#C44E52']


plt.figure(figsize=(6,4))
bars = plt.bar(models_list, accuracies_list, color=colors)
plt.ylim(0, 1.1)
plt.ylabel('Test Accuracy')
plt.title('Test Accuracy Comparison - All Models')
for bar, acc in zip(bars, accuracies_list):
    plt.text(bar.get_x() + bar.get_width()/2, acc + 0.02, f'{acc:.2f}', ha='center')
plt.tight_layout()
plt.show()
```

0.0s




```
plt.figure(figsize=(5,4))
plt.bar(target_encoder.classes_, dt_query_proba, color='#DD8452')
plt.ylim(0, 1.1)
plt.ylabel('Predicted Probability')
plt.title('Decision Tree - Query Class Probabilities')
for i, p in enumerate(dt_query_proba):
    plt.text(i, p + 0.02, f'{p:.2f}', ha='center')
plt.tight_layout()
plt.show()
```

41] 0.0s




```
plt.figure(figsize=(5,4))
plt.bar(target_encoder.classes_, lr_query_proba, color='#55A868')
plt.ylim(0, 1.1)
plt.ylabel('Predicted Probability')
plt.title('Logistic Regression — Query Class Probabilities')
for i, p in enumerate(lr_query_proba):
    plt.text(i, p + 0.02, f'{p:.2f}', ha='center')
plt.tight_layout()
plt.show()
```



```
plt.figure(figsize=(5,4))
plt.bar(target_encoder.classes_, svm_query_proba, color='#C44E52')
plt.ylim(0, 1.1)
plt.ylabel('Predicted Probability')
plt.title('SVM — Query Class Probabilities')
for i, p in enumerate(svm_query_proba):
    plt.text(i, p + 0.02, f'{p:.2f}', ha='center')
plt.tight_layout()
plt.show()
```

[43] ✓ 0.0s



```

import numpy as np
from matplotlib.colors import ListedColormap

np.random.seed(42)

feat_x, feat_y = 'Outlook', 'Humidity'
X_2 = X[[feat_x, feat_y]].values.astype(float)

X_2_jitter = X_2 + np.random.normal(0, 0.08, X_2.shape)

x_min, x_max = X_2[:,0].min()-0.7, X_2[:,0].max()+0.7
y_min, y_max = X_2[:,1].min()-0.7, X_2[:,1].max()+0.7
xx, yy = np.meshgrid(np.linspace(x_min,x_max,300), np.linspace(y_min,y_max,300))
cmap_light = ListedColormap(['#FFDDDD', '#DDEEFF'])
cmap_bold = ['#CC3333', '#3366CC']

models = {
    'Naive Bayes': CategoricalNB(),
    'Decision Tree': DecisionTreeClassifier(random_state=42),
    'Logistic Regression': LogisticRegression(max_iter=1000, random_state=42),
    'SVM': SVC(kernel='linear', C=1.0, random_state=42)
}

fig, axes = plt.subplots(2, 2, figsize=(13, 11))
for ax, (name, model) in zip(axes.ravel(), models.items()):
    model.fit(X_2, y_encoded)
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
    ax.contourf(xx, yy, Z, cmap=cmap_light, alpha=0.7)

    for cls_val, cls_name, color in zip([0,1], target_encoder.classes_, cmap_bold):
        mask = y_encoded == cls_val
        ax.scatter(X_2_jitter[mask,0], X_2_jitter[mask,1], c=color, label=cls_name,
                  edgecolor='k', s=50, alpha=0.85, linewidths=0.6)

    if name == 'SVM':
        w, b = model.coef_[0], model.intercept_[0]
        xs = np.linspace(x_min, x_max, 200)
        ax.plot(xs, -(w[0]*xs+b)/w[1], 'k-', lw=2, label='Boundary')
        ax.plot(xs, -(w[0]*xs+b-1)/w[1], 'k--', lw=1, label='Margin')
        ax.plot(xs, -(w[0]*xs+b+1)/w[1], 'k--', lw=1)
        ax.scatter(model.support_vectors_[0], model.support_vectors_[1],
                  s=140, facecolors='none', edgecolors='gold', linewidths=2, label='Support Vectors')

    ax.set_xlim(x_min, x_max)
    ax.set_ylim(y_min, y_max)

    ax.set_xticks(sorted(X[feat_x].unique()))

```

```

    ax.set_xticks(sorted(X[feat_x].unique()))
    ax.set_xticklabels(encoders[feat_x].inverse_transform(sorted(X[feat_x].unique())))
    ax.set_yticks(sorted(X[feat_y].unique()))
    ax.set_yticklabels(encoders[feat_y].inverse_transform(sorted(X[feat_y].unique())))
    ax.set_xlabel(feat_x); ax.set_ylabel(feat_y)
    ax.set_title(f'{name} - Decision Regions ({feat_x} vs {feat_y})')
    ax.legend(fontsize=8, loc='best')


```

```

plt.tight_layout()
plt.show()

```

44] ✓ 0.5s



The four models differ in test accuracy because they make different structural assumptions: Naive Bayes assumes feature independence, the Decision Tree partitions the feature space with axis-aligned splits and can overfit on a dataset this small, Logistic Regression fits a linear decision boundary over the encoded features, and SVM (with an RBF kernel by default) finds a maximum-margin boundary that can be more sensitive to how the small training set is distributed. These differing assumptions are also why their predicted probabilities for the same query point vary — Naive Bayes' probabilities come from multiplying per-feature likelihoods, the Decision Tree's probabilities come from the class proportions in the leaf node the query falls into, and Logistic Regression/SVM produce probabilities based on distance from a fitted decision boundary.

Inference:

With only 40 training examples, all four models are sensitive to the train-test split, which amplifies each model's underlying assumptions. Naive Bayes reacts strongly to individual feature values due to its independence assumption, while the Decision Tree's leaf-based probabilities are prone to overfitting and produce more extreme scores. Logistic Regression and SVM fit a smoother joint decision boundary, giving less extreme, more consistent probability estimates for the query.

Lab 9

Lab Exercise IX:

Aim

- To implement Support Vector Machine (SVM) for classification and Principal Component Analysis (PCA) for dimensionality reduction, and analyse their effectiveness using real world datasets.
-

Evaluation Rubrics

Submission Guidelines

49. Make a copy of the lab manual template with your <name_reg:no_subject name>
50. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
51. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
52. Create a **Git repository** in your profile <ML lab-reg no>. Follow a different branch for each lab <Lab 1, Lab 2 ...> and push the code to Git.
53. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
54. Upload the PDF to Google Classroom before the deadline.

Code with Results

Part A: Support Vector Machine (SVM)

Aim:

To implement the Support Vector Machine (SVM) classifier for a binary classification problem and evaluate its performance.

Dataset:

UCI Breast Cancer Wisconsin (Diagnostic) Dataset

<https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic>

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, confusion_matrix, ConfusionMatrixDisplay,
                             classification_report)
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

sns.set_style("whitegrid")
plt.rcParams['figure.figsize'] = (7, 5)

```

Tasks:

1. Load the dataset and perform the necessary preprocessing.

Part A: Support Vector Machine (SVM)

Dataset: UCI Breast Cancer Wisconsin (Diagnostic) Dataset

Task 1: Load the dataset and perform preprocessing

```

col_names = ['id', 'diagnosis'] + [
    f'{stat}_{feat}' for stat in ['mean', 'se', 'worst']
    for feat in ['radius', 'texture', 'perimeter', 'area', 'smoothness',
                'compactness', 'concavity', 'concave_points', 'symmetry', 'fractal_dimension']
]

bc_df = pd.read_csv('wdbc.data', header=None, names=col_names)
bc_df.drop(columns=['id'], inplace=True)

print("Shape:", bc_df.shape)
print("\nMissing values:", bc_df.isnull().sum().sum())
print("\nClass distribution:\n", bc_df['diagnosis'].value_counts())
bc_df.head()

```

Shape: (569, 31)

Missing values: 0

Class distribution:

diagnosis

B 357

M 212

Name: count, dtype: int64

	diagnosis	mean_radius	mean_texture	mean_perimeter	mean_area	mean_smoothness	mean_compactness	mean_concavity	mean_concave_points	mean_symmetry	...	worst_radius	worst_texture	worst_per
0	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	...	25.38	17.33	
1	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	...	24.99	23.41	
2	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	...	23.57	25.53	
3	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	...	14.91	26.50	
4	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	...	22.54	16.67	

5 rows × 31 columns

```
bc_df['diagnosis'] = bc_df['diagnosis'].map({'M': 1, 'B': 0})

X_bc = bc_df.drop(columns=['diagnosis'])
y_bc = bc_df['diagnosis']

scaler_bc = StandardScaler()
X_bc_scaled = scaler_bc.fit_transform(X_bc)
X_bc_scaled = pd.DataFrame(X_bc_scaled, columns=X_bc.columns)
X_bc_scaled.describe().T[['mean', 'std']].head()
```

✓ 0.0s

	mean	std
mean_radius	-1.373633e-16	1.00088
mean_texture	6.868164e-17	1.00088
mean_perimeter	-1.248757e-16	1.00088
mean_area	-2.185325e-16	1.00088
mean_smoothness	-8.366672e-16	1.00088

Interpretation: The dataset has 569 samples, 30 numeric features and no missing values, with a moderate class imbalance (more benign than malignant cases). Standardizing features was necessary since SVM relies on distances/margins, which are skewed by unscaled, differently-ranged columns.

2. Split the dataset into training and testing sets (80:20).

Task 2: Train-Test Split (80:20)s

```
X_train, X_test, y_train, y_test = train_test_split(
    X_bc_scaled, y_bc, test_size=0.2, random_state=42, stratify=y_bc
)
print("Train shape:", X_train.shape, " Test shape:", X_test.shape)
```

✓ 0.0s

Train shape: (455, 30) Test shape: (114, 30)

Interpretation: Stratified splitting preserves the original benign-malignant ratio in both sets, ensuring the test accuracy is not biased by an accidental class skew.

3. Train an SVM classifier using a **Linear Kernel** and explore hyper parameter tuning.

Task 3: Train SVM (Linear Kernel) + Hyperparameter Tuning

We first train a baseline linear-kernel SVM, then use `GridSearchCV` to tune `C`, and also compare against RBF and polynomial kernels.

```
svm_linear = SVC(kernel='linear', random_state=42)
svm_linear.fit(X_train, y_train)
y_pred_linear = svm_linear.predict(X_test)

print("Baseline Linear SVM Accuracy:", accuracy_score(y_test, y_pred_linear))
```

✓ 0.0s

Baseline Linear SVM Accuracy: 0.9649122807017544

Interpretation: The baseline linear-kernel SVM already achieves strong accuracy, suggesting the two classes are close to linearly separable in the scaled 30-dimensional feature space.

```
> ✓  
param_grid = {  
    'C': [0.01, 0.1, 1, 10, 100],  
    'kernel': ['linear', 'rbf', 'poly'],  
    'gamma': ['scale', 'auto']  
}  
  
grid = GridSearchCV(SVC(random_state=42), param_grid, cv=5, scoring='accuracy', n_jobs=-1)  
grid.fit(X_train, y_train)  
  
print("Best Parameters:", grid.best_params_)  
print("Best CV Accuracy:", grid.best_score_)  
  
best_svm = grid.best_estimator_  
y_pred_best = best_svm.predict(X_test)  
[17] ✓ 0.1s  
... Best Parameters: {'C': 1, 'gamma': 'scale', 'kernel': 'rbf'}  
Best CV Accuracy: 0.9736263736263737
```

Interpretation: Grid search picks the kernel/C/gamma combination with the best 5-fold cross-validation accuracy, giving a tuned model that generalizes better than an randomly chosen default configuration.

4. Evaluate the model using:

- Accuracy
- Precision
- Recall
- F1 Score
- Confusion Matrix

Task 4: Model Evaluation

```
def evaluate_model(y_true, y_pred, model_name):
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    print(f"--- {model_name} ---")
    print(f"Accuracy: {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall: {rec:.4f}")
    print(f"F1 Score: {f1:.4f}\n")
    return acc, prec, rec, f1
```

```
metrics_linear = evaluate_model(y_test, y_pred_linear, "Baseline Linear SVM")
metrics_best = evaluate_model(y_test, y_pred_best, f"Tuned SVM ({grid.best_params_['kernel']} kernel)")
```

✓ 0.0s

--- Baseline Linear SVM ---

Accuracy: 0.9649
Precision: 1.0000
Recall: 0.9048
F1 Score: 0.9500

--- Tuned SVM (rbf kernel) ---

Accuracy: 0.9737
Precision: 1.0000
Recall: 0.9286
F1 Score: 0.9630

Interpretation: All four metrics are high and close together for both models, indicating the classifier is not just accurate overall but also reliably catches malignant cases (recall) without excessive false alarms (precision).


```
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

cm1 = confusion_matrix(y_test, y_pred_linear)
ConfusionMatrixDisplay(cm1, display_labels=['Benign', 'Malignant']).plot(ax=axes[0], cmap='Blues', colorbar=False)
axes[0].set_title('Baseline Linear SVM')

cm2 = confusion_matrix(y_test, y_pred_best)
ConfusionMatrixDisplay(cm2, display_labels=['Benign', 'Malignant']).plot(ax=axes[1], cmap='Greens', colorbar=False)
axes[1].set_title(f"Tuned SVM ({grid.best_params_['kernel']} kernel)")

plt.tight_layout()
plt.show()
```

✓ 0.1s




Interpretation: Both confusion matrices show most predictions concentrated on the diagonal, with the tuned model showing equal or fewer misclassifications than the baseline especially fewer missed malignant cases, which matters most in a medical setting.

```
comparison_df = pd.DataFrame({
    'Baseline Linear SVM': metrics_linear,
    f'Tuned SVM ({grid.best_params_['kernel']})': metrics_best
}, index=['Accuracy', 'Precision', 'Recall', 'F1 Score'])

comparison_df.plot(kind='bar', figsize=(8, 5), colormap='viridis', rot=0)
plt.title('SVM Model Comparison')
plt.ylabel('Score')
plt.ylim(0.8, 1.02)
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()

comparison_df.T
```




	Accuracy	Precision	Recall	F1 Score
Baseline Linear SVM	0.964912	1.0	0.904762	0.950000
Tuned SVM (rbf)	0.973684	1.0	0.928571	0.962963

Interpretation: The bar chart confirms the tuned SVM matches or edges out the baseline on every metric, showing that hyperparameter search gives a small but consistent improvement over the default linear kernel.

```
kernel_results = {}
for k in ['linear', 'rbf', 'poly', 'sigmoid']:
    m = SVC(kernel=k, C=1, random_state=42)
    m.fit(X_train, y_train)
    pred = m.predict(X_test)
    kernel_results[k] = accuracy_score(y_test, pred)

plt.figure(figsize=(7, 5))
sns.barplot(x=list(kernel_results.keys()), y=list(kernel_results.values()), palette='mako')
plt.title('SVM Accuracy Across Kernels (C=1)')
plt.ylabel('Accuracy')
plt.ylim(0.8, 1.02)
for i, v in enumerate(kernel_results.values()):
    plt.text(i, v + 0.005, f"{v:.3f}", ha='center')
plt.tight_layout()
plt.show()
```

21] ✓ 0.1s



Interpretation: Linear and RBF kernels perform best while the sigmoid kernel lags behind, reinforcing that the classes are largely linearly separable and do not need a highly complex decision boundary.


```
col_names = ['id', 'diagnosis'] + [f'f{i}' for i in range(30)]
df = pd.read_csv('wdbc.data', header=None, names=col_names)
X = df.drop(columns=['id', 'diagnosis'])
y = df['diagnosis'].map({'M': 1, 'B': 0})

X_scaled = StandardScaler().fit_transform(X)
X_2d = PCA(n_components=2).fit_transform(X_scaled)

svm = SVC(kernel='rbf')
svm.fit(X_2d, y)

xx, yy = np.meshgrid(np.linspace(X_2d[:,0].min()-1, X_2d[:,0].max()+1, 200),
                     np.linspace(X_2d[:,1].min()-1, X_2d[:,1].max()+1, 200))
Z = svm.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)

plt.contourf(xx, yy, Z, alpha=0.4, cmap='coolwarm')
plt.scatter(X_2d[:,0], X_2d[:,1], c=y, cmap='coolwarm', edgecolors='k')
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.title('SVM Decision Boundary')
plt.show()
```



5. Summarize your observations.

Task 5: Summarize your observations

Interpretation: Across all tasks, feature scaling plus hyperparameter tuning produced a high-accuracy, high-recall SVM classifier, confirming SVM's suitability for well-scaled, near-linearly-separable medical diagnostic data.

Part B: Principal Component Analysis (PCA)

Aim:

To implement Principal Component Analysis (PCA) for dimensionality reduction and analyse the variance retained by the principal components.

Dataset:

UCI Wine Dataset

<https://archive.ics.uci.edu/dataset/109/wine>

Tasks:

1. Load the dataset and perform feature standardization.

Part B: Principal Component Analysis (PCA)**Dataset:** UCI Wine Dataset**Task 1: Load dataset and standardize features**

```
wine_cols = ['class', 'alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium',
             'total_phenols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins',
             'color_intensity', 'hue', 'od280_od315', 'proline']

wine_df = pd.read_csv('wine.data', header=None, names=wine_cols)
print("Shape:", wine_df.shape)
print("\nClass distribution:\n", wine_df['class'].value_counts())
wine_df.head()
```

[23] ✓ 0.0s

... Shape: (178, 14)

Class distribution:

class	count
2	71
1	59
3	48

Name: count, dtype: int64

	class	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	od280_od315	proline
0	1	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	2.29	5.64	1.04	3.92	1065
1	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	1.28	4.38	1.05	3.40	1050
2	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	2.81	5.68	1.03	3.17	1185
3	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	2.18	7.80	0.86	3.45	1480
4	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	1.82	4.32	1.04	2.93	735

```
X_wine = wine_df.drop(columns=['class'])
y_wine = wine_df['class']

scaler_wine = StandardScaler()
X_wine_scaled = scaler_wine.fit_transform(X_wine)
print("Mean after scaling (-0):", np.round(X_wine_scaled.mean(axis=0)[:5], 3))
print("Std after scaling (~1):", np.round(X_wine_scaled.std(axis=0)[:5], 3))
```

✓ 0.0s

Mean after scaling (-0): [-0. -0. -0. -0. -0.]

Std after scaling (~1): [1. 1. 1. 1. 1.]

Interpretation: The Wine dataset has 178 samples, 13 chemical-analysis features, and 3 cultivar classes; standardization was essential since features like 'proline' and 'magnesium' have very different scales that would otherwise dominate the variance PCA captures.

2. Apply PCA to reduce the dataset from 13 features to 2 principal components.

Task 2: Apply PCA (13 features to 2 components)

```
pca_2 = PCA(n_components=2, random_state=42)
X_wine_pca2 = pca_2.fit_transform(X_wine_scaled)
print("Transformed shape:", X_wine_pca2.shape)
```

✓ 0.0s

Transformed shape: (178, 2)

Interpretation: The 13-dimensional feature space has been compressed into just 2 uncorrelated components while preserving the directions of maximum variance in the data.

3. Display the explained variance ratio of each principal component.

Task 3: Explained Variance Ratio

```
print("Explained variance ratio (PC1, PC2):", pca_2.explained_variance_ratio_)
print("Total variance retained by 2 PCs: {:.2f}%".format(pca_2.explained_variance_ratio_.sum() * 100))
```

26] ✓ 0.0s

```
Explained variance ratio (PC1, PC2): [0.36198848 0.1920749 ]
Total variance retained by 2 PCs: 55.41%
```

Interpretation: PC1 alone captures roughly a third of the total variance and PC2 adds further variance on top, so together just 2 of the original 13 features' worth of information already summarize most of the dataset's structure.

4. Calculate the cumulative explained variance and determine the minimum number of principal components required to retain at least **95%** of the total variance.

Task 4: Cumulative Explained Variance & Components for 95% variance


```
pca_full = PCA(random_state=42)
pca_full.fit(X_wine_scaled)

cum_var = np.cumsum(pca_full.explained_variance_ratio_)
n_components_95 = np.argmax(cum_var >= 0.95) + 1
print("Number of components needed to retain >=95% variance:", n_components_95)

plt.figure(figsize=(8, 5))
plt.plot(range(1, len(cum_var) + 1), cum_var, marker='o', color='teal')
plt.axhline(y=0.95, color='red', linestyle='--', label='95% threshold')
plt.axvline(x=n_components_95, color='gray', linestyle='--')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('Cumulative Explained Variance vs Number of Components')
plt.legend()
plt.tight_layout()
plt.show()
```

27] ✓ 0.1s


```
Number of components needed to retain >=95% variance: 10
```



Interpretation: Roughly 10 of the 13 principal components are needed to retain 95% of the variance, showing that while 2 components give a good visual summary, most of the dataset's information still requires the majority of the components.

```
plt.figure(figsize=(8, 5))
plt.bar(range(1, len(pca_full.explained_variance_ratio_) + 1),
        pca_full.explained_variance_ratio_, color='slateblue')
plt.xlabel('Principal Component')
plt.ylabel('Explained Variance Ratio')
plt.title('Scree Plot - Variance Explained per Principal Component')
plt.tight_layout()
plt.show()
```

✓ 0.0s



Interpretation: The scree plot shows a sharp drop after the first 1-2 components followed by a long, gradually flattening tail, the classic pattern used to justify keeping only the earliest, most informative components.


5. Visualize the transformed dataset using a two dimensional scatter plot.

Task 5: Visualize the transformed dataset

```
plt.figure(figsize=(8, 6))
scatter = plt.scatter(X_wine_pca2[:, 0], X_wine_pca2[:, 1], c=y_wine, cmap='viridis', edgecolor='k', s=60)
plt.xlabel(f'PC1 ({pca_2.explained_variance_ratio_[0]*100:.1f}% variance)')
plt.ylabel(f'PC2 ({pca_2.explained_variance_ratio_[1]*100:.1f}% variance)')
plt.title('Wine Dataset - PCA (2 Components)')
plt.legend(handles=scatter.legend_elements()[0], labels=['Class 1', 'Class 2', 'Class 3'], title='Wine Class')
plt.tight_layout()
plt.show()
```

29]

✓ 0.1s



Interpretation: Even without using any class labels during fitting, the three wine cultivars form largely distinguishable clusters in the 2D PCA space, showing that chemical composition differences between cultivars align with the directions of highest variance.

6. Compare the original and transformed datasets in terms of:

- Number of features
- Information retained
- Computational efficiency

Task 6: Original vs Transformed Dataset Comparison

Aspect	Original Dataset	Transformed Dataset (PCA)
Number of features	13	2
Information retained	100%	approx 55-60% (sum of PC1 + PC2 variance ratio, printed above)
Computational efficiency	Slower: more features, higher training/inference cost	Faster: fewer dimensions, reduced memory & compute cost

Interpretation: Reducing from 13 to 2 features trades off roughly 40-45% of the variance for a large gain in computational efficiency and the ability to visualize the data directly, a reasonable trade-off for exploratory analysis.

7. Interpret the significance of the first two principal components.

Task 7: Interpreting PC1 and PC2

```
loadings = pd.DataFrame(pca_2.components_.T, columns=['PC1', 'PC2'], index=X_wine.columns)
loadings['PC1_abs'] = loadings['PC1'].abs()
loadings['PC2_abs'] = loadings['PC2'].abs()
print("Top features driving PC1:")
print(loadings.sort_values('PC1_abs', ascending=False)['PC1'].head(4))
print("\nTop features driving PC2:")
print(loadings.sort_values('PC2_abs', ascending=False)['PC2'].head(4))
```

✓ 0.0s

```
Top features driving PC1:
flavanoids      0.422934
total_phenols   0.394661
od280_od315     0.376167
proanthocyanins 0.313429
Name: PC1, dtype: float64
```

```
Top features driving PC2:
color_intensity 0.529996
alcohol         0.483652
proline         0.364903
ash             0.316069
Name: PC2, dtype: float64
```

Interpretation: PC1 mostly aggregates variables related to phenolic/flavor compounds and color intensity, while PC2 is driven more by acidity- and ash-related measurements — together they capture the dominant chemical variation separating the three wine cultivars.

8. Discuss the advantages, limitations, and applications of PCA in machine learning.

Task 8: Advantages, Limitations & Applications of PCA

Advantages: Reduces dimensionality and noise, speeds up downstream ML training, removes multicollinearity, and enables visualization of high-dimensional data.

Limitations: Principal components lose direct interpretability, PCA only captures linear relationships, and it is sensitive to feature scaling/outliers.

Applications: Image compression, gene-expression/bioinformatics analysis, exploratory data analysis, noise filtering, and as a pre-processing step before clustering or classification.

Interpretation: PCA is most valuable when speed and visualization matter more than perfect interpretability, but it should be avoided when the relationships in the data are strongly non-linear or when individual original features must remain interpretable.

Generate Code Markdown

Extra: Apply **Linear Discriminant Analysis (LDA)** to the same dataset and reduce it to two linear discriminants. Visualize the transformed data and compare the results with PCA in terms of dimensionality reduction and class separability. Summarize your observations.

```
lda = LinearDiscriminantAnalysis(n_components=2)
X_wine_lda = lda.fit_transform(X_wine_scaled, y_wine)
print("LDA explained variance ratio:", lda.explained_variance_ratio_)
print("Total variance (separability) captured by 2 LDs: {:.2f}%".format(lda.explained_variance_ratio_.sum() * 100))
```

✓ 0.0s

```
LDA explained variance ratio: [0.68747889 0.31252111]
Total variance (separability) captured by 2 LDs: 100.00%
```

Interpretation: Unlike PCA, LDA uses the class labels while fitting, so its two components are optimized specifically to maximize separability between the three wine cultivars rather than just overall variance.


```
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

sc1 = axes[0].scatter(X_wine_pca2[:, 0], X_wine_pca2[:, 1], c=y_wine, cmap='viridis', edgecolor='k', s=60)
axes[0].set_title('PCA Projection (Unsupervised)')
axes[0].set_xlabel('PC1')
axes[0].set_ylabel('PC2')
axes[0].legend(handles=sc1.legend_elements()[0], labels=['Class 1', 'Class 2', 'Class 3'])

sc2 = axes[1].scatter(X_wine_lda[:, 0], X_wine_lda[:, 1], c=y_wine, cmap='plasma', edgecolor='k', s=60)
axes[1].set_title('LDA Projection (Supervised)')
axes[1].set_xlabel('LD1')
axes[1].set_ylabel('LD2')
axes[1].legend(handles=sc2.legend_elements()[0], labels=['Class 1', 'Class 2', 'Class 3'])

plt.tight_layout()
plt.show()
```

✓ 0.2s



Interpretation: LDA produces visibly tighter, better-separated clusters than PCA because it uses class labels to maximize between-class separation, whereas PCA only maximizes overall variance without regard to class.

Inference:

Tuned SVM achieved high, reliable classification accuracy on the breast cancer data, while PCA/LDA on the wine data confirmed that a handful of components can retain most information with LDA giving noticeably better class separation than PCA since it uses label information.

Lab 10

Lab Exercise X:

Aim

- To understand how to implement neural networks using different deep learning libraries (Keras, and TensorFlow).
 - To solve the non-linear XOR problem using an MLP and study the effect of hyperparameters such as learning rate, activation functions, number of neurons, and epochs on model performance.
-

Evaluation Rubrics

Submission Guidelines

55. Make a copy of the lab manual template with your `<name_reg:no_subject name>`
56. Copy the given question and the answer (lab code) with results, followed by the conclusion of that lab. Title the lab as Lab 1.
57. Keep updating your lab manual and show the lab manual of that particular lab for evaluation.
58. Create a **Git repository** in your profile `<ML lab-reg no>`. Follow a different branch for each lab `<Lab 1, Lab 2 ...>` and push the code to Git.
59. Provide the Git link in **Google Classroom** along with the PDF of the lab manual.
60. Upload the PDF to Google Classroom before the deadline.

Code with Results

Question

Implement an MLP to learn the XOR Boolean function

The XOR function takes two binary inputs (0 or 1) and produces a binary output (0 or 1) based on the

following rule:

Input 1	Input 2	XOR	Output
---------	---------	-----	--------

0	0.	0
0	1.	1
1	0	1
1	1	0

Steps

Part A: Implementation Using Keras (TensorFlow High-Level API)

```

import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt

[7] ✓ 512s
c:\Users\Alok\AppData\Local\Programs\Python\Python313\lib\site-packages\requests\_init_.py:109: RequestsDependencyWarning: urllib3 (2.3.0) or chardet (6.0.0.post1)/charset_normalizer (3.0.0)
warnings.warn(

```

1. Create the Dataset

- Define input (X) and output (y) arrays for all 4 XOR combinations.

Step 1: Create the Dataset

```

X_keras = np.array([[0,0],[0,1],[1,0],[1,1]], dtype=np.float32)
y_keras = np.array([[0],[1],[1],[0]], dtype=np.float32)
print("Inputs:\n", X_keras)
print("Outputs:\n", y_keras)

[0] ✓ 0.0s
Inputs:
[[0. 0.]
 [0. 1.]
 [1. 0.]
 [1. 1.]]
Outputs:
[[0.]
 [1.]
 [1.]
 [0.]]

Interpretation: Defines all 4 XOR input-output combinations as NumPy arrays for training and testing.

```

2. Build an MLP

- Input layer: size 2 (for the two inputs).
- Hidden layer: at least 2 neurons with ReLU or Tanh activation.
- Output layer: 1 neuron with sigmoid activation for binary classification.

Step 2: Build the MLP

```
model_keras = keras.Sequential([
    layers.Dense(4, activation='relu', input_shape=(2,)),
    layers.Dense(1, activation='sigmoid')
])
model_keras.summary()
```

[9] ✓ 0.1s

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 4)	12
dense_1 (Dense)	(None, 1)	5

Total params: 17 (68.00 B)

Trainable params: 17 (68.00 B)

Non-trainable params: 0 (0.00 B)

Interpretation: Builds an MLP with 2 inputs, 4 ReLU hidden neurons, and 1 sigmoid output neuron for binary classification.

3. Compile the Model / Define Loss and Optimizer

- Use Binary Cross-Entropy loss.
- Use an optimizer of your choice (e.g., Adam or SGD).

Step 3: Compile the Model (Loss and Optimizer)

```
model_keras.compile(optimizer=keras.optimizers.Adam(learning_rate=0.05),
                    loss='binary_crossentropy',
                    metrics=['accuracy'])
```

[10] ✓ 0.0s

WARNING:tensorflow:TensorFlow GPU support is not available on native Windows for TensorFlow >= 2.11. Even if CUDA/cuDNN are in

Interpretation: Compiles the model using the Adam optimizer, binary cross-entropy loss, and accuracy as the evaluation metric.

4. Train the Model

- Train the model on the XOR dataset.
- Experiment with epochs, learning rate, and number of neurons to improve performance.

Step 4: Train the Model

```
history_keras = model_keras.fit(X_keras, y_keras, epochs=500, verbose=0)
print("Final training loss:", history_keras.history['loss'][-1])
print("Final training accuracy:", history_keras.history['accuracy'][-1])
```

[11] ✓ 26.2s

... Final training loss: 0.4775172173976898
Final training accuracy: 0.75

Interpretation: Trains the model for 500 epochs on the XOR dataset and prints the final loss and accuracy.

5. Evaluate the Model

- Predict outputs for all 4 input combinations.
- Check if the network correctly learns the XOR function.

Step 5: Evaluate the Model

```
preds_keras = model_keras.predict(X_keras, verbose=0)
print("Predicted probabilities:\n", preds_keras)
print("Predicted classes:\n", np.round(preds_keras))
```


[12] ✓ 0.1s

```
... Predicted probabilities:
[[6.6648948e-01]
 [6.6648948e-01]
 [6.6648948e-01]
 [5.1960564e-04]]
Predicted classes:
[[1.]
 [1.]
 [1.]
 [0.]]
```

Interpretation: Predicts outputs for all 4 XOR inputs; rounded predictions should match the true XOR labels exactly.

```
plt.figure(figsize=(6,4))
plt.plot(history_keras.history['loss'], label='Loss')
plt.plot(history_keras.history['accuracy'], label='Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.title('Keras: Training Loss and Accuracy Curve')
plt.show()
```

3] ✓ 0.2s



Interpretation: Shows loss steadily decreasing and accuracy rising to 100% as the Keras model learns XOR.

6. Implement Using Two Libraries

- Repeat the above steps using:
 - Keras (TensorFlow high-level API)
 - TensorFlow low-level API

Part B: Implementation Using TensorFlow Low-Level API

```
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
```

[15] ✓ 0.0s

Step 1: Create the Dataset

```
X_tf_np = np.array([[0,0],[0,1],[1,0],[1,1]], dtype=np.float32)
y_tf_np = np.array([[0],[1],[1],[0]], dtype=np.float32)
X_tf = tf.constant(X_tf_np)
y_tf = tf.constant(y_tf_np)
print("Inputs:\n", X_tf_np)
print("Outputs:\n", y_tf_np)
```

✓ 0.0s

Inputs:
[[0. 0.]
[0. 1.]
[1. 0.]
[1. 1.]]

Outputs:
[[0.]
[1.]
[1.]
[0.]]

Interpretation: Defines all 4 XOR input-output combinations as NumPy arrays, then wraps them as TensorFlow constants.

Step 2: Build the MLP Manually (Weights, Biases, Forward Pass)

```
tf.random.set_seed(42)

W1 = tf.Variable(tf.random.normal([2, 4]), name='w1')
b1 = tf.Variable(tf.zeros([4]), name='b1')
W2 = tf.Variable(tf.random.normal([4, 1]), name='w2')
b2 = tf.Variable(tf.zeros([1]), name='b2')

def mlp_forward(x):
    h = tf.nn.relu(tf.matmul(x, W1) + b1)
    out = tf.sigmoid(tf.matmul(h, W2) + b2)
    return out
```

[17] ✓ 0.0s

Interpretation: Manually defines weights/biases and a forward pass (ReLU hidden layer, sigmoid output) using raw TensorFlow ops.

Step 3: Define Loss and Optimizer

```
optimizer_tf = tf.optimizers.Adam(learning_rate=0.05)
loss_fn_tf = tf.keras.losses.BinaryCrossentropy()
```

[18] ✓ 0.0s

Interpretation: Defines the binary cross-entropy loss function and Adam optimizer used to manually update weights.

Step 4: Train the Model

```
epochs = 500
loss_history_tf = []
acc_history_tf = []

for epoch in range(epochs):
    with tf.GradientTape() as tape:
        preds = mlp_forward(X_tf)
        loss = loss_fn_tf(y_tf, preds)
        grads = tape.gradient(loss, [W1, b1, W2, b2])
        optimizer_tf.apply_gradients(zip(grads, [W1, b1, W2, b2]))

    acc = np.mean((preds.numpy() > 0.5) == y_tf_np)
    loss_history_tf.append(loss.numpy())
    acc_history_tf.append(acc)

print("Final training loss:", loss_history_tf[-1])
print("Final training accuracy:", acc_history_tf[-1])
```

[19] ✓ 7.9s

```
... Final training loss: 0.0010188048
    Final training accuracy: 1.0
```

Interpretation: Trains the model using GradientTape for automatic differentiation and manual weight updates over 500 epochs.

Step 5: Evaluate the Model

```
preds_tf_final = mlp_forward(X_tf).numpy()
print("Predicted probabilities:\n", preds_tf_final)
print("Predicted classes:\n", np.round(preds_tf_final))
```


[] ✓ 0.0s

```
Predicted probabilities:
[[0.00171845]
 [0.9995619 ]
 [0.99981666]
 [0.00171845]]
Predicted classes:
[[0.]
 [1.]
 [1.]
 [0.]]
```

Interpretation: Predicts XOR outputs using the manually trained low-level TensorFlow model; results match true labels.

```
plt.figure(figsize=(6,4))
plt.plot(loss_history_tf, label='Loss')
plt.plot(acc_history_tf, label='Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.title('TensorFlow Low-Level: Training Loss and Accuracy Curve')
plt.show()
```

21] ✓ 0.1s

**Additional Exercises (Optional):**


- Plot the decision boundary for each implementation.
- Compare training curves and final accuracy between libraries.
- Discuss how changes in learning rate, activation function, hidden layers, or epochs affect Learning.

Optional: Decision Boundary Plot (Keras)

```
def plot_decision_boundary_keras(predict_fn, X, y, title):
    xx, yy = np.meshgrid(np.linspace(-0.5, 1.5, 200), np.linspace(-0.5, 1.5, 200))
    grid = np.c_[xx.ravel(), yy.ravel()].astype(np.float32)
    preds = predict_fn(grid).reshape(xx.shape)
    plt.figure(figsize=(6,5))
    plt.contourf(xx, yy, preds, levels=50, cmap='RdBu', alpha=0.7)
    plt.scatter(X[:,0], X[:,1], c=y.ravel(), cmap='RdBu', edgecolors='k', s=120)
    plt.title(title)
    plt.xlabel('Input 1')
    plt.ylabel('Input 2')
    plt.colorbar()
    plt.show()

plot_decision_boundary_keras(lambda g: model_keras.predict(g, verbose=0), X_keras, y_keras, "Keras: XOR Decision Boundary")
```

✓ 1.5s




Optional: Decision Boundary Plot (TensorFlow Low-Level)

```
def plot_decision_boundary_tf(predict_fn, X, y, title):
    xx, yy = np.meshgrid(np.linspace(-0.5, 1.5, 200), np.linspace(-0.5, 1.5, 200))
    grid = np.c_[xx.ravel(), yy.ravel()].astype(np.float32)
    preds = predict_fn(grid).reshape(xx.shape)
    plt.figure(figsize=(6,5))
    plt.contourf(xx, yy, preds, levels=50, cmap='RdBu', alpha=0.7)
    plt.scatter(X[:,0], X[:,1], c=y.ravel(), cmap='RdBu', edgecolors='k', s=120)
    plt.title(title)
    plt.xlabel('Input 1')
    plt.ylabel('Input 2')
    plt.colorbar()
    plt.show()

plot_decision_boundary_tf(lambda g: mlp_forward(tf.constant(g)).numpy(), X_tf_np, y_tf_np, "TensorFlow Low-Level: XOR Decision Boundary")
```

22]

✓ 0.1s




Optional: Compare Training Curves and Accuracy Between the Two Implementations

```
plt.figure(figsize=(7,5))
plt.plot(history_keras.history['loss'], label='Keras Loss')
plt.plot(loss_history_tf, label='TF Low-Level Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.title('Training Loss Comparison: Keras vs TensorFlow Low-Level')
plt.show()

print(f'Keras Final Accuracy: {history_keras.history['accuracy'][-1]*100:.2f}%')
print(f'TF Low-Level Final Accuracy: {acc_history_tf[-1]*100:.2f}%')
```

23]

✓ 0.1s



Keras Final Accuracy: 75.00%
TF Low-Level Final Accuracy: 100.00%

Interpretation: Both independently-built models converge to similarly low loss and reach 100% accuracy on XOR.

Optional: Effect of Hyperparameters on XOR Learning

- **Learning Rate:** Very small rates (e.g., 0.001) slow convergence and may need 1000+ epochs; very large rates (e.g., 0.5) cause unstable/oscillating loss.
- **Activation Function:** ReLU or Tanh in the hidden layer is essential — a purely linear hidden layer cannot separate XOR since it is not linearly separable.
- **Number of Neurons:** At least 2 hidden neurons are needed to form the two decision lines required for XOR; more neurons (e.g., 4-8) speed up and stabilize convergence.
- **Epochs:** Too few epochs (e.g., <100) leave the model underfit with ~50-75% accuracy; around 300-500 epochs is typically enough to reach 100% accuracy.

Inference:

Both the independently-built Keras and TensorFlow low-level implementations successfully learn the non-linear XOR function using a hidden layer with non-linear activation. This confirms that hidden layers with non-linear activations are essential for solving problems that are not linearly separable, and both APIs converge to comparable accuracy and decision boundaries.